

Modeling Non-Determinism in HPC Applications

Dylan Chapp

Advisor: Michela Taufer



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

BIG ORANGE. BIG IDEAS.[®]

Non-Determinism and Correctness in HPC

- **HPC Community Position:** “Non-determinism control” and “anomaly detection” identified in DOE report on the 2017 HPC Correctness Summit as key challenges in bug detection and localization [1]

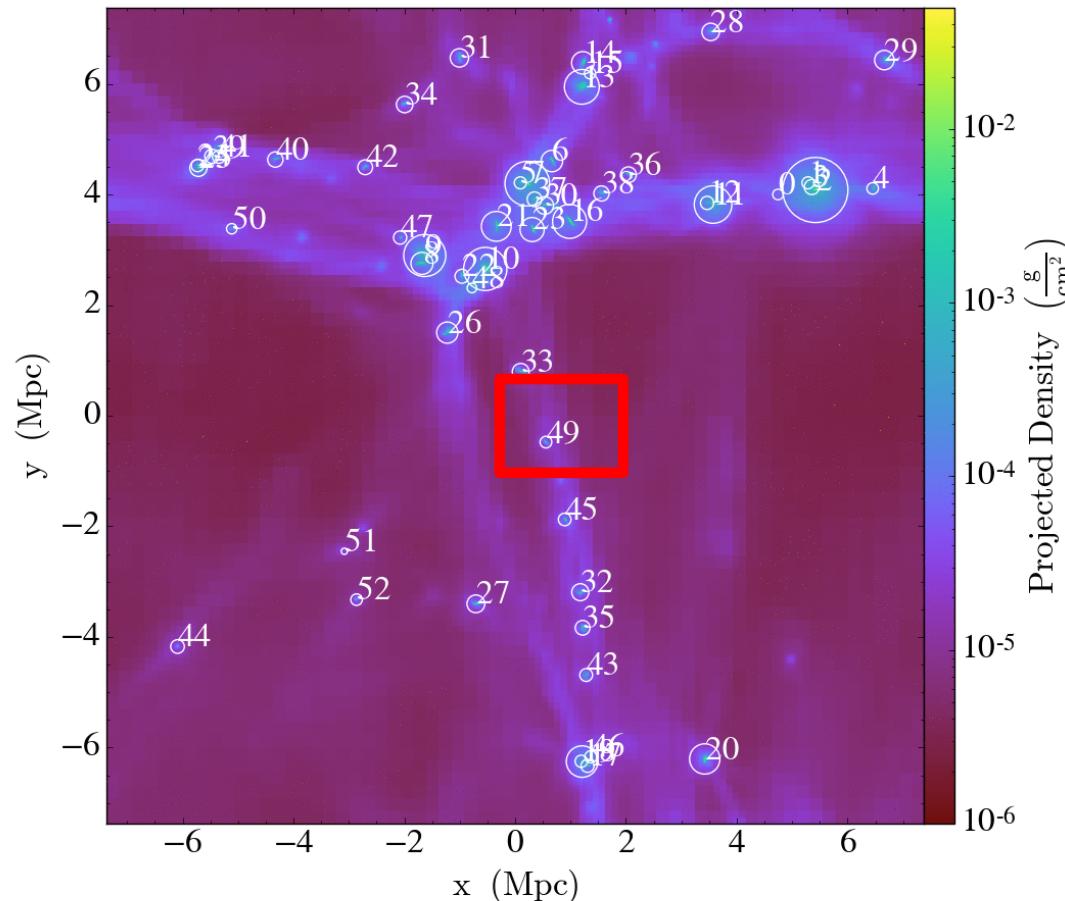


Non-Determinism and Correctness in HPC

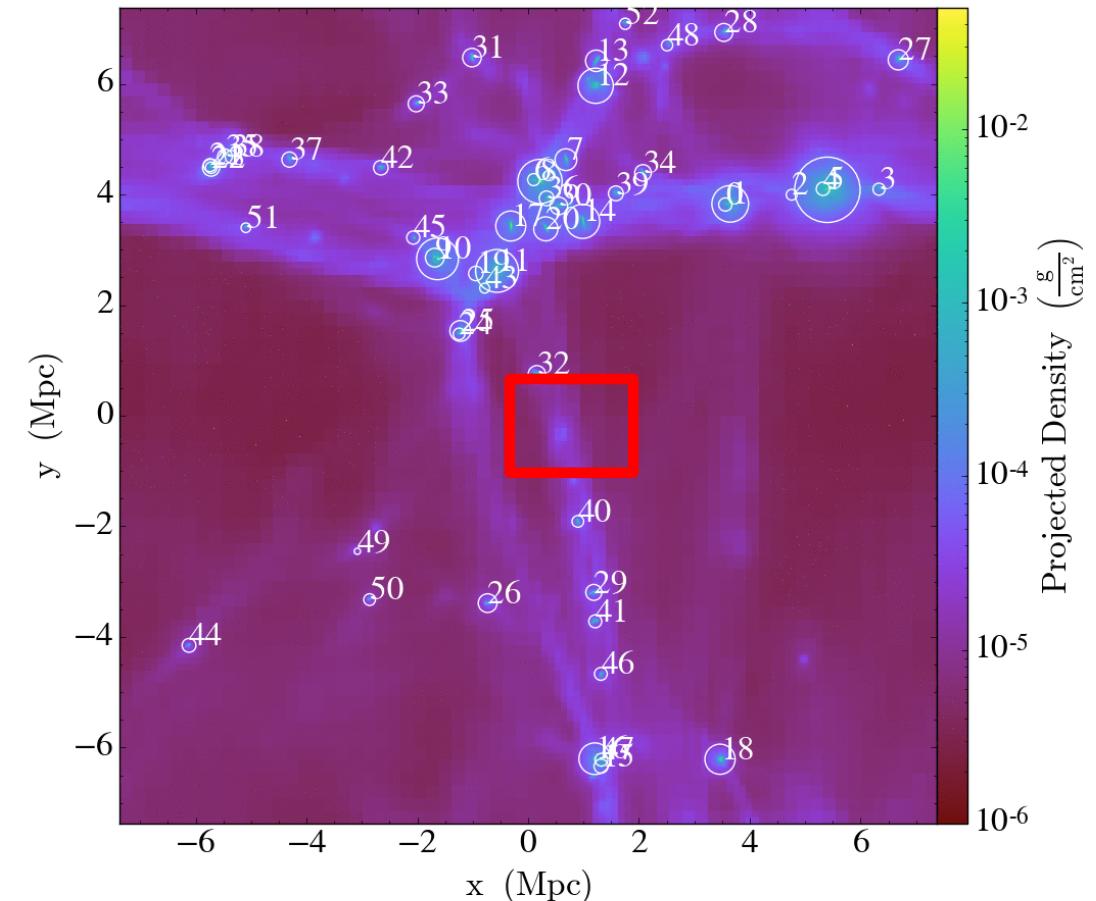
- **HPC Community Position:** “Non-determinism control” and “anomaly detection” identified in DOE report on the 2017 HPC Correctness Summit as key challenges in bug detection and localization [1]
- **Our Position:**
 - These challenges go hand-in-hand.
 - Detecting when and how applications act non-deterministically **in anomalous ways** is critical



Impacts of Non-Determinism on Scientific Outcomes



Run where galactic halo **was detected**



Run where galactic halo was **not detected**



Interaction Between Non-Associativity and Non-Determinism

$$a = 10^9, b = -10^9, c = 10^{-9}$$

Summation order 1

$$(a + b) + c = (10^9 - 10^9) + 10^{-9} = 10^{-9}$$

Summation order 2

$$a + (b + c) = 10^9 + (-10^9 + 10^{-9}) = 0$$

Interaction Between Non-Associativity and Non-Determinism

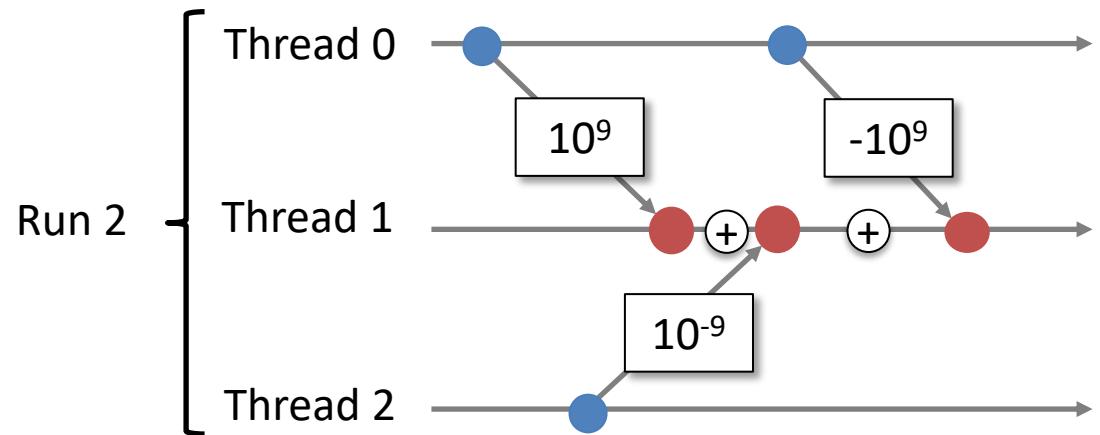
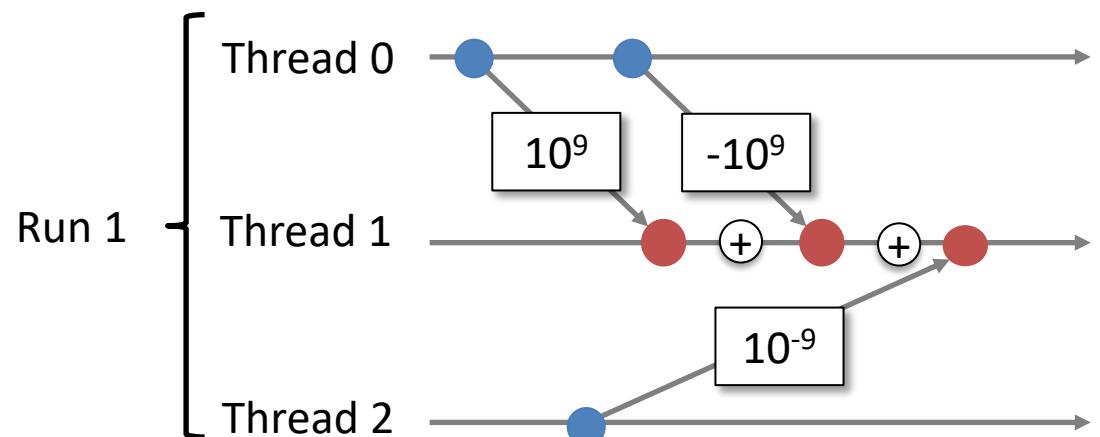
$$a = 10^9, b = -10^9, c = 10^{-9}$$

Summation order 1

$$(a + b) + c = (10^9 - 10^9) + 10^{-9} = 10^{-9}$$

Summation order 2

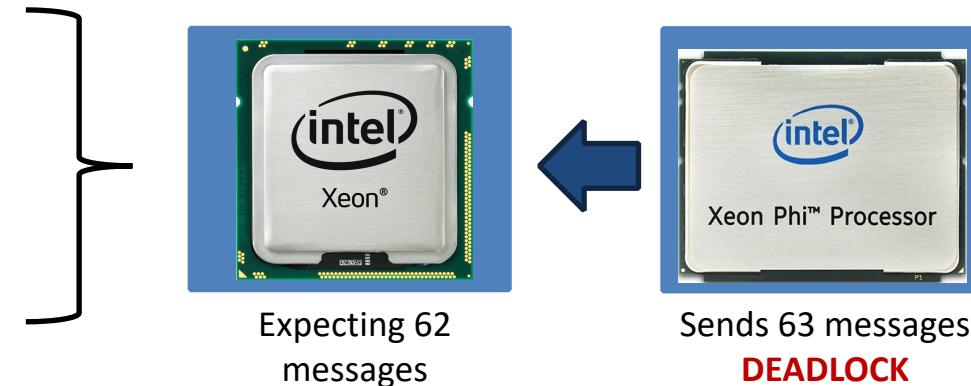
$$a + (b + c) = 10^9 + (-10^9 + 10^{-9}) = 0$$



Impacts of Non-Determinism on Correctness

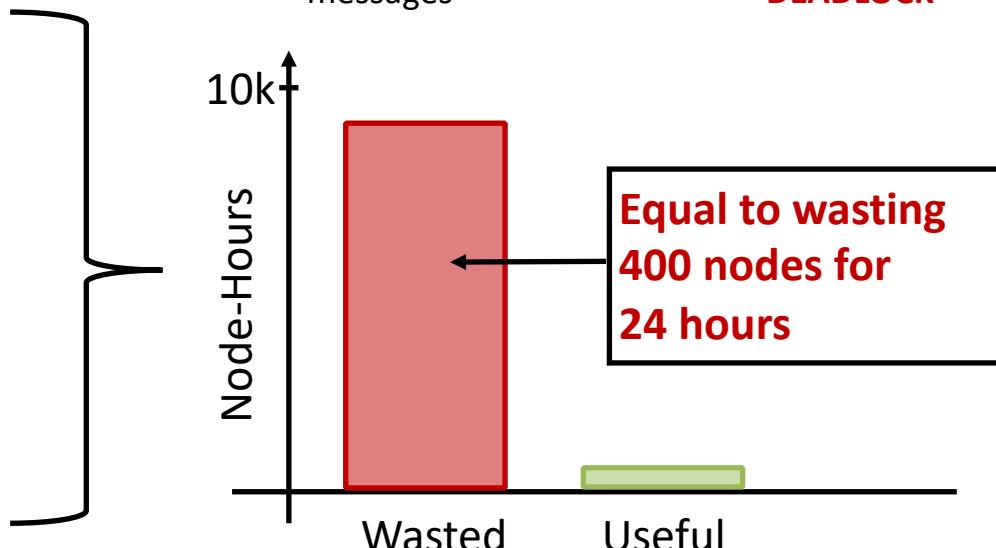
Example 1:

- Rounding difference between Xeon CPU and Xeon Phi caused message count to differ on CPU code vs. accelerator code [1]
- Message count difference induced deadlock



Example 2:

- A non-deterministic bug in Diablo/HYPRE 2.10.1 [2]
- Application hung after several hours, in approximately 1/50 runs
- Cost of debugging effort:
 - **18 months of scientists' time**
 - **9560 node-hours**



Summary of Challenges

- HPC applications involve the interaction between:
 - Floating-point non-associativity
 - Communication non-determinism

Summary of Challenges

- HPC applications involve the interaction between:
 - Floating-point non-associativity
 - Communication non-determinism
- Impacts can range from program incorrectness to irreproducibility of scientific outputs

Summary of Challenges

- HPC applications involve the interaction between:
 - Floating-point non-associativity
 - Communication non-determinism
- Impacts can range from program incorrectness to irreproducibility of scientific outputs
- Mitigation strategies exist, but can be costly, limited in scope, and fail to address a key need:

Summary of Challenges

- HPC applications involve the interaction between:
 - Floating-point non-associativity
 - Communication non-determinism
- Impacts can range from program incorrectness to irreproducibility of scientific outputs
- Mitigation strategies exist, but can be costly, limited in scope, and fail to address a key need:

Linking observable non-determinism to its root causes

Our Approach

- Need to link observations of potentially harmful non-determinism to potential root causes

Our Approach

- Need to link observations of potentially harmful non-determinism to potential root causes
- Need to distinguish between anomalous non-deterministic application behaviors and expected ones

Our Approach

- Need to link observations of potentially harmful non-determinism to potential root causes
- Need to distinguish between anomalous non-deterministic application behaviors and expected ones
- Need a metric for execution similarity

Our Approach

- Need to link observations of potentially harmful non-determinism to potential root causes
- Need to distinguish between anomalous non-deterministic application behaviors and expected ones
- Need a metric for execution similarity
- Need a model of executions that supports such a metric

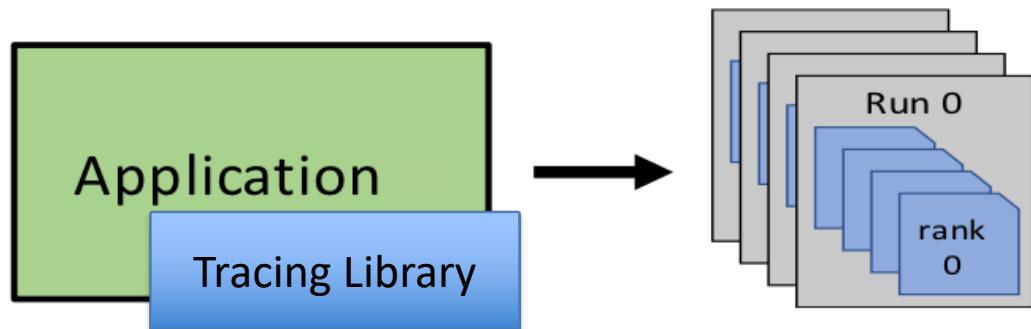
Workflow for Non-Determinism Characterization

- **Phase 1:** Build graph-structured models of executions
- **Phase 2:** Quantify cross-execution trends in non-deterministic communication via graph similarity
- **Phase 3:** Detect periods of anomalous execution dissimilarity and localize potential root causes

Workflow for Non-Determinism Characterization

- **Phase 1:** Build graph-structured models of executions
- **Phase 2:** Quantify cross-execution trends in non-deterministic communication via graph similarity
- **Phase 3:** Detect periods of anomalous execution dissimilarity and localize potential root causes

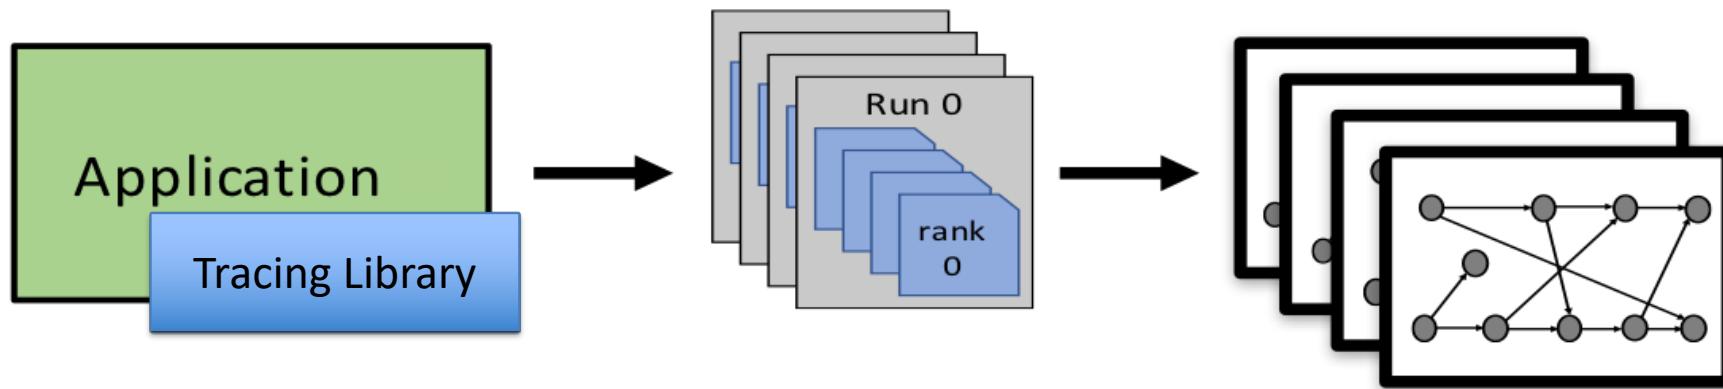
From Traces to Event Graphs



- We trace a non-deterministic application multiple times, capturing a record of communication events



From Traces to Event Graphs



- We trace a non-deterministic application multiple times, capturing a record of communication events
- We convert the set of trace files to a graph-structured model of the inter-process communication that occurred during the execution—i.e., the event graph [3]

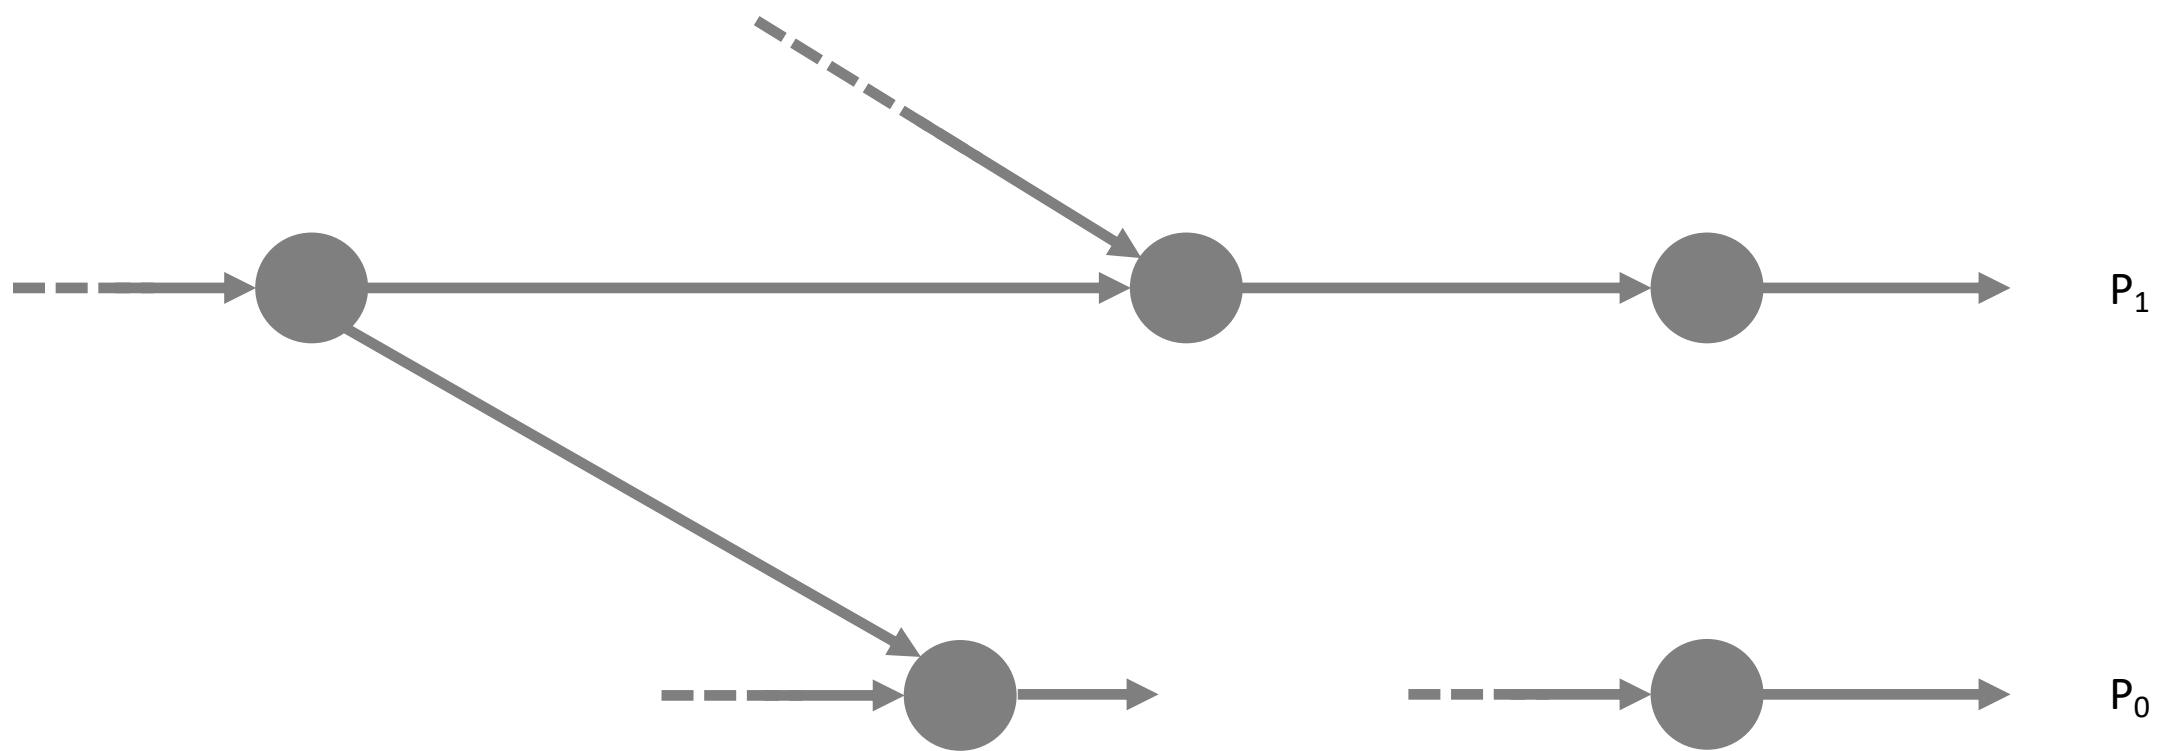
1. Rodrigues, A.F., Voskuilen, G.R., Hammond, S.D. and Hemmert, K.S., 2016. *Structural Simulation Toolkit (SST)* (No. SAND2016-3693PE). Sandia National Lab.(SNL-NM), Albuquerque, NM (United States).

2. <https://github.com/sstsimulator/sst-dump>

3. Kranzlmüller, D., 2000. *Event graph analysis for debugging massively parallel programs*.

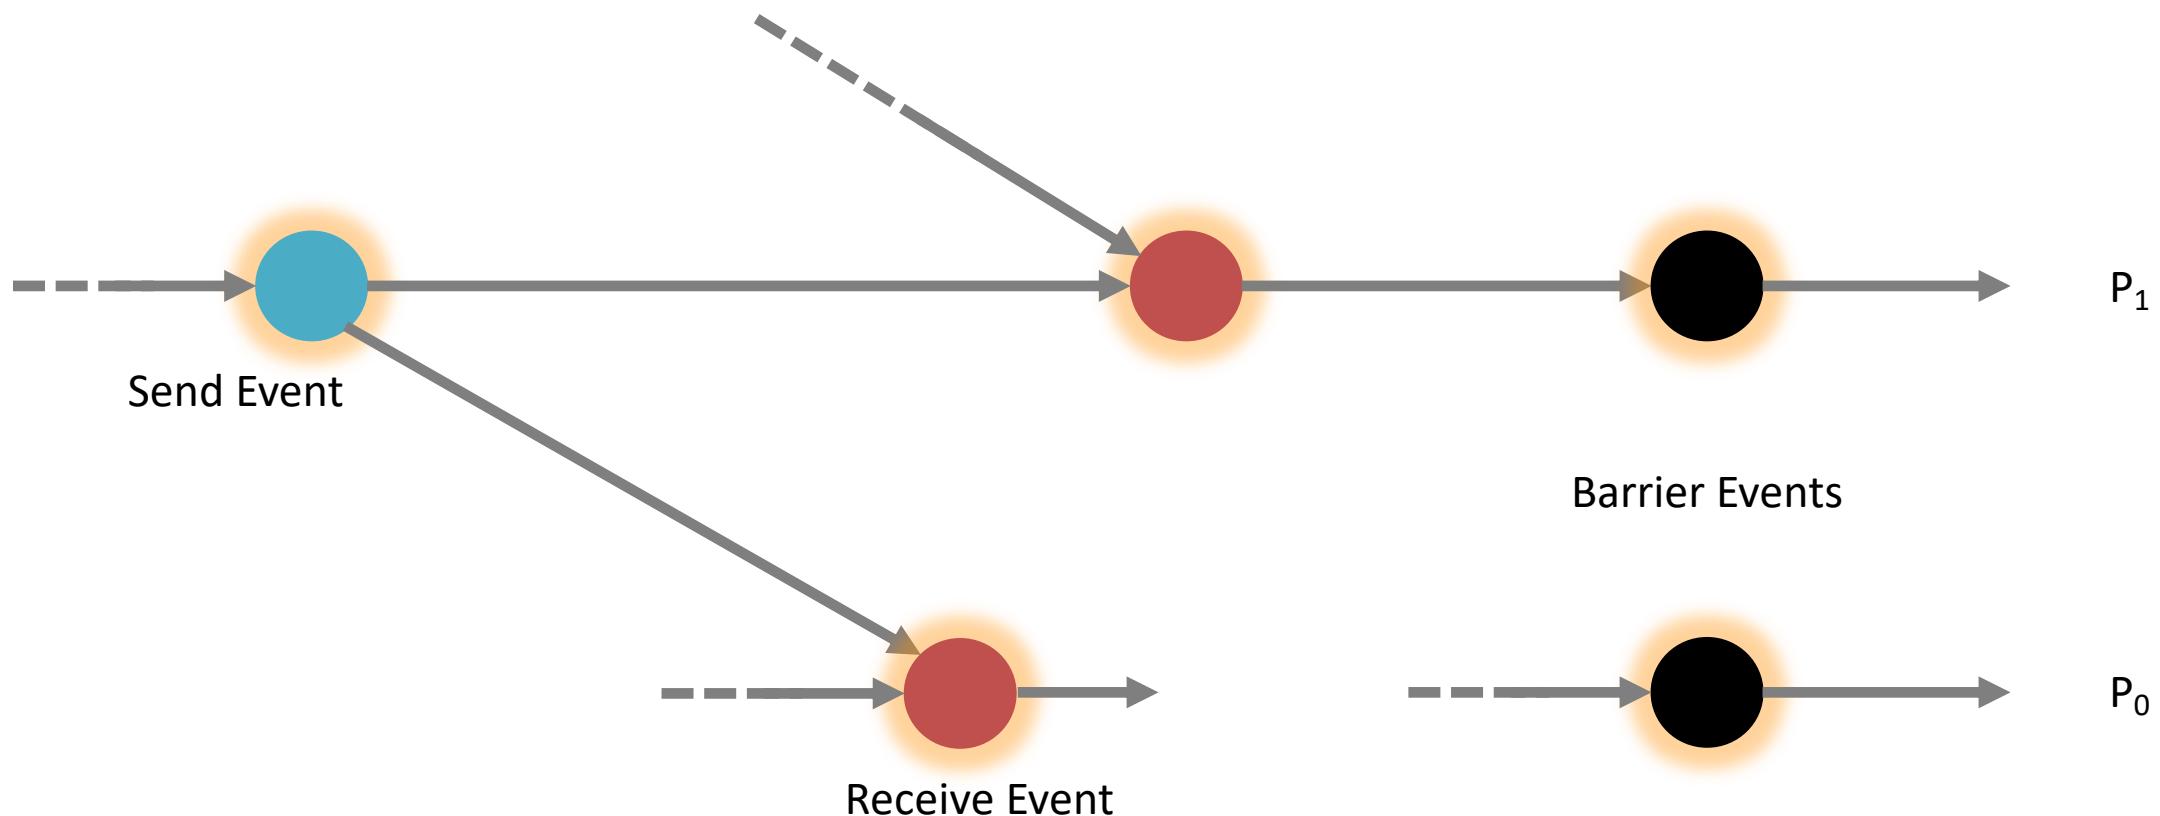
Event Graph Structure

- Directed, acyclic graph (DAG) representing a message-passing program execution [1]



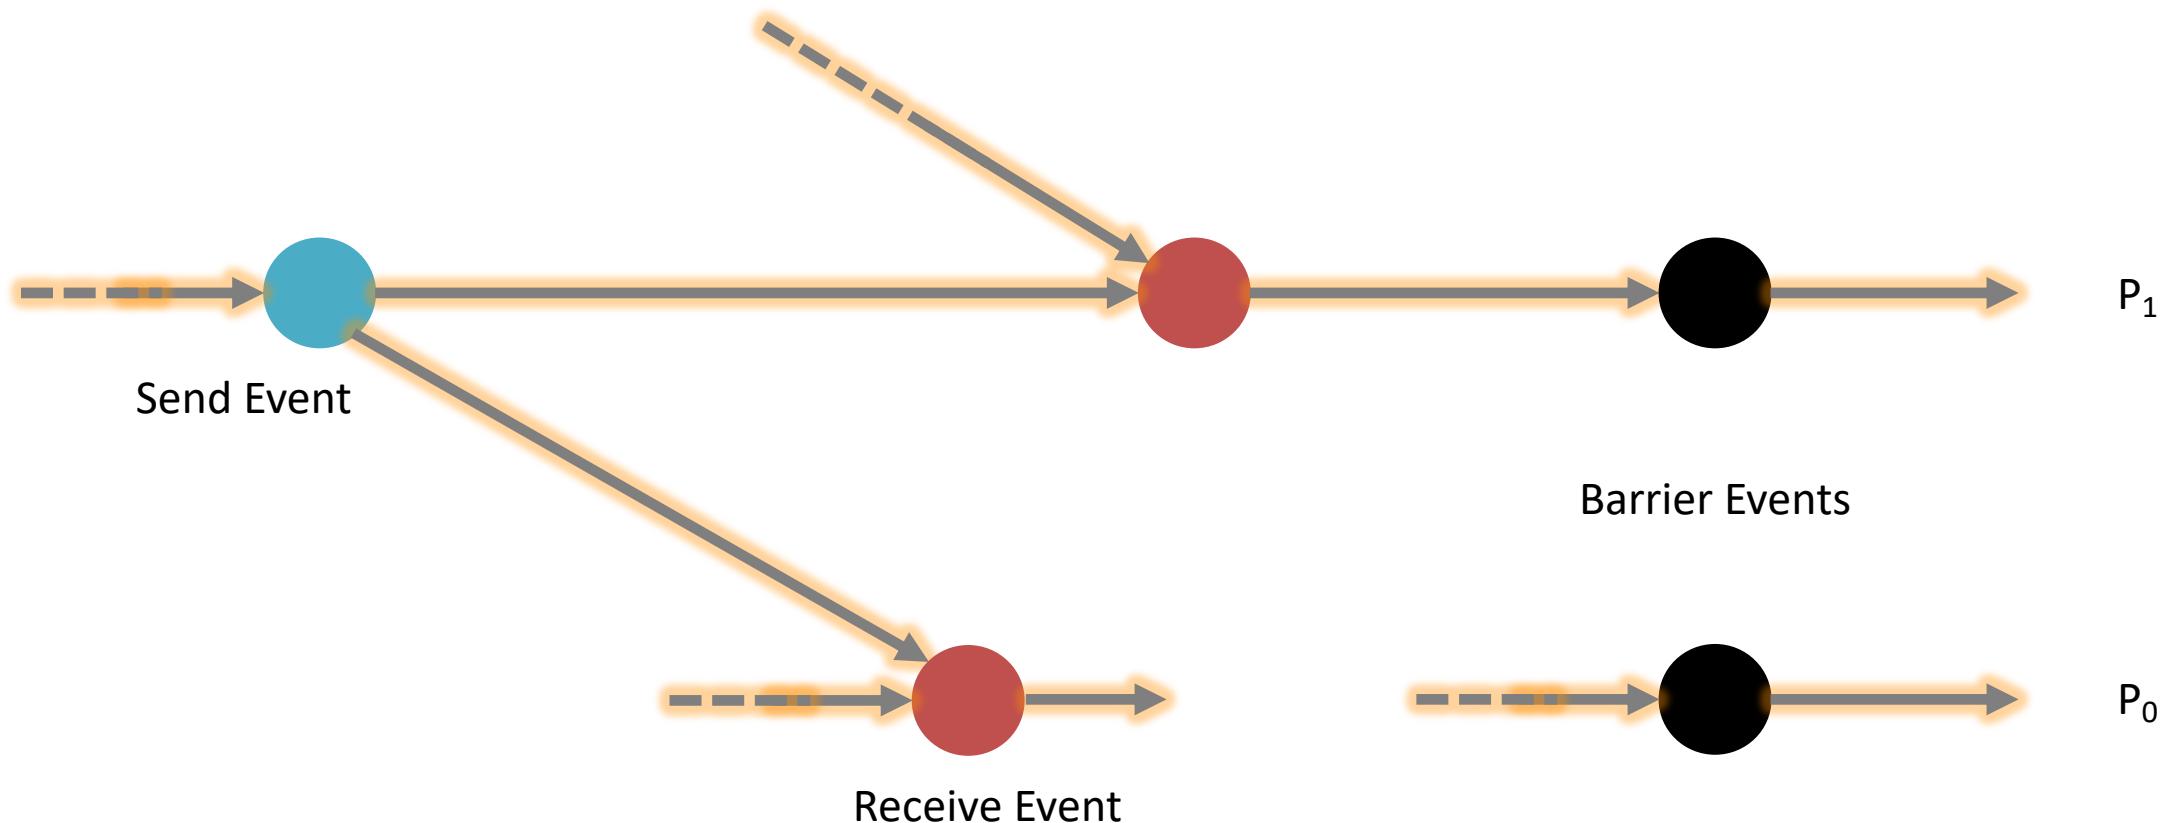
Event Graph Structure

- Directed, acyclic graph (DAG) representing a message-passing program execution [1]
- Vertices represent communication events (e.g., message sends, receives, and barriers)



Event Graph Structure

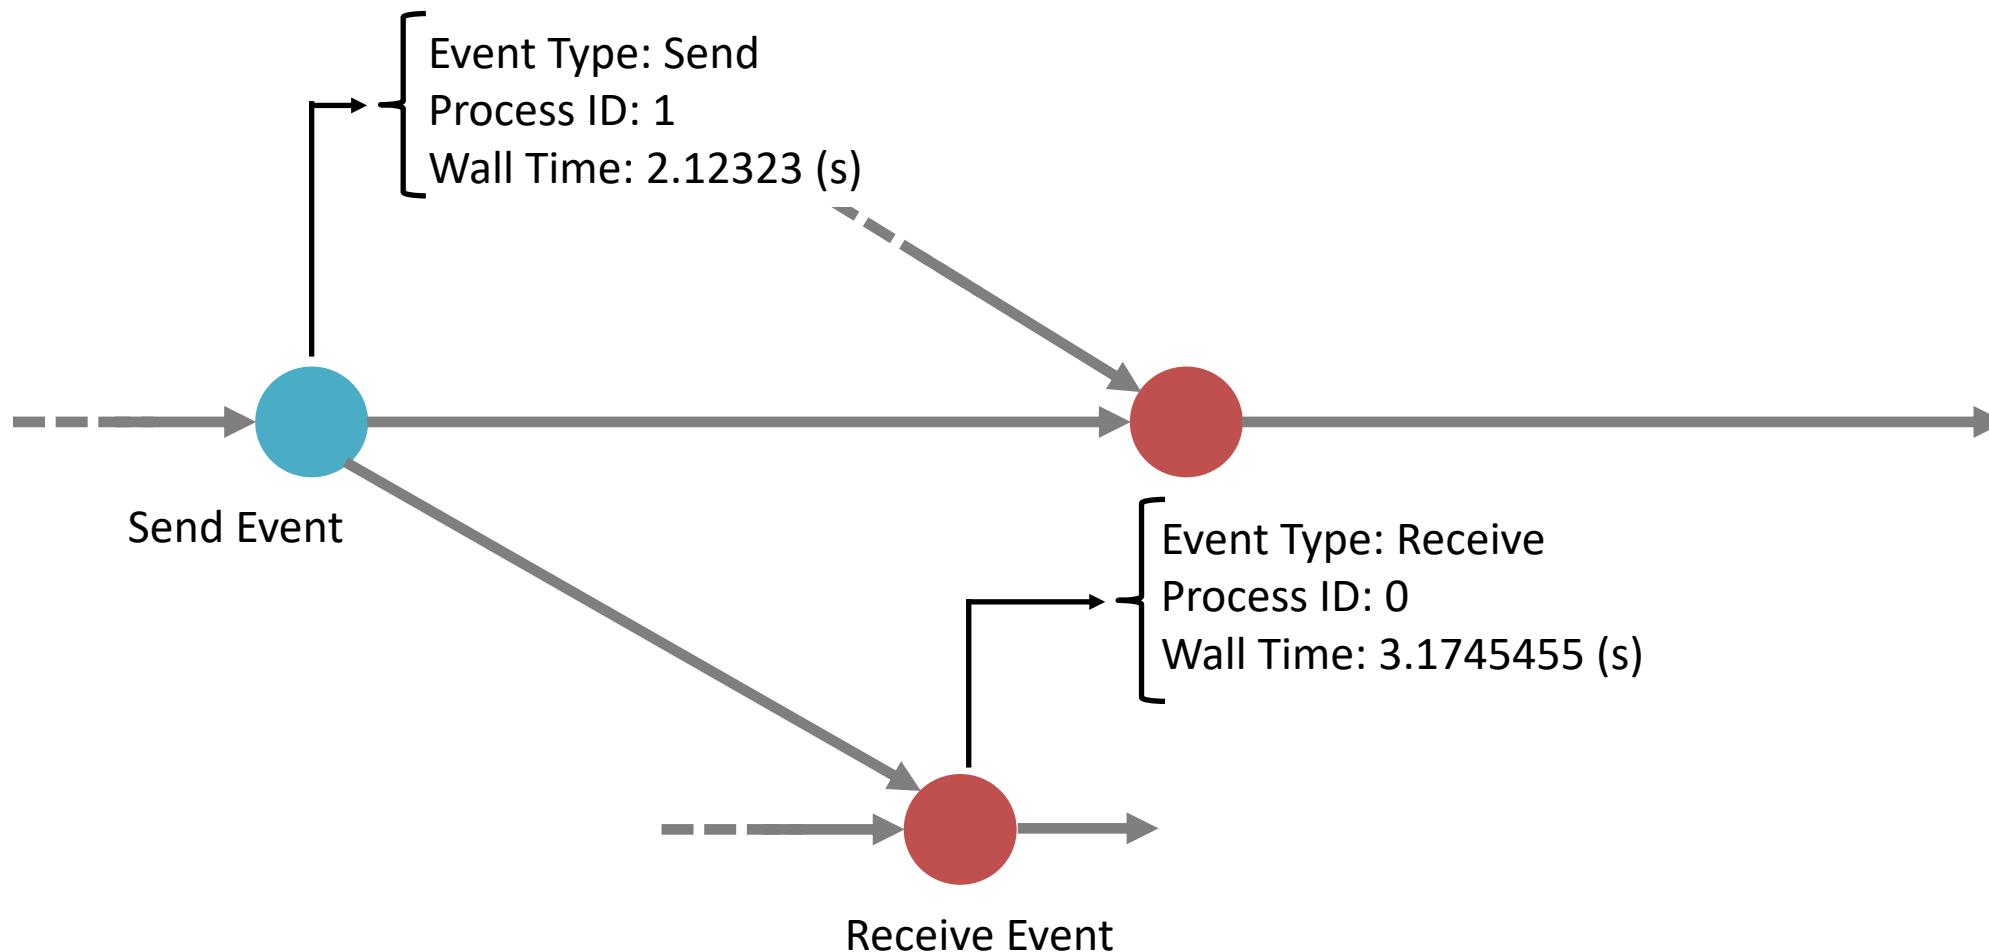
- Directed, acyclic graph (DAG) representing a message-passing program execution [1]
- Vertices represent communication events (e.g., message sends, receives, and barriers)
- Edges represent “happens-before”[2] relationship between events



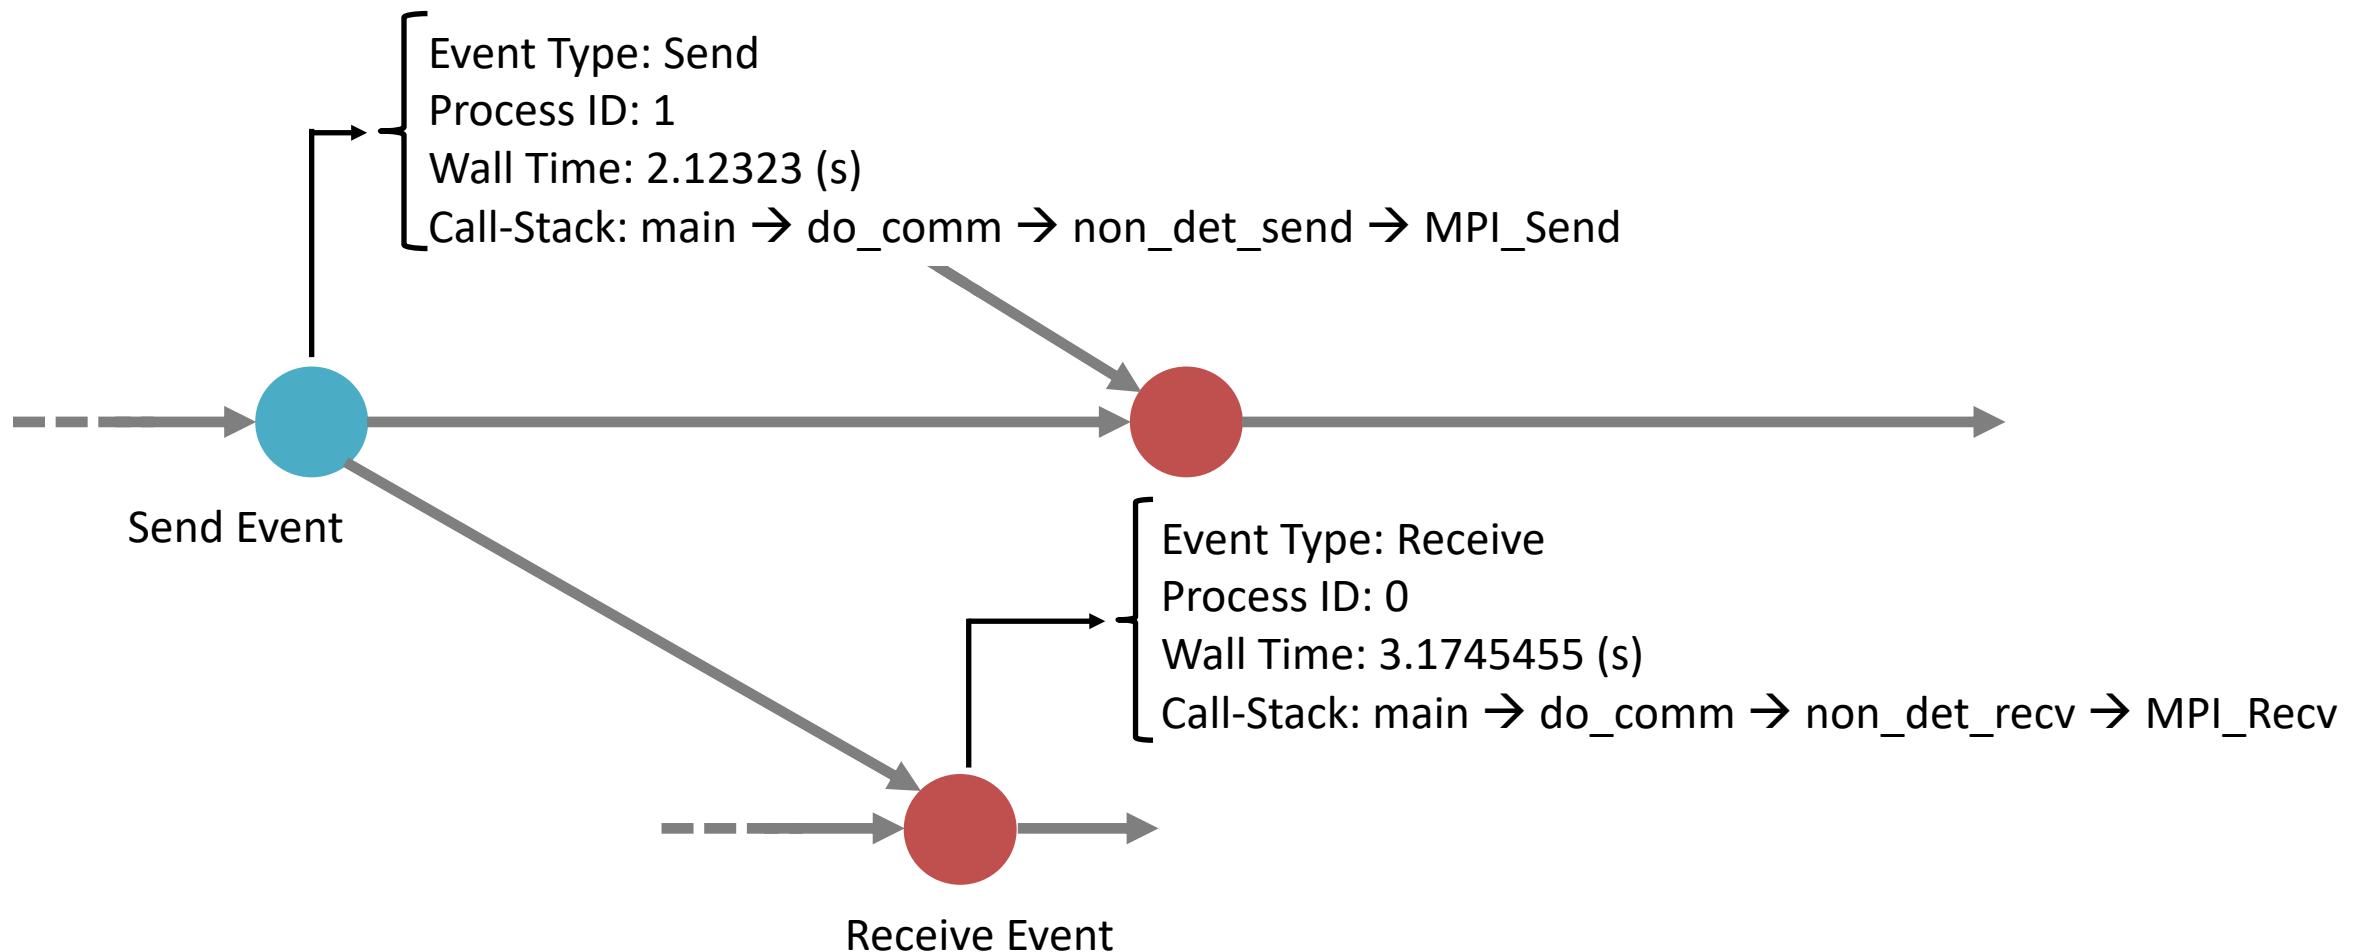
1. Kranzmüller, D., 2000. *Event graph analysis for debugging massively parallel programs*.

2. Lamport, L., 1978. Time, clocks, and the ordering of events in a distributed system. *Communications of the ACM*, 21(7), pp.558-565.

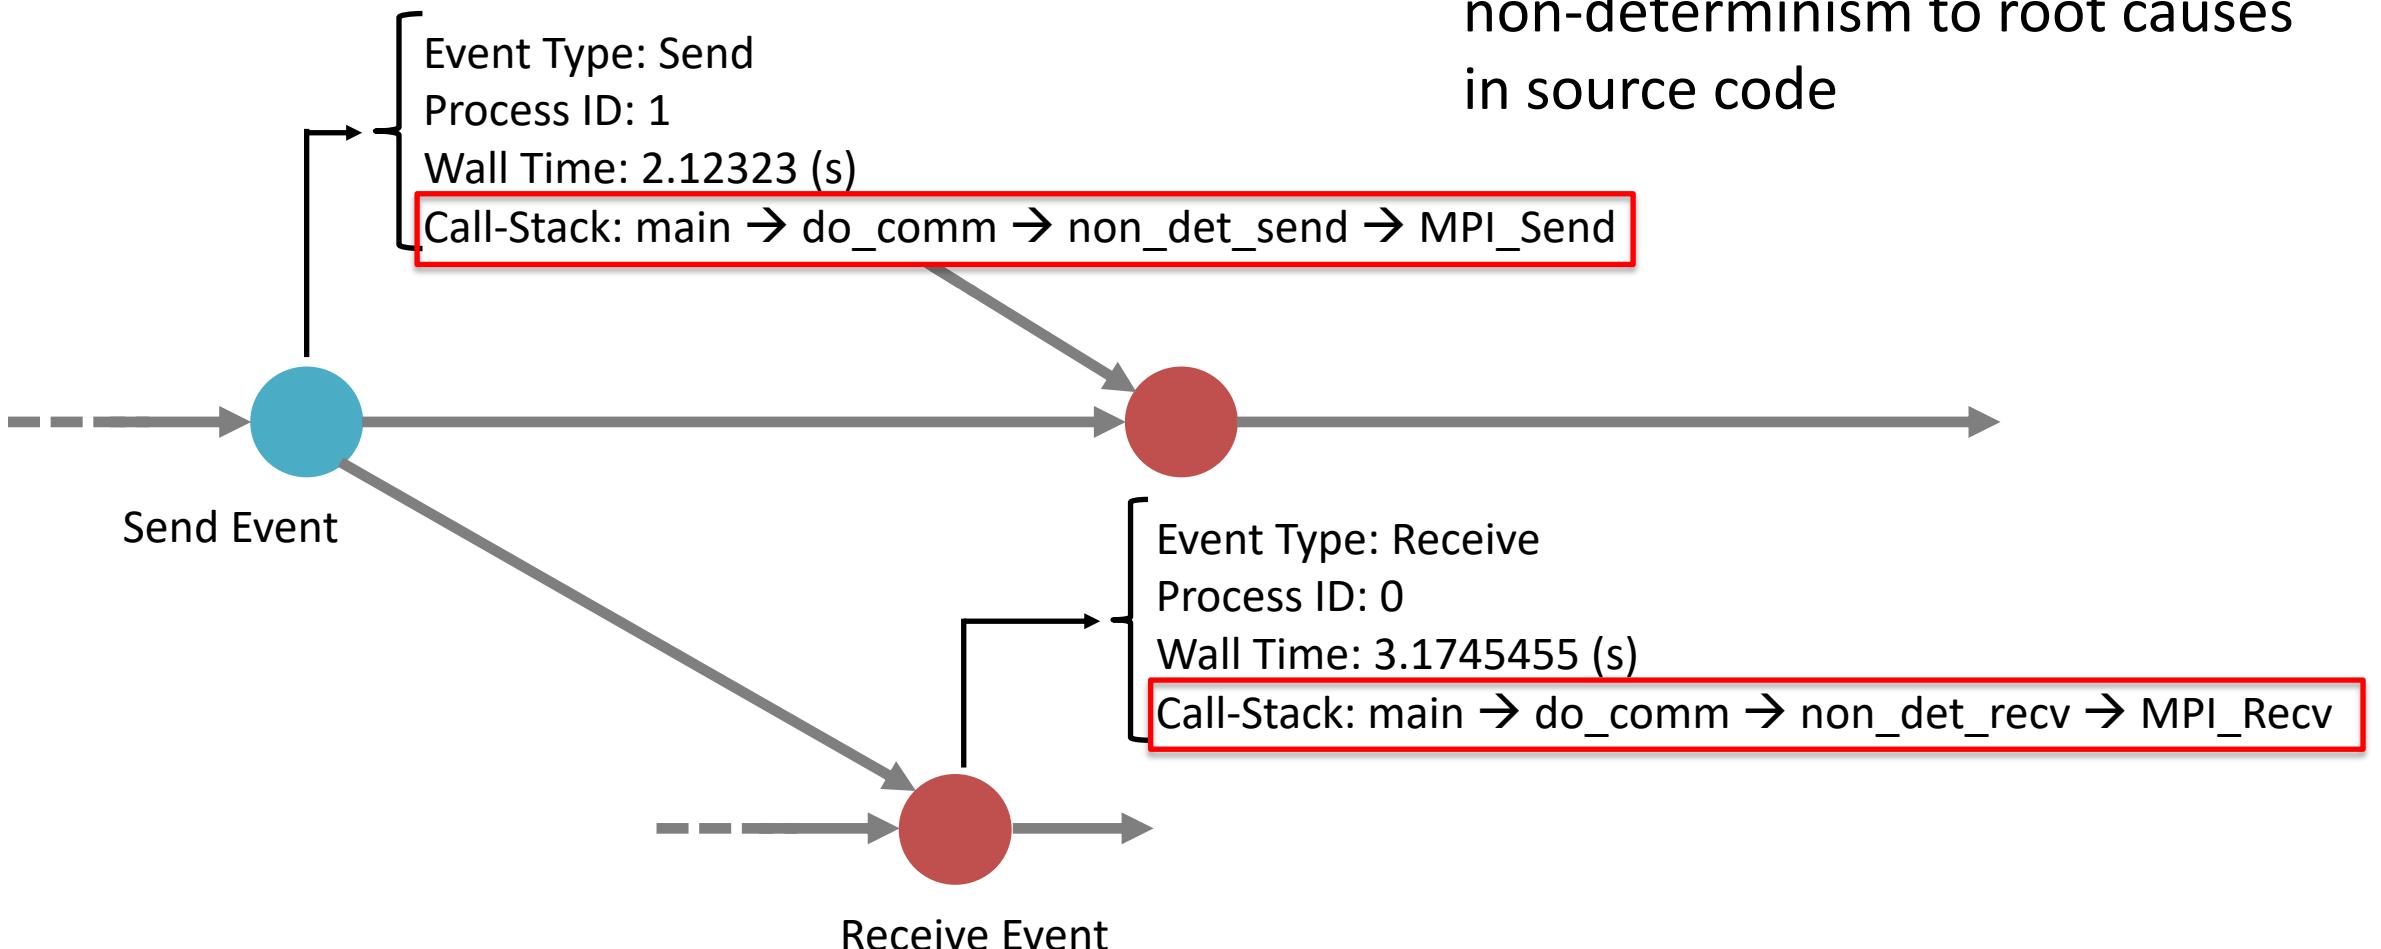
Event Graph Vertex Labels



Event Graph Vertex Labels



Event Graph Vertex Labels



Call-stack labels link run-time non-determinism to root causes in source code

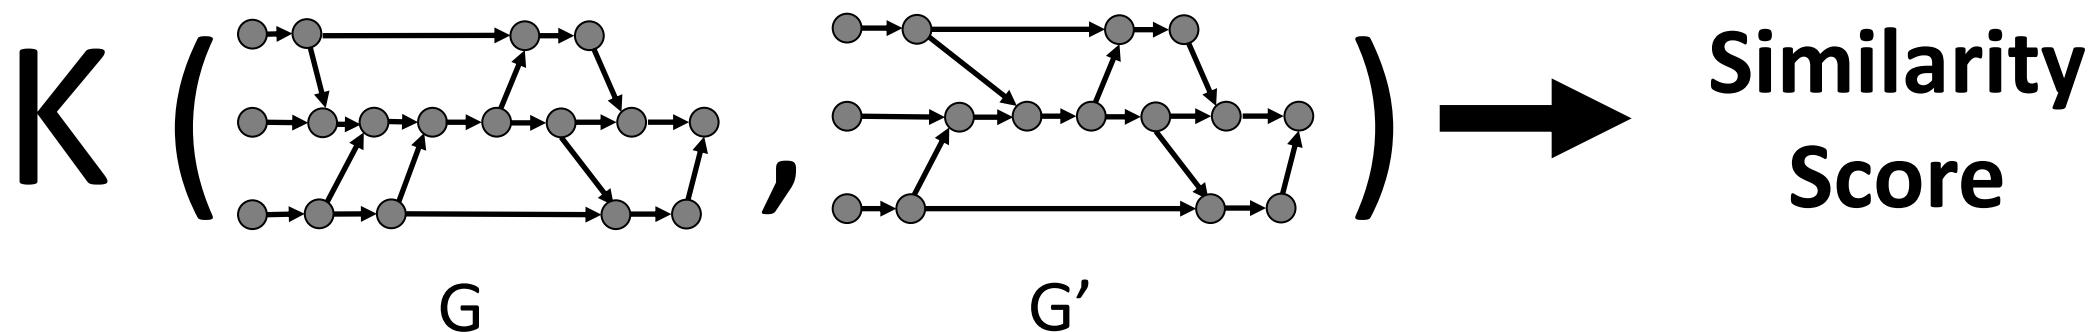
Summary of Event Graph Model

- Traces of MPI application → DAG
- Vertices → communication events
- Edges → happens-before orders
- Vertex Labels:
 - Event type → What happened?
 - Process ID → In which process did it happen?
 - Timestamp → When did it happen?
 - Call-Stack → Where in the code did it come from?

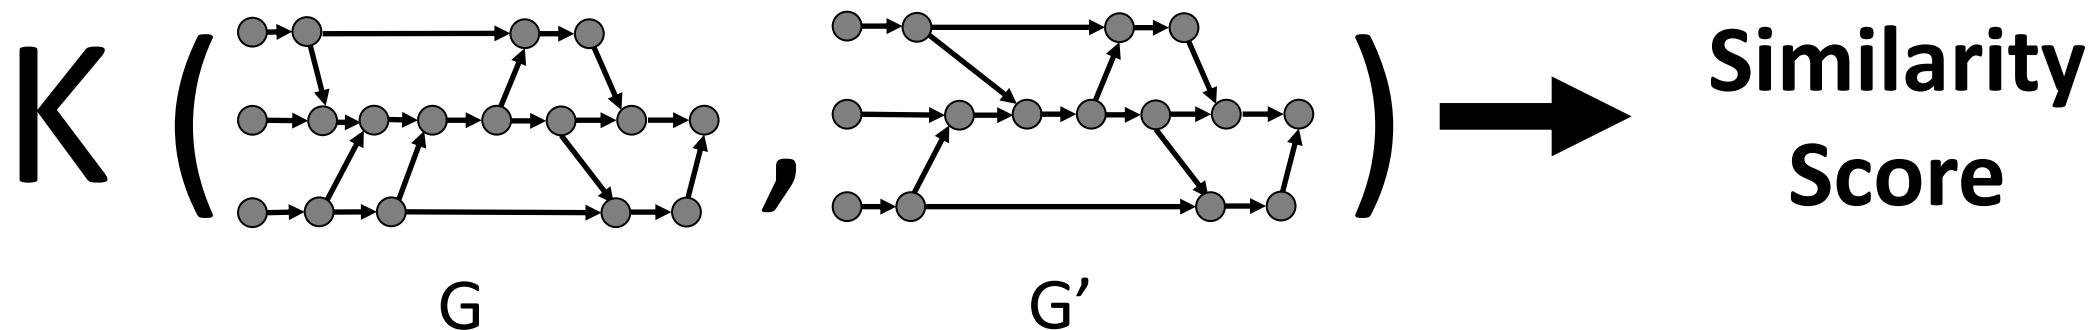
Workflow for Non-Determinism Characterization

- **Phase 1:** Build graph-structured models of executions
- **Phase 2:** Quantify cross-execution trends in non-deterministic communication via graph similarity
- **Phase 3:** Detect periods of anomalous execution dissimilarity and localize potential root causes

Graph Similarity via Graph Kernels

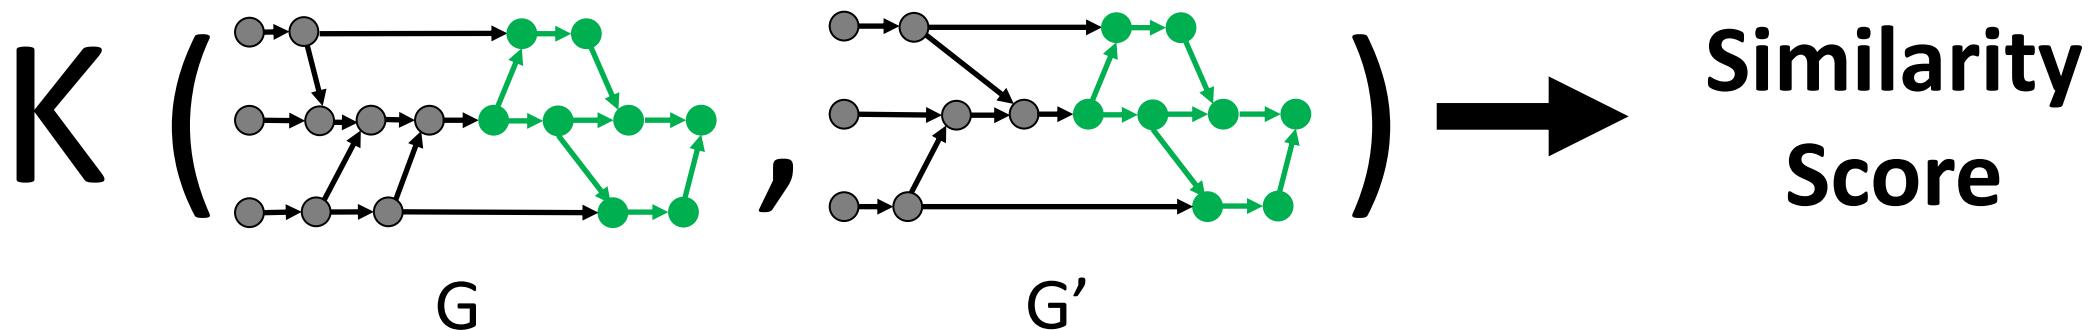


Graph Similarity via Graph Kernels



Intuition: K counts matching substructures

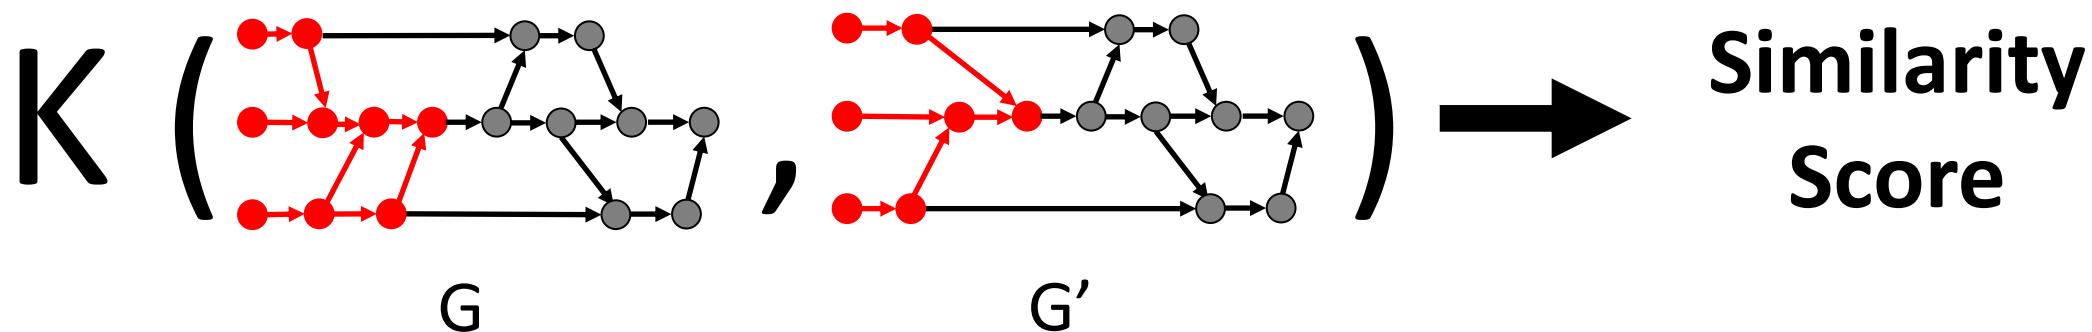
Graph Similarity via Graph Kernels



Intuition: K counts matching substructures

- **Matches** between G and G' increase score

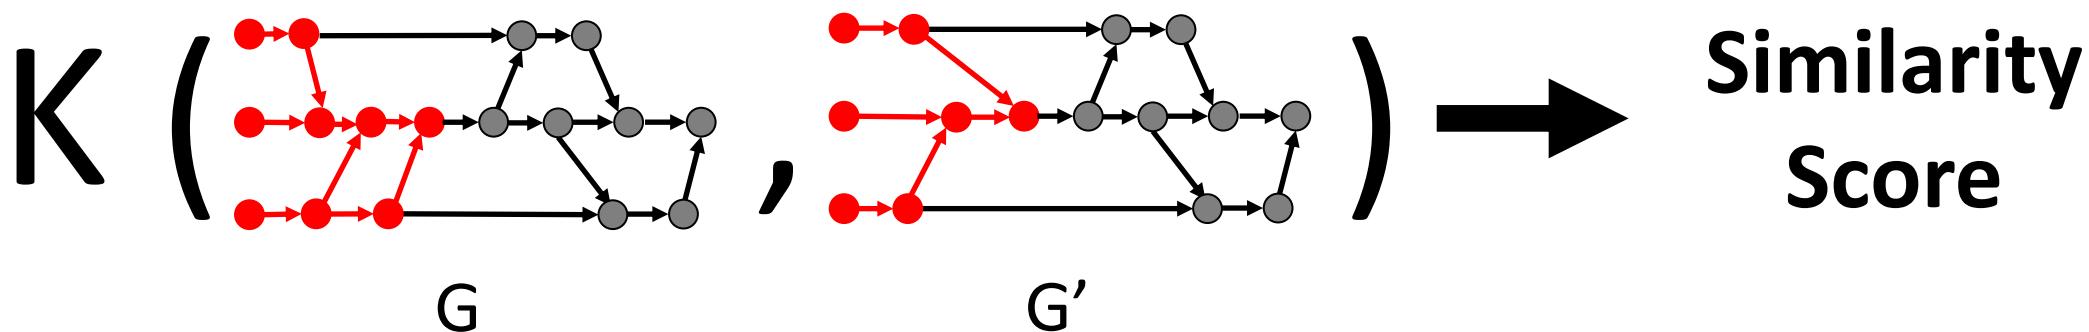
Graph Similarity via Graph Kernels



Intuition: K counts matching substructures

- Matches between G and G' increase score
- **Differences** do not

Graph Similarity via Graph Kernels



Intuition: K counts matching substructures

- Matches between G and G' increase score
- **Differences** do not
- Graph kernel K **induces a metric** → the graph kernel distance [1]

$$\text{Formula: } D(G, G') = \sqrt{K(G, G') + K(G, G') - 2K(G, G')}$$

Quantifying Non-Determinism via Graph Kernel Distance

We evaluate the Weisfeiler-Lehman Subtree Pattern Kernel for quantifying dissimilarity between event graphs:

1. Shervashidze, N., Schweitzer, P., Leeuwen, E.J.V., Mehlhorn, K. and Borgwardt, K.M., 2011. Weisfeiler-lehman graph kernels. *Journal of Machine Learning Research*, 12(Sep), pp.2539-2561.
2. Yanardag, P. and Vishwanathan, S.V.N., 2015, August. Deep graph kernels. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (pp. 1365-1374). ACM.
3. Ghosh, S., Das, N., Gonçalves, T., Quaresma, P. and Kundu, M., 2018. The journey of graph kernels through two decades. *Computer Science Review*, 27, pp.88-111.



Quantifying Non-Determinism via Graph Kernel Distance

We evaluate the Weisfeiler-Lehman Subtree Pattern Kernel for quantifying dissimilarity between event graphs:

- Demonstrated performance on other graph classification tasks [1]



Quantifying Non-Determinism via Graph Kernel Distance

We evaluate the Weisfeiler-Lehman Subtree Pattern Kernel for quantifying dissimilarity between event graphs:

- Demonstrated performance on other graph classification tasks [1]
- Scalability compared to other graph kernels



Quantifying Non-Determinism via Graph Kernel Distance

We evaluate the Weisfeiler-Lehman Subtree Pattern Kernel for quantifying dissimilarity between event graphs:

- Demonstrated performance on other graph classification tasks [1]
- Scalability compared to other graph kernels
- Incorporation of arbitrary vertex label data

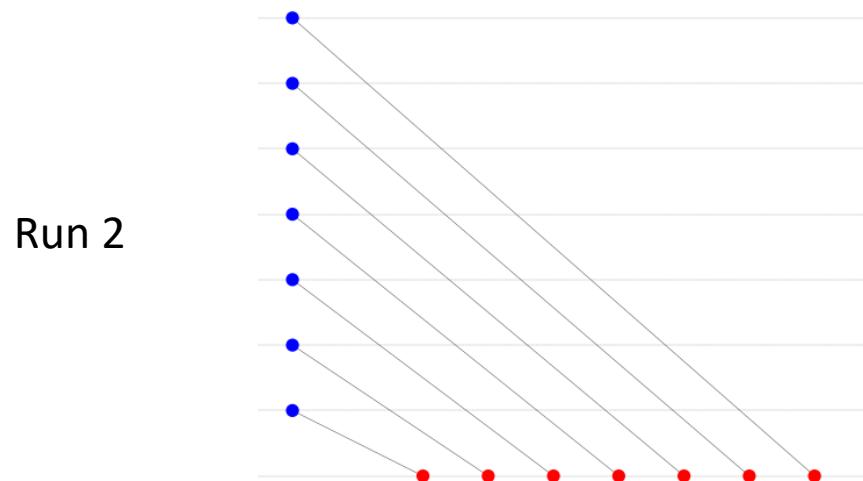
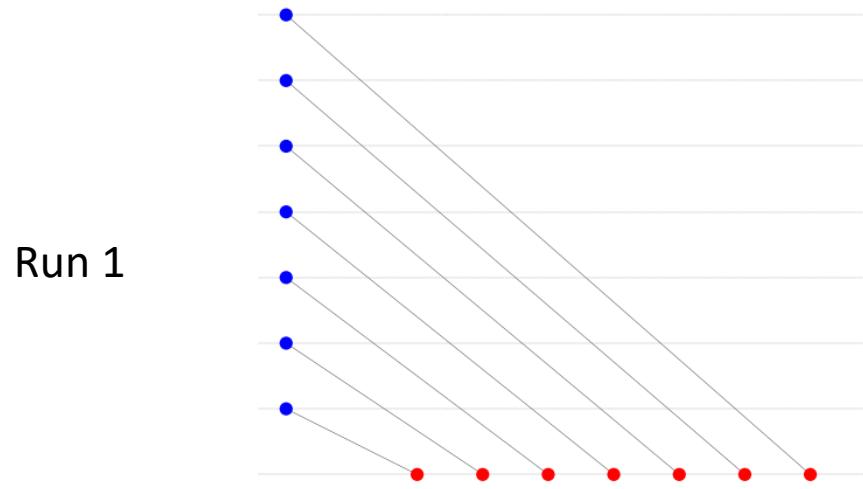


Kernel Distance Evaluation Methodology

- Construct event graphs for common communication patterns with:
 - Controlled degree of non-determinism
 - Fixed amount of communication volume

Kernel Distance Evaluation Methodology

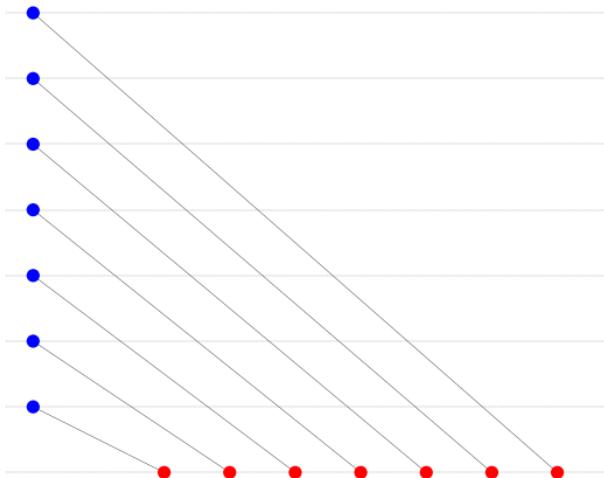
- Construct event graphs for common communication patterns with:
 - Controlled degree of non-determinism
 - Fixed amount of communication volume
- **Hypothesis:** As greater non-determinism is permitted in the runs, the graph kernel distances between the event graphs representing those runs will increase



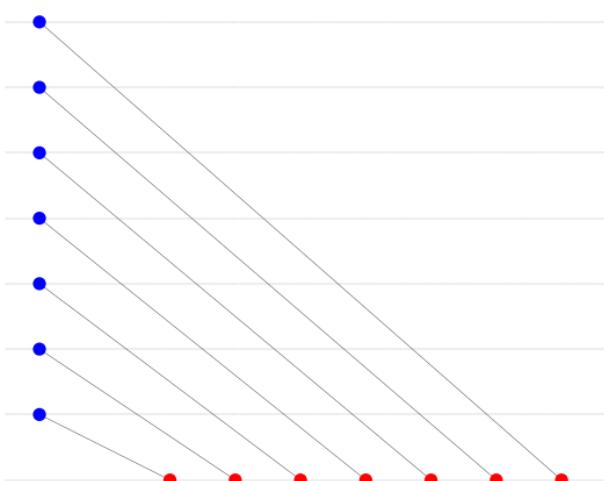
Deterministic

All receives guaranteed to occur
in same order run-to-run

Run 1

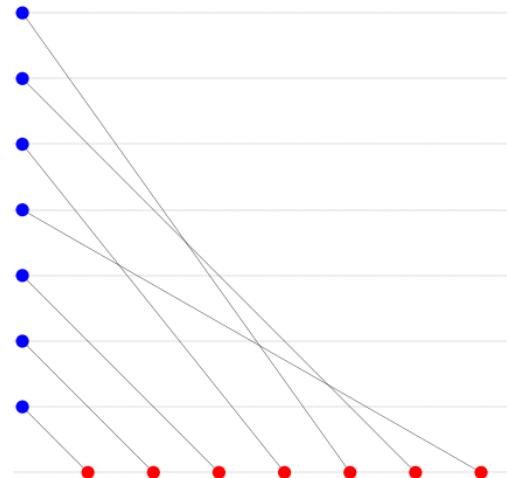


Run 2



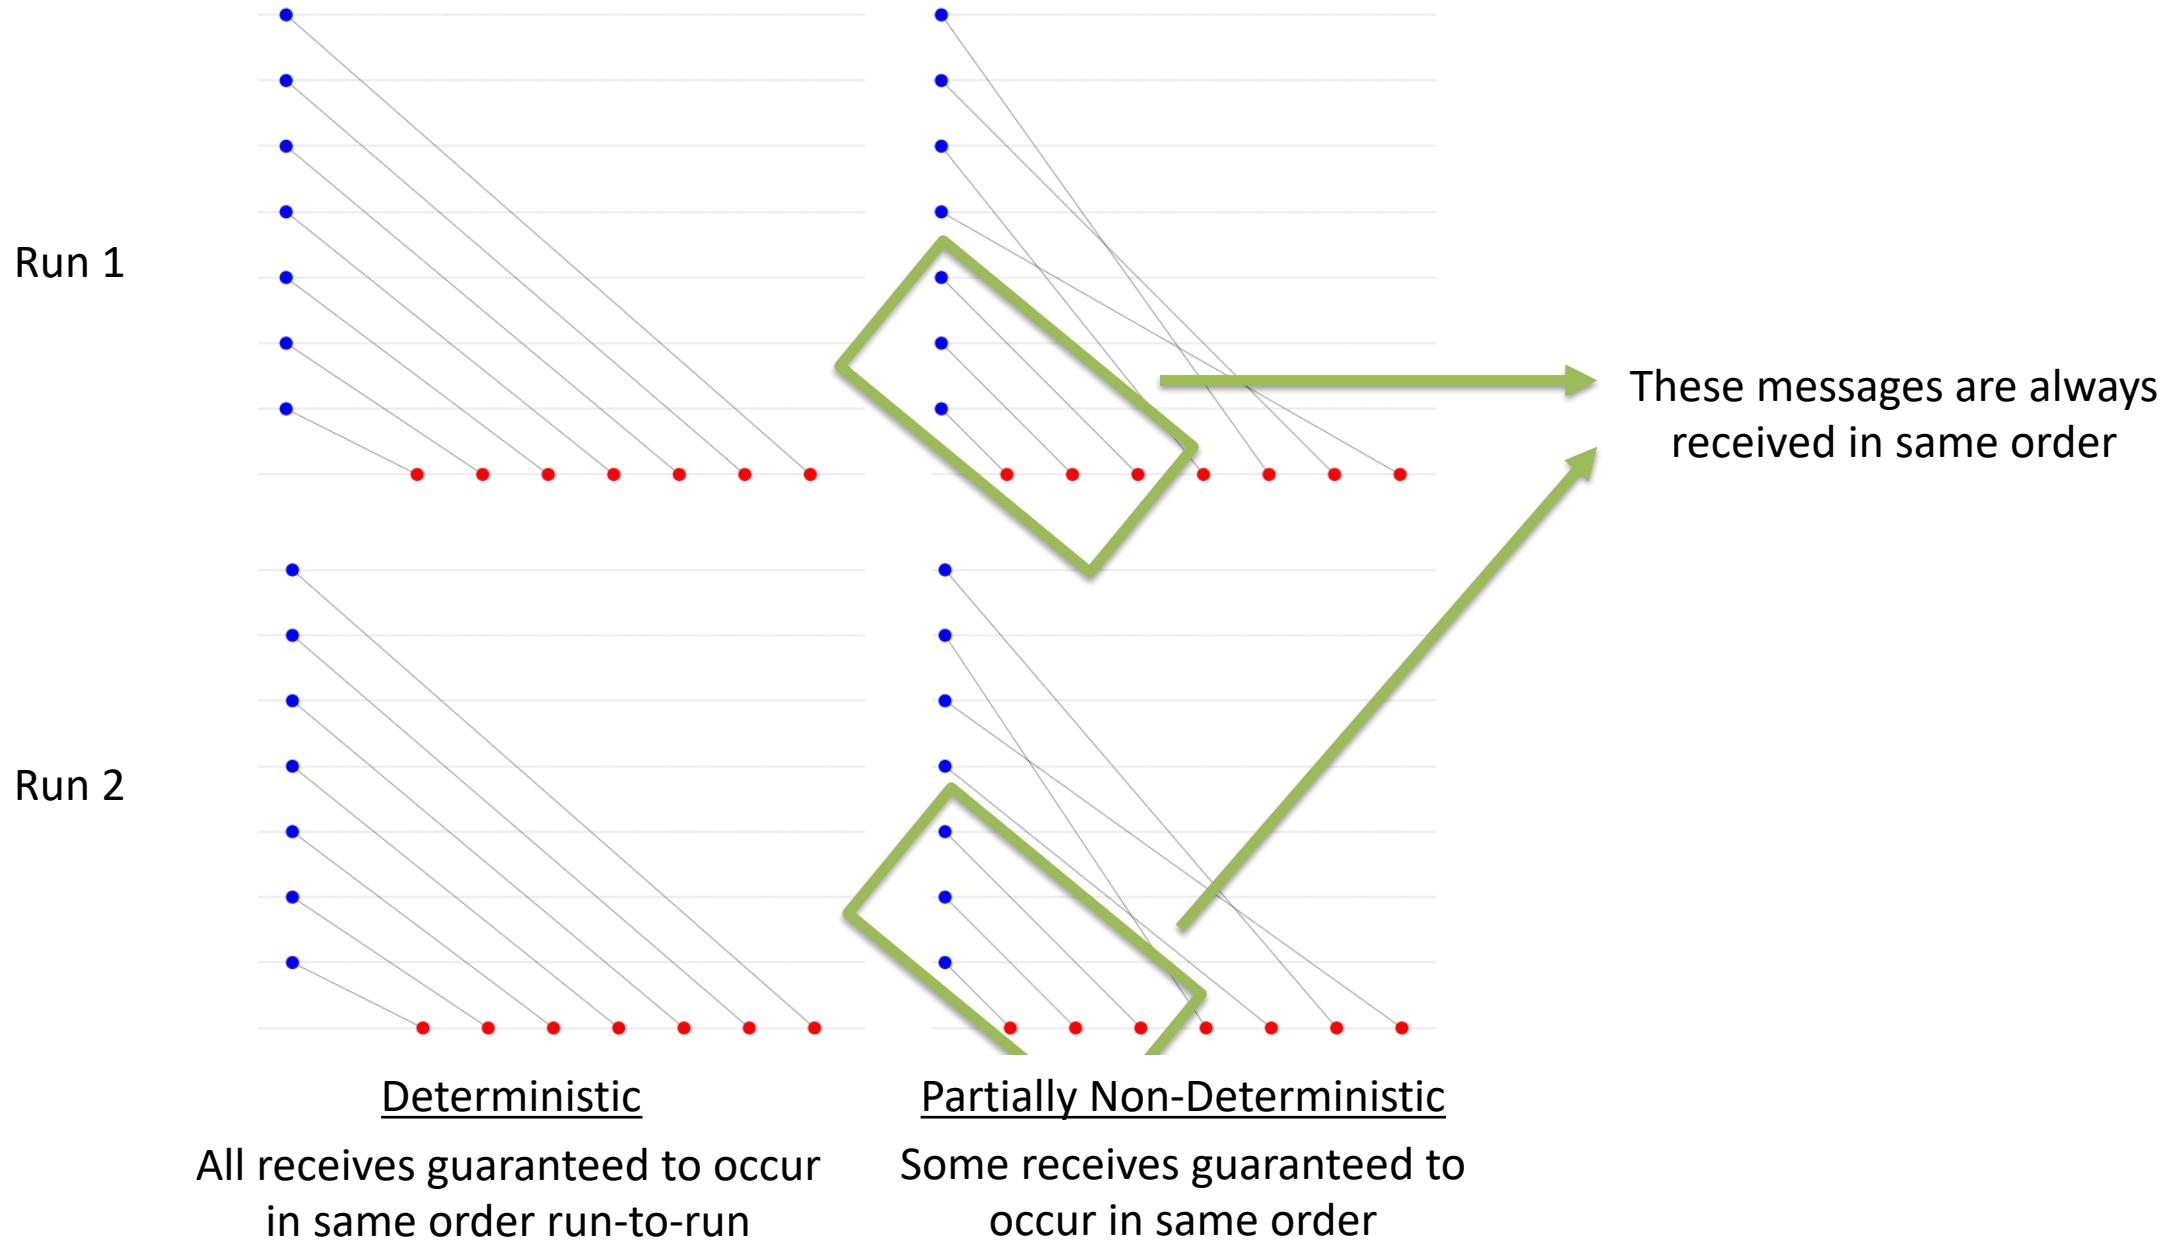
Deterministic

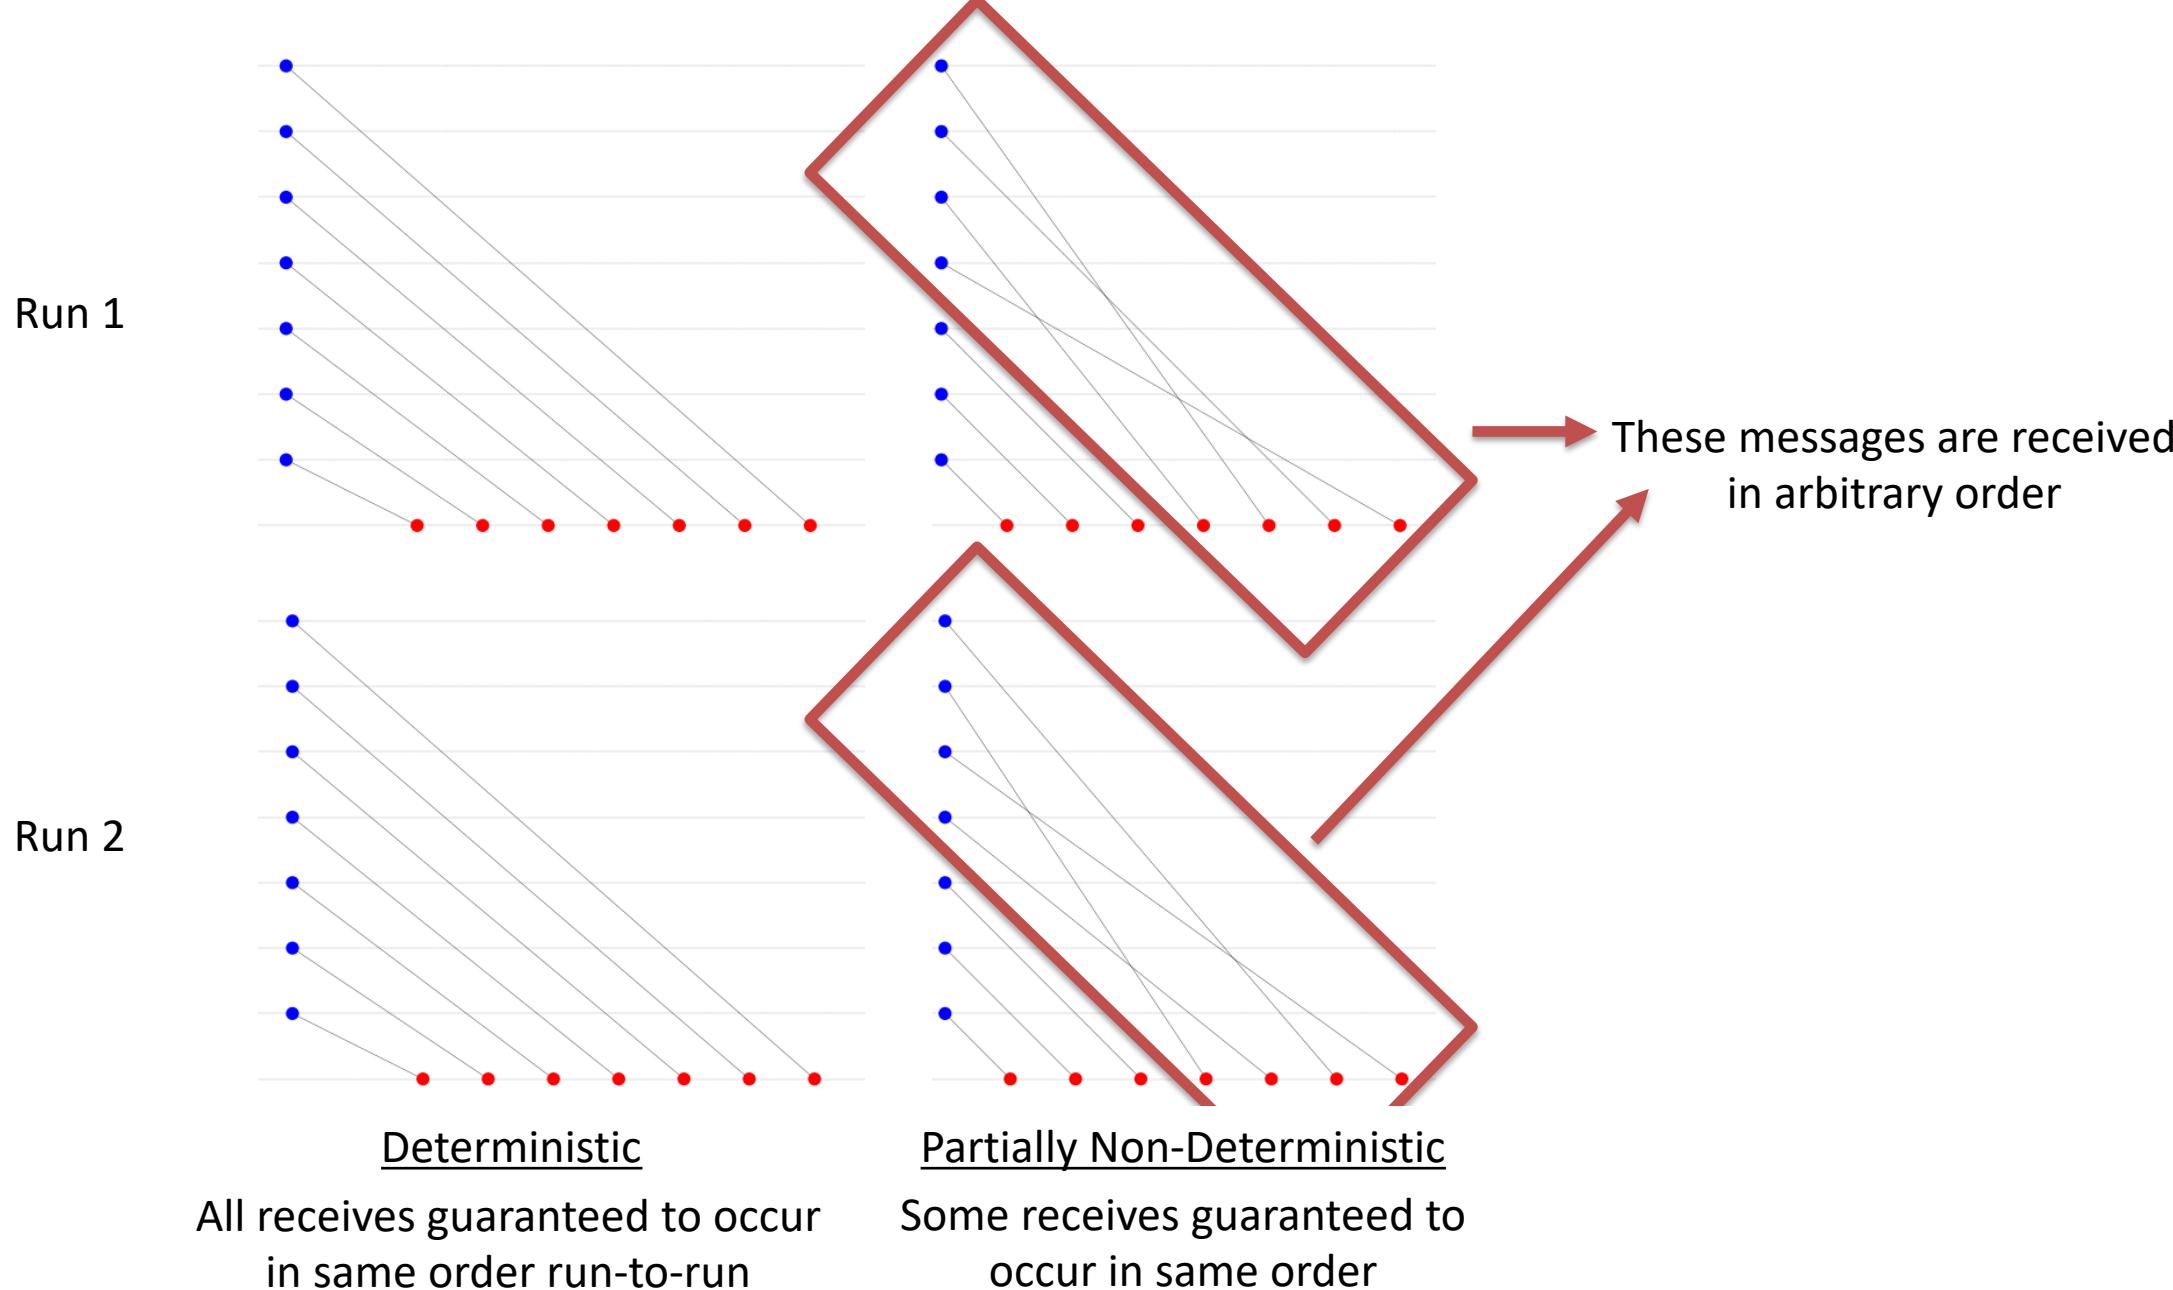
All receives guaranteed to occur
in same order run-to-run



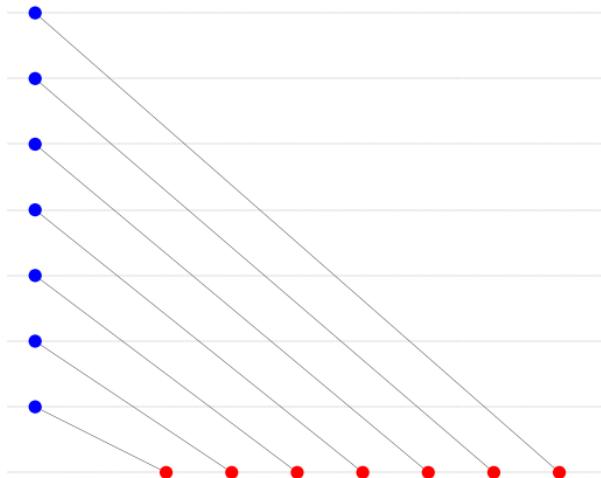
Partially Non-Deterministic

Some receives guaranteed to
occur in same order

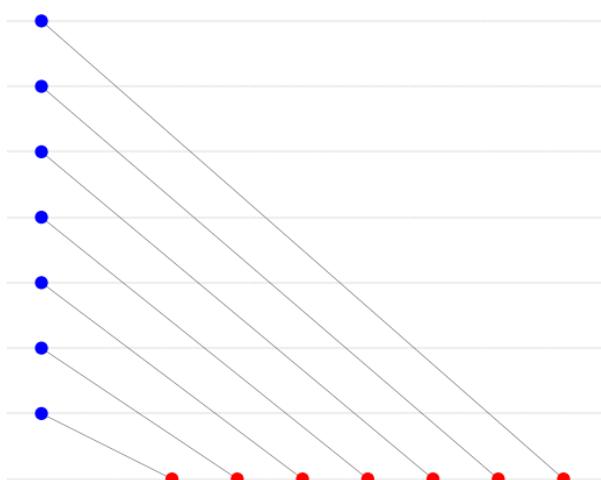




Run 1

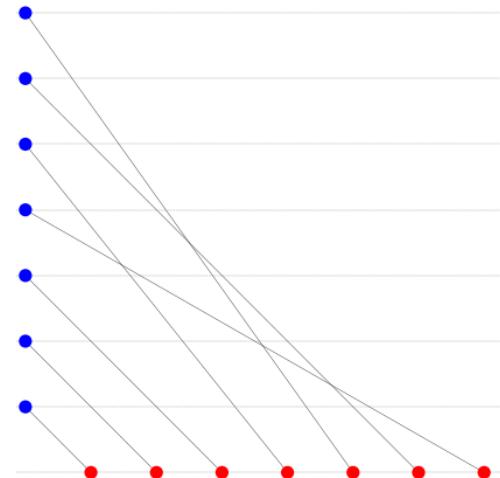


Run 2



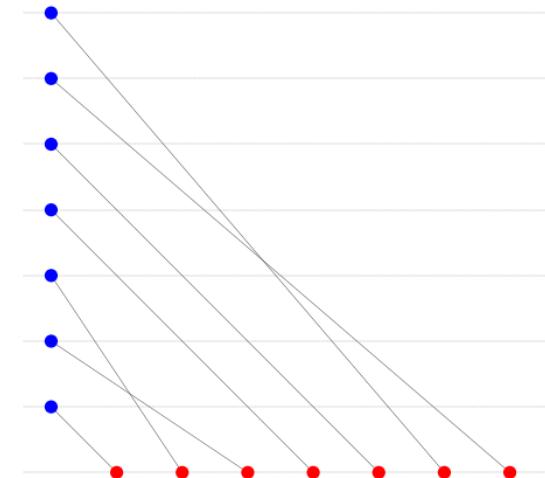
Deterministic

All receives guaranteed to occur
in same order run-to-run



Partially Non-Deterministic

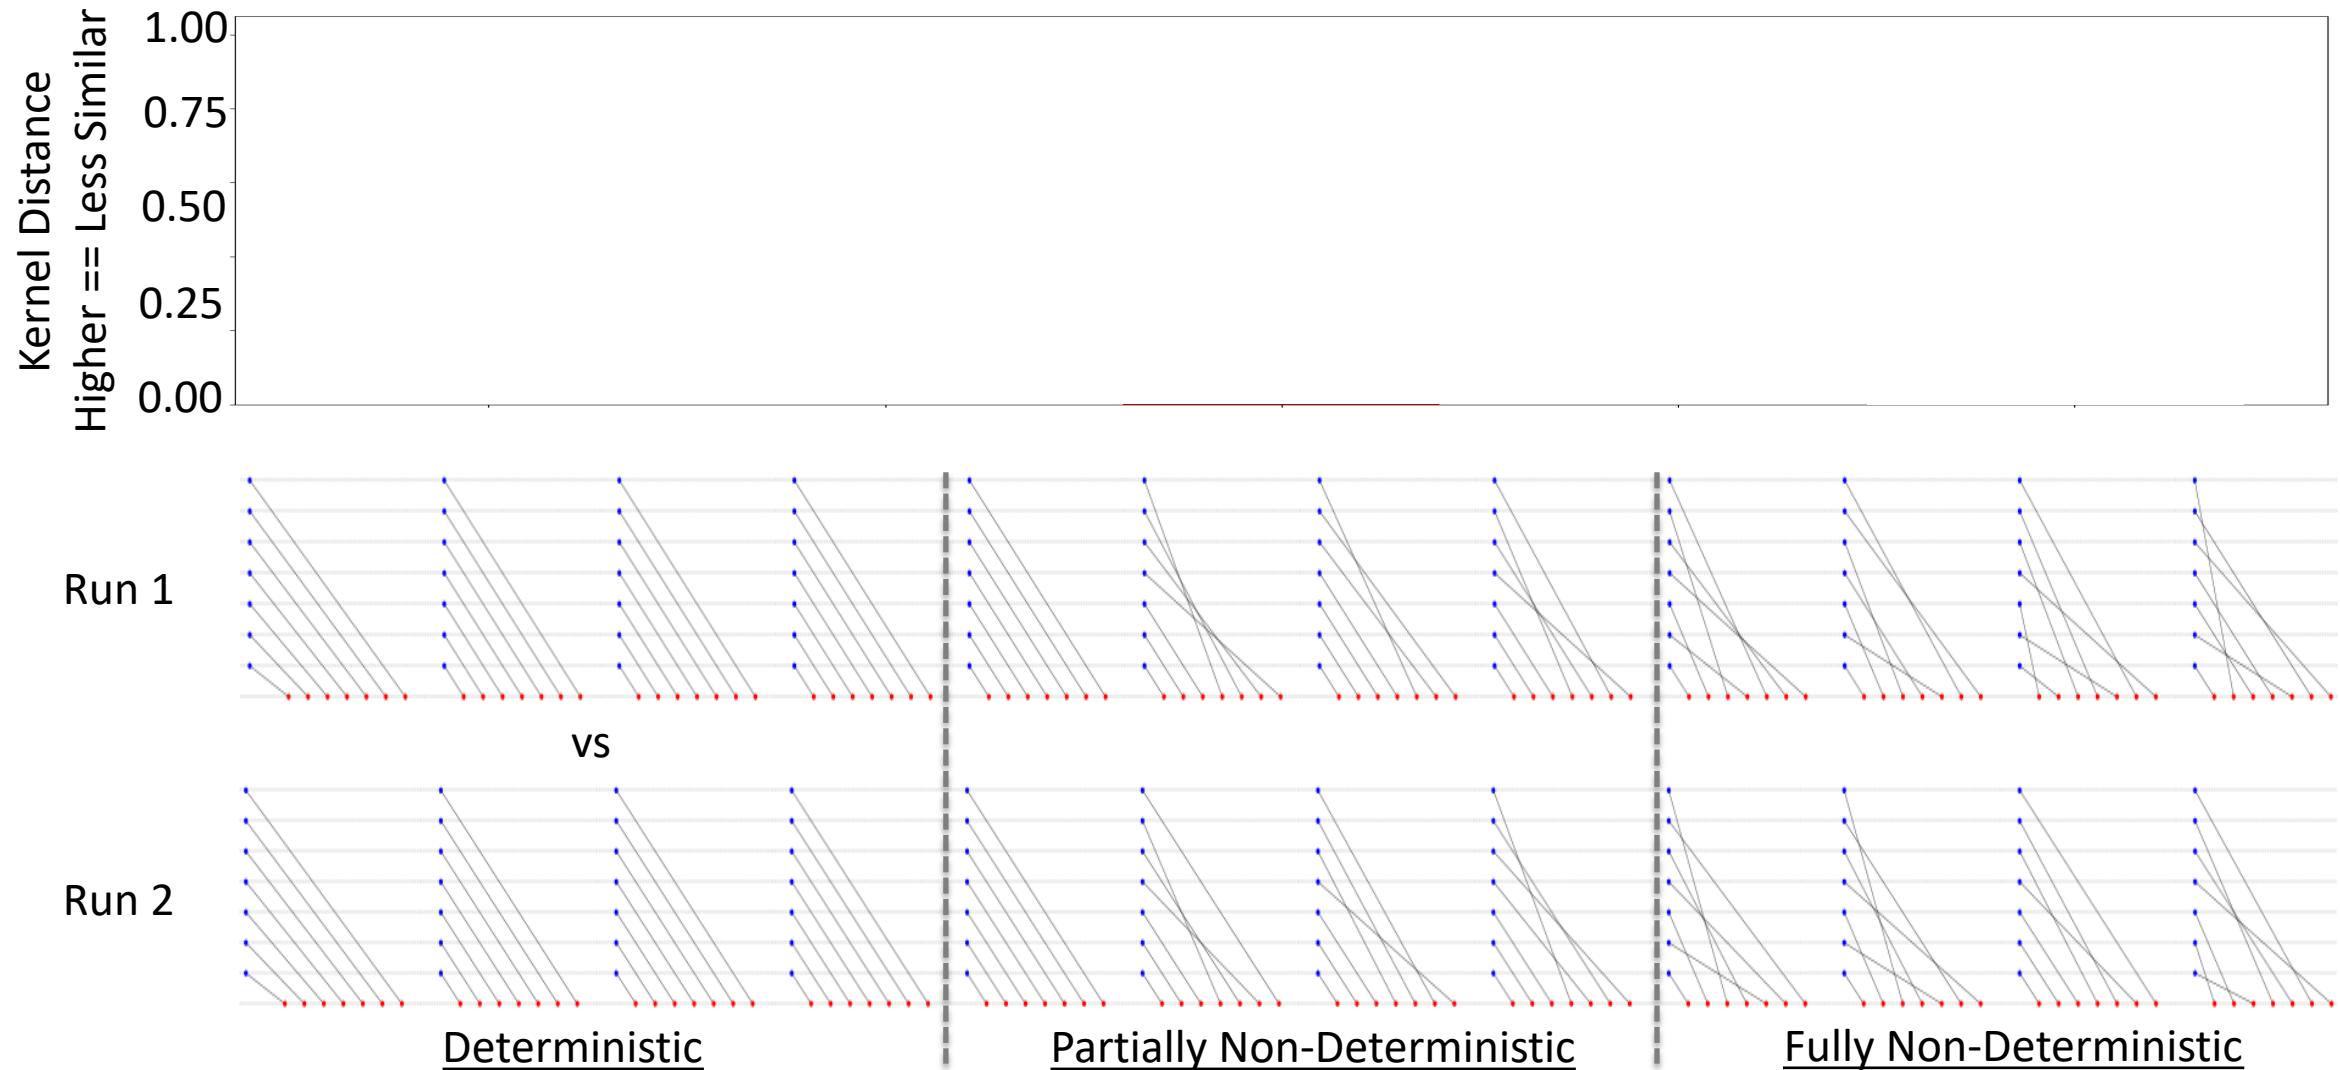
Some receives guaranteed to
occur in same order



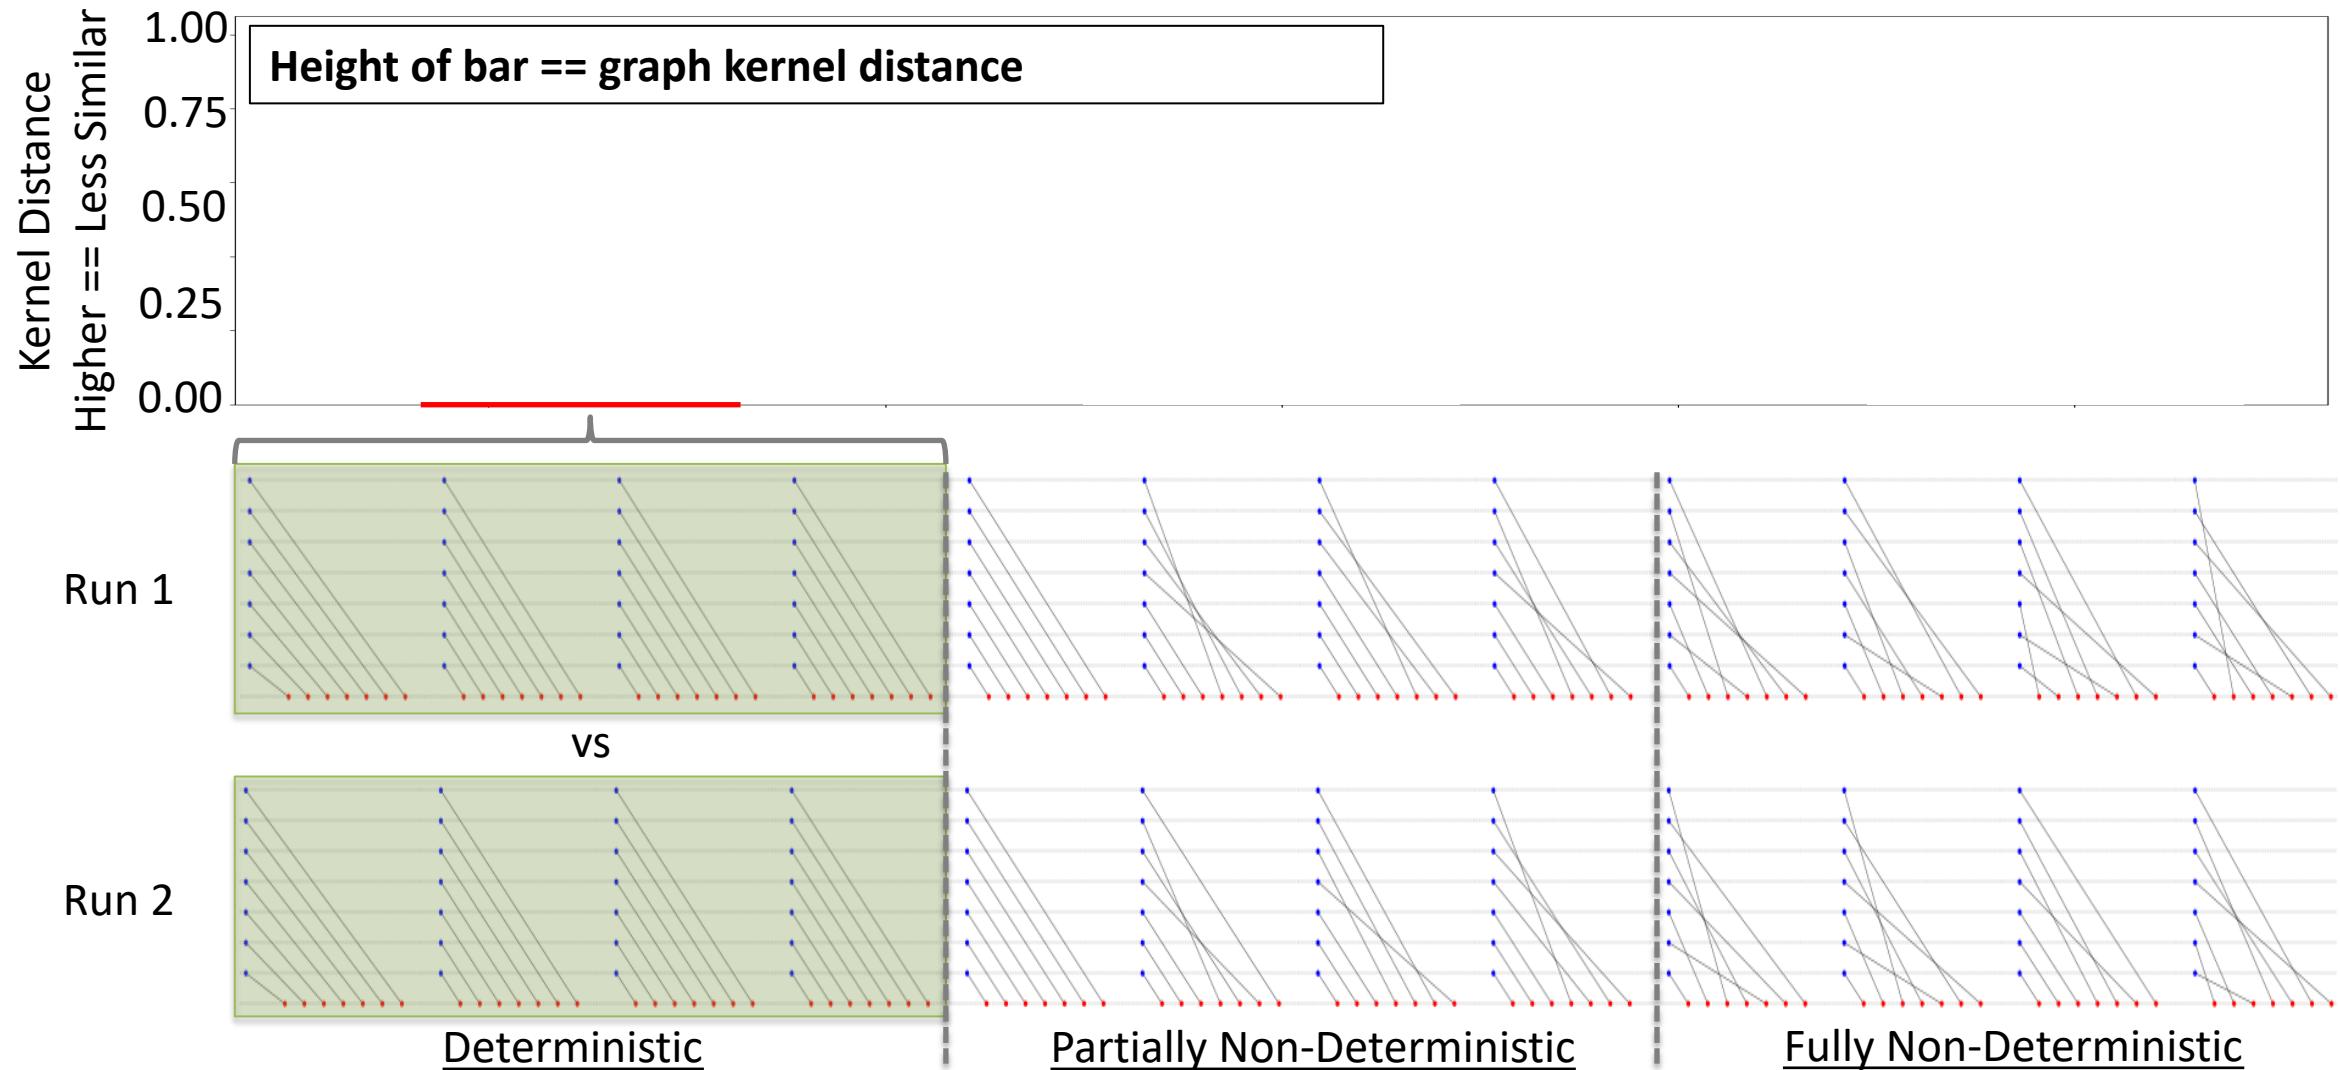
Fully Non-Deterministic

All receives occur in
arbitrary order run-to-run

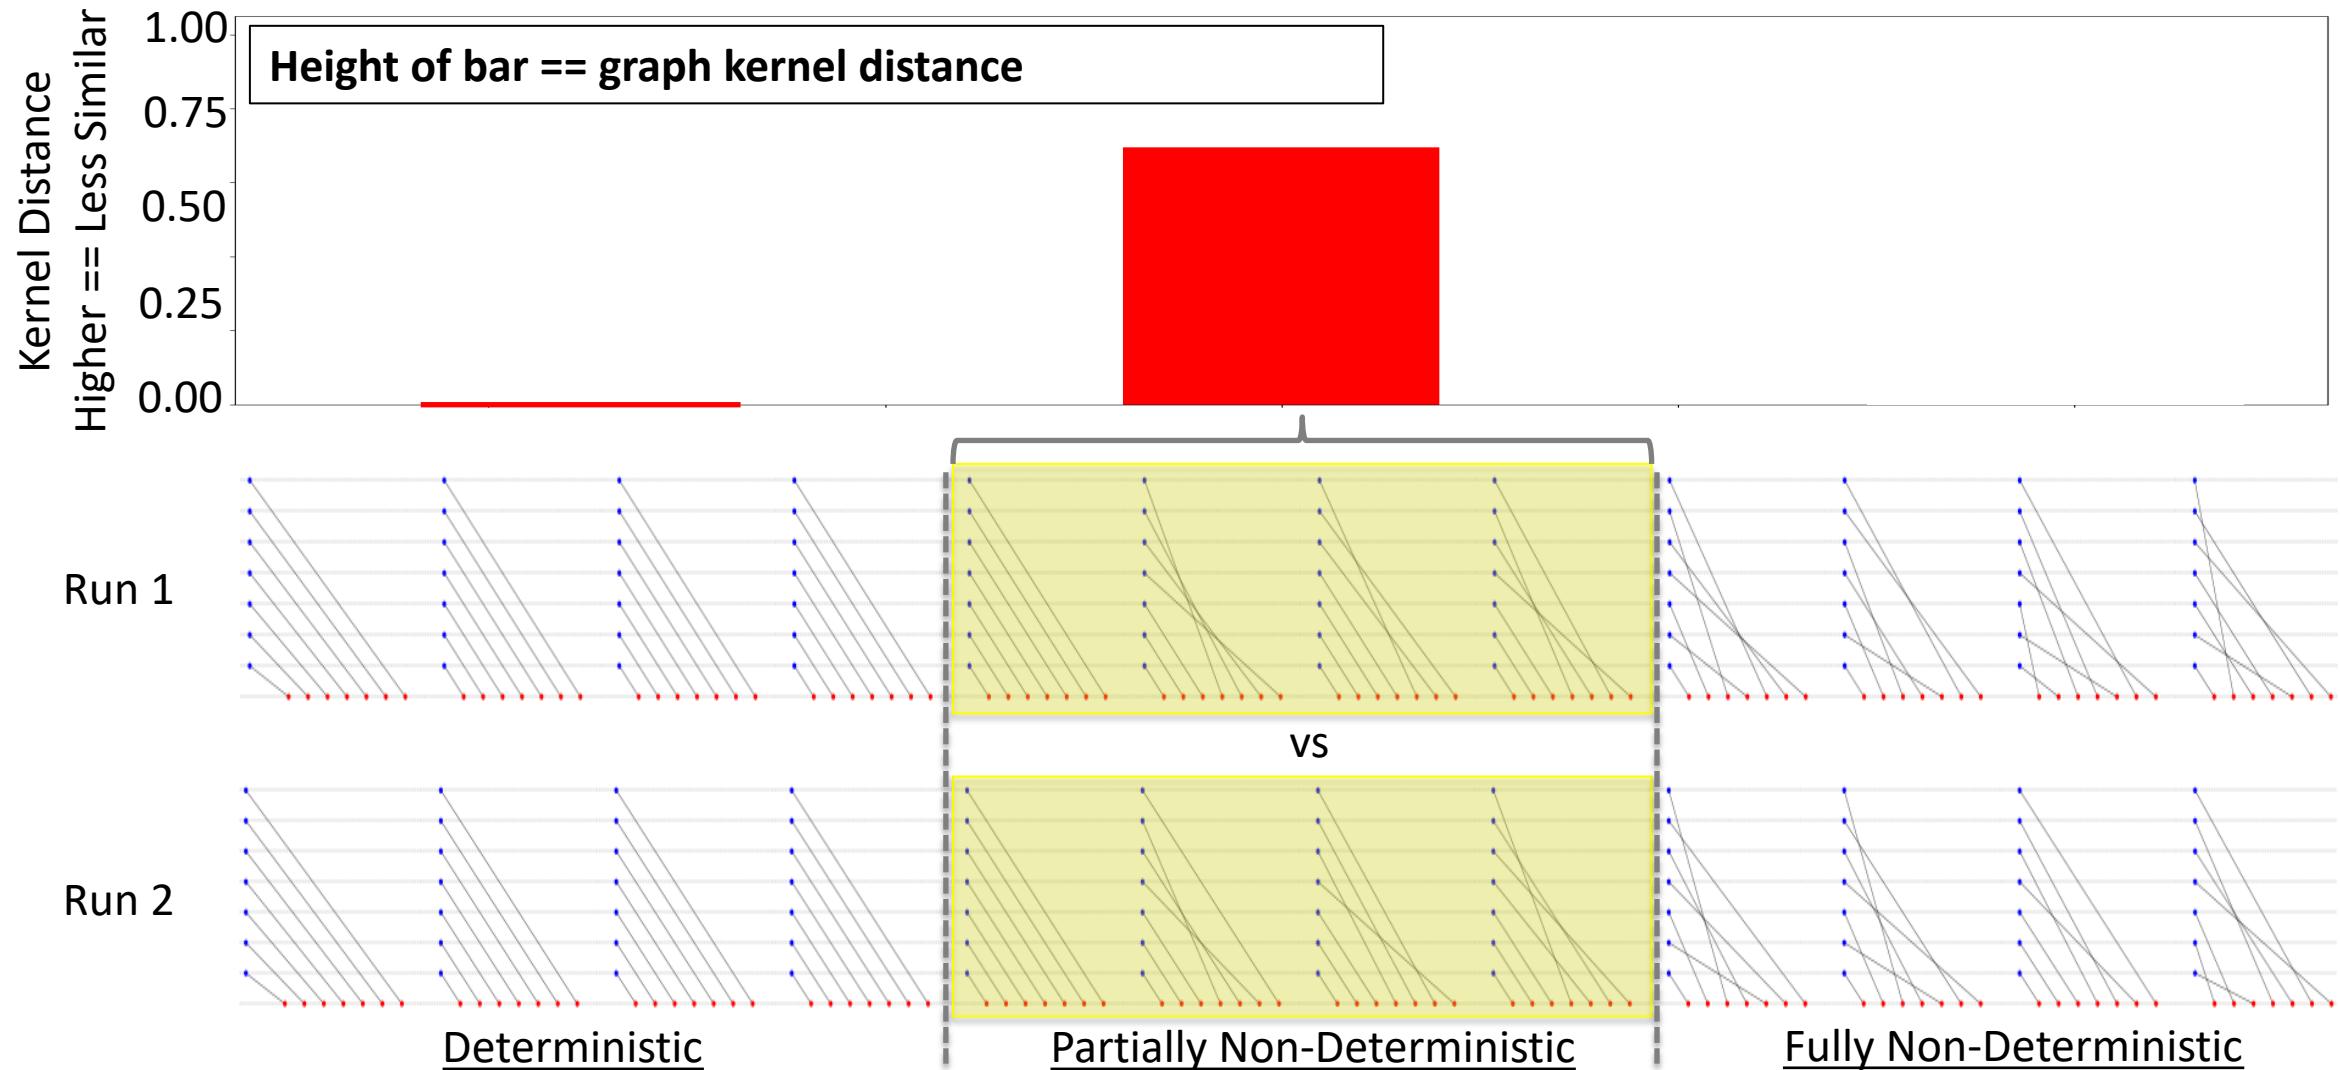
Comparing Runs via Kernel Distance



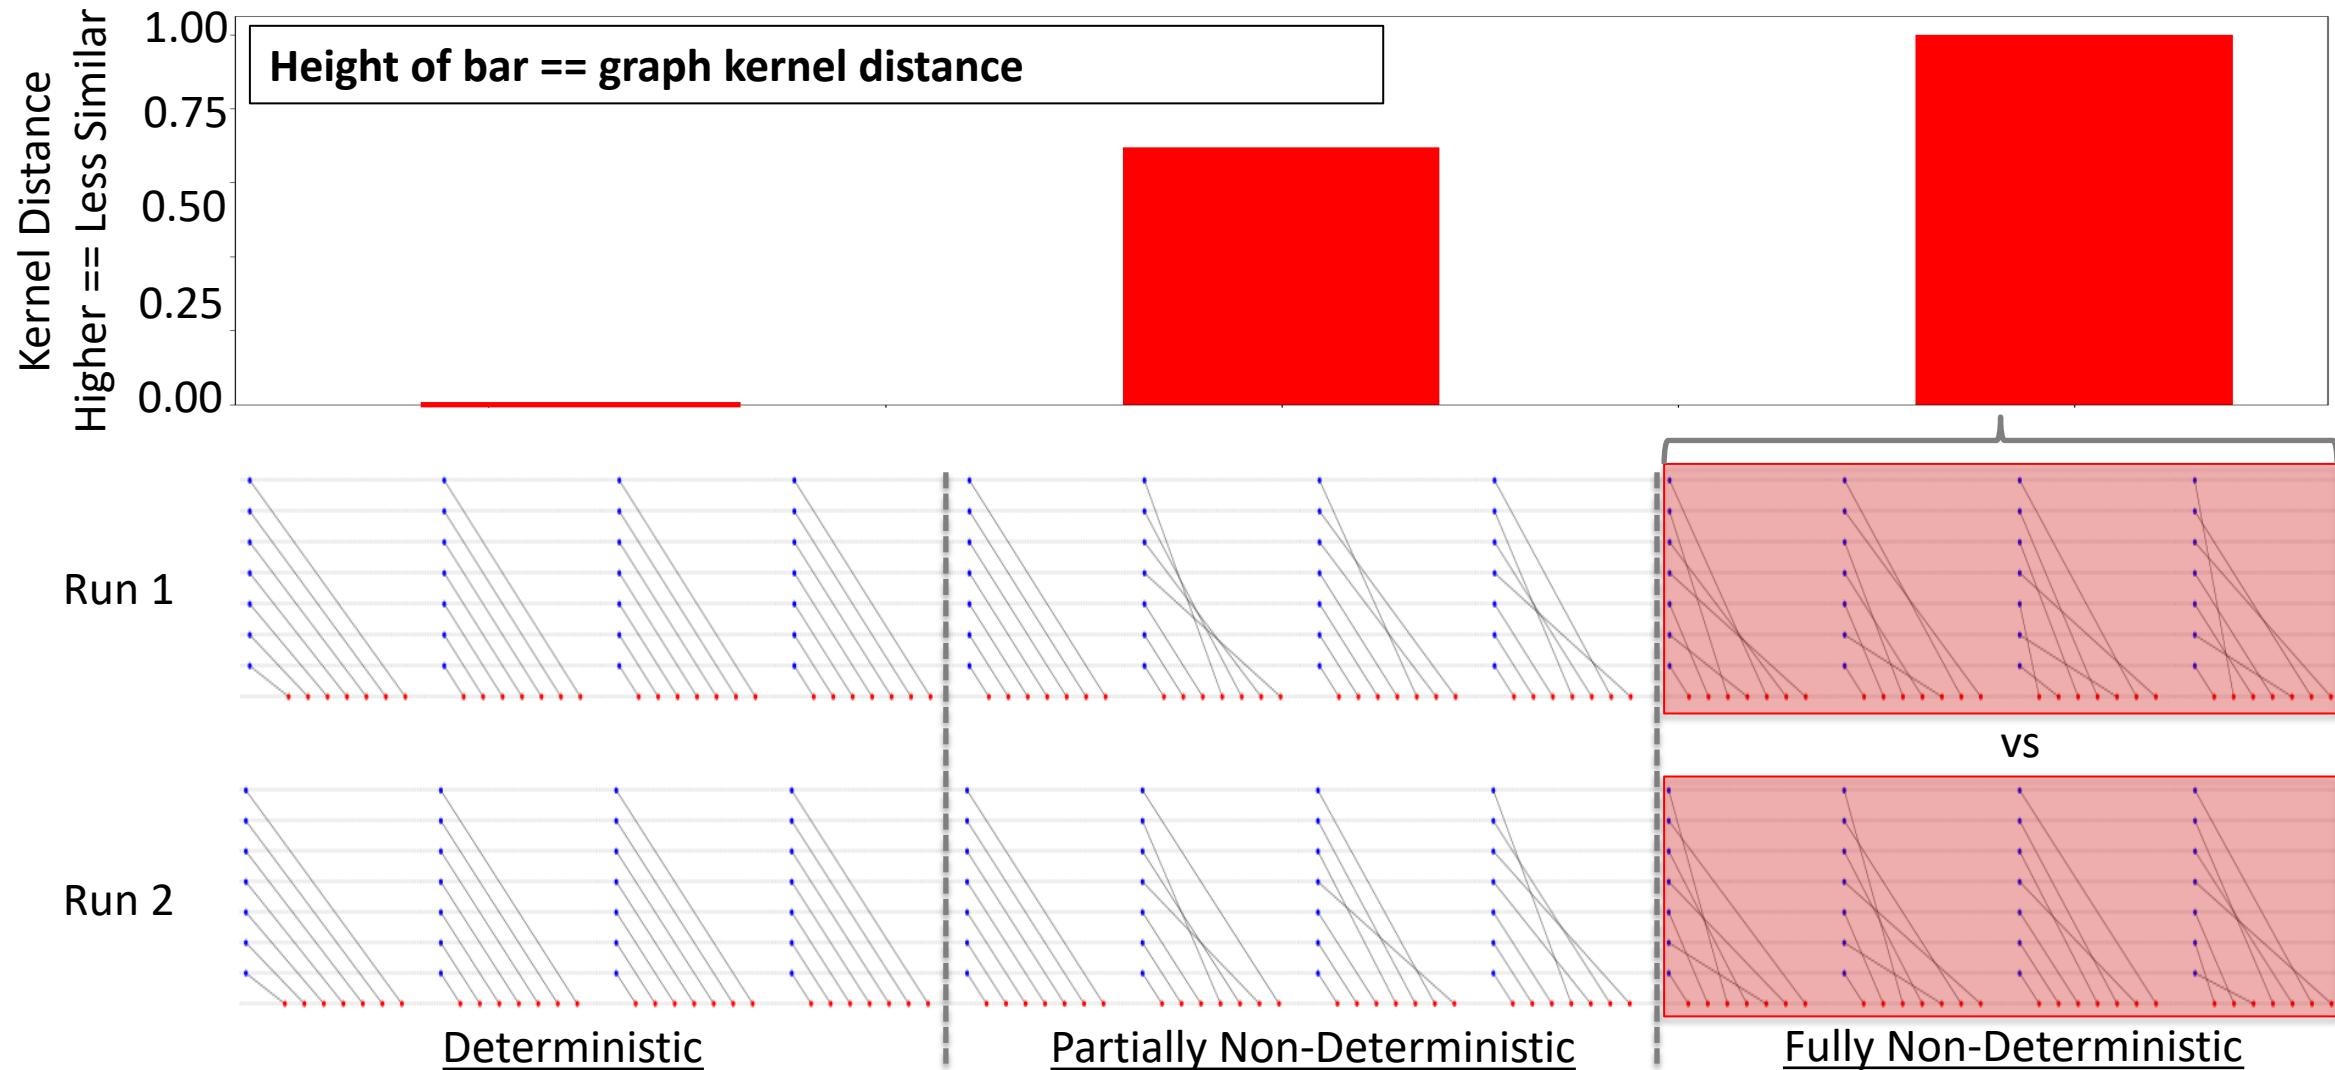
Comparing Runs via Kernel Distance



Comparing Runs via Kernel Distance

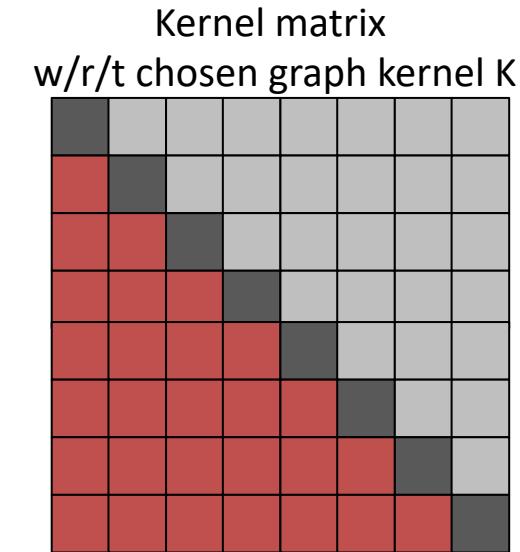
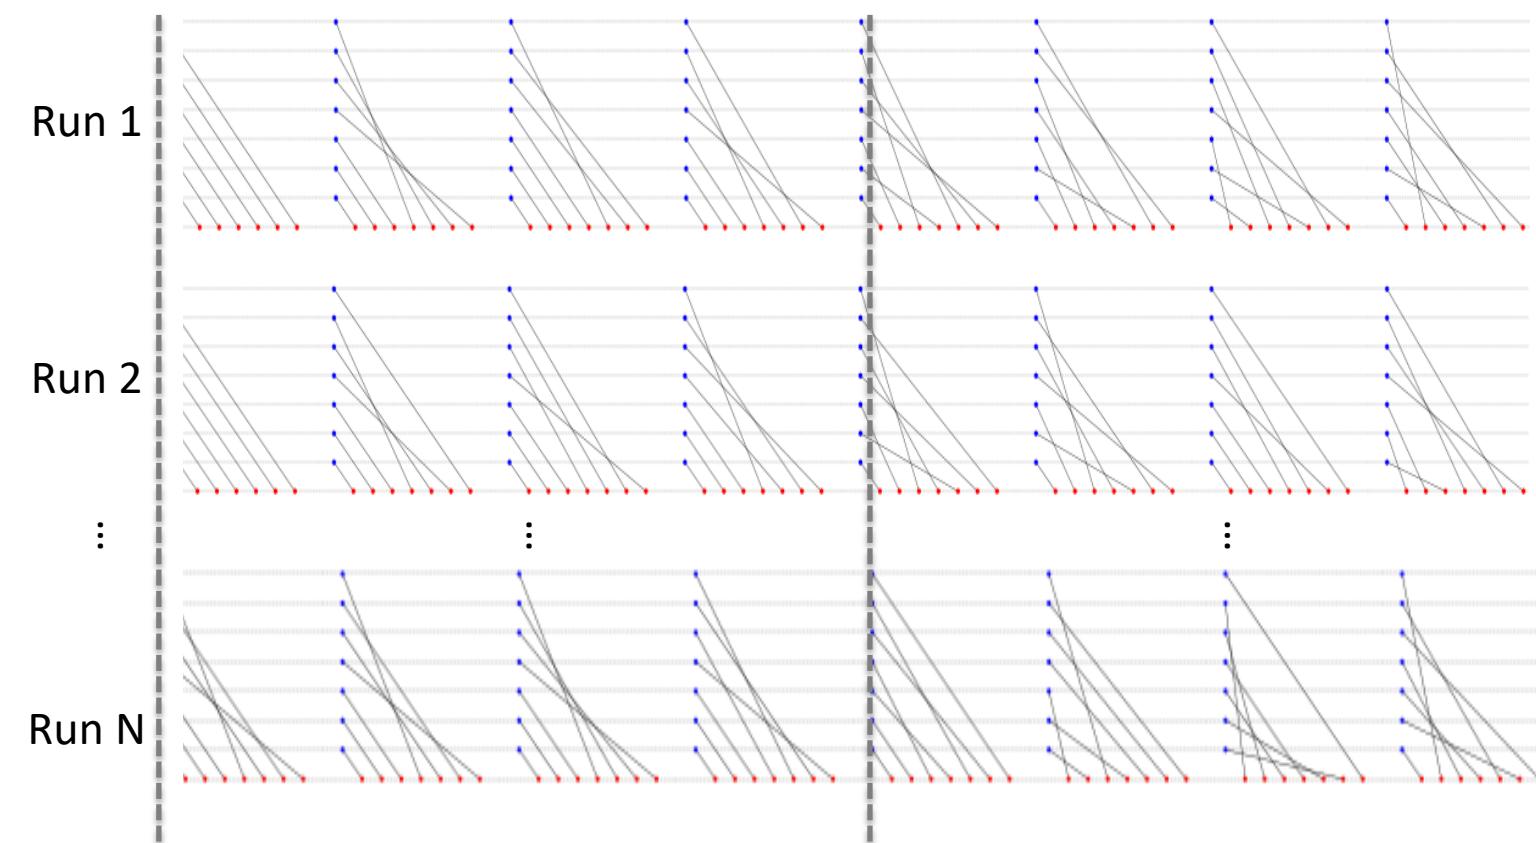


Comparing Runs via Kernel Distance



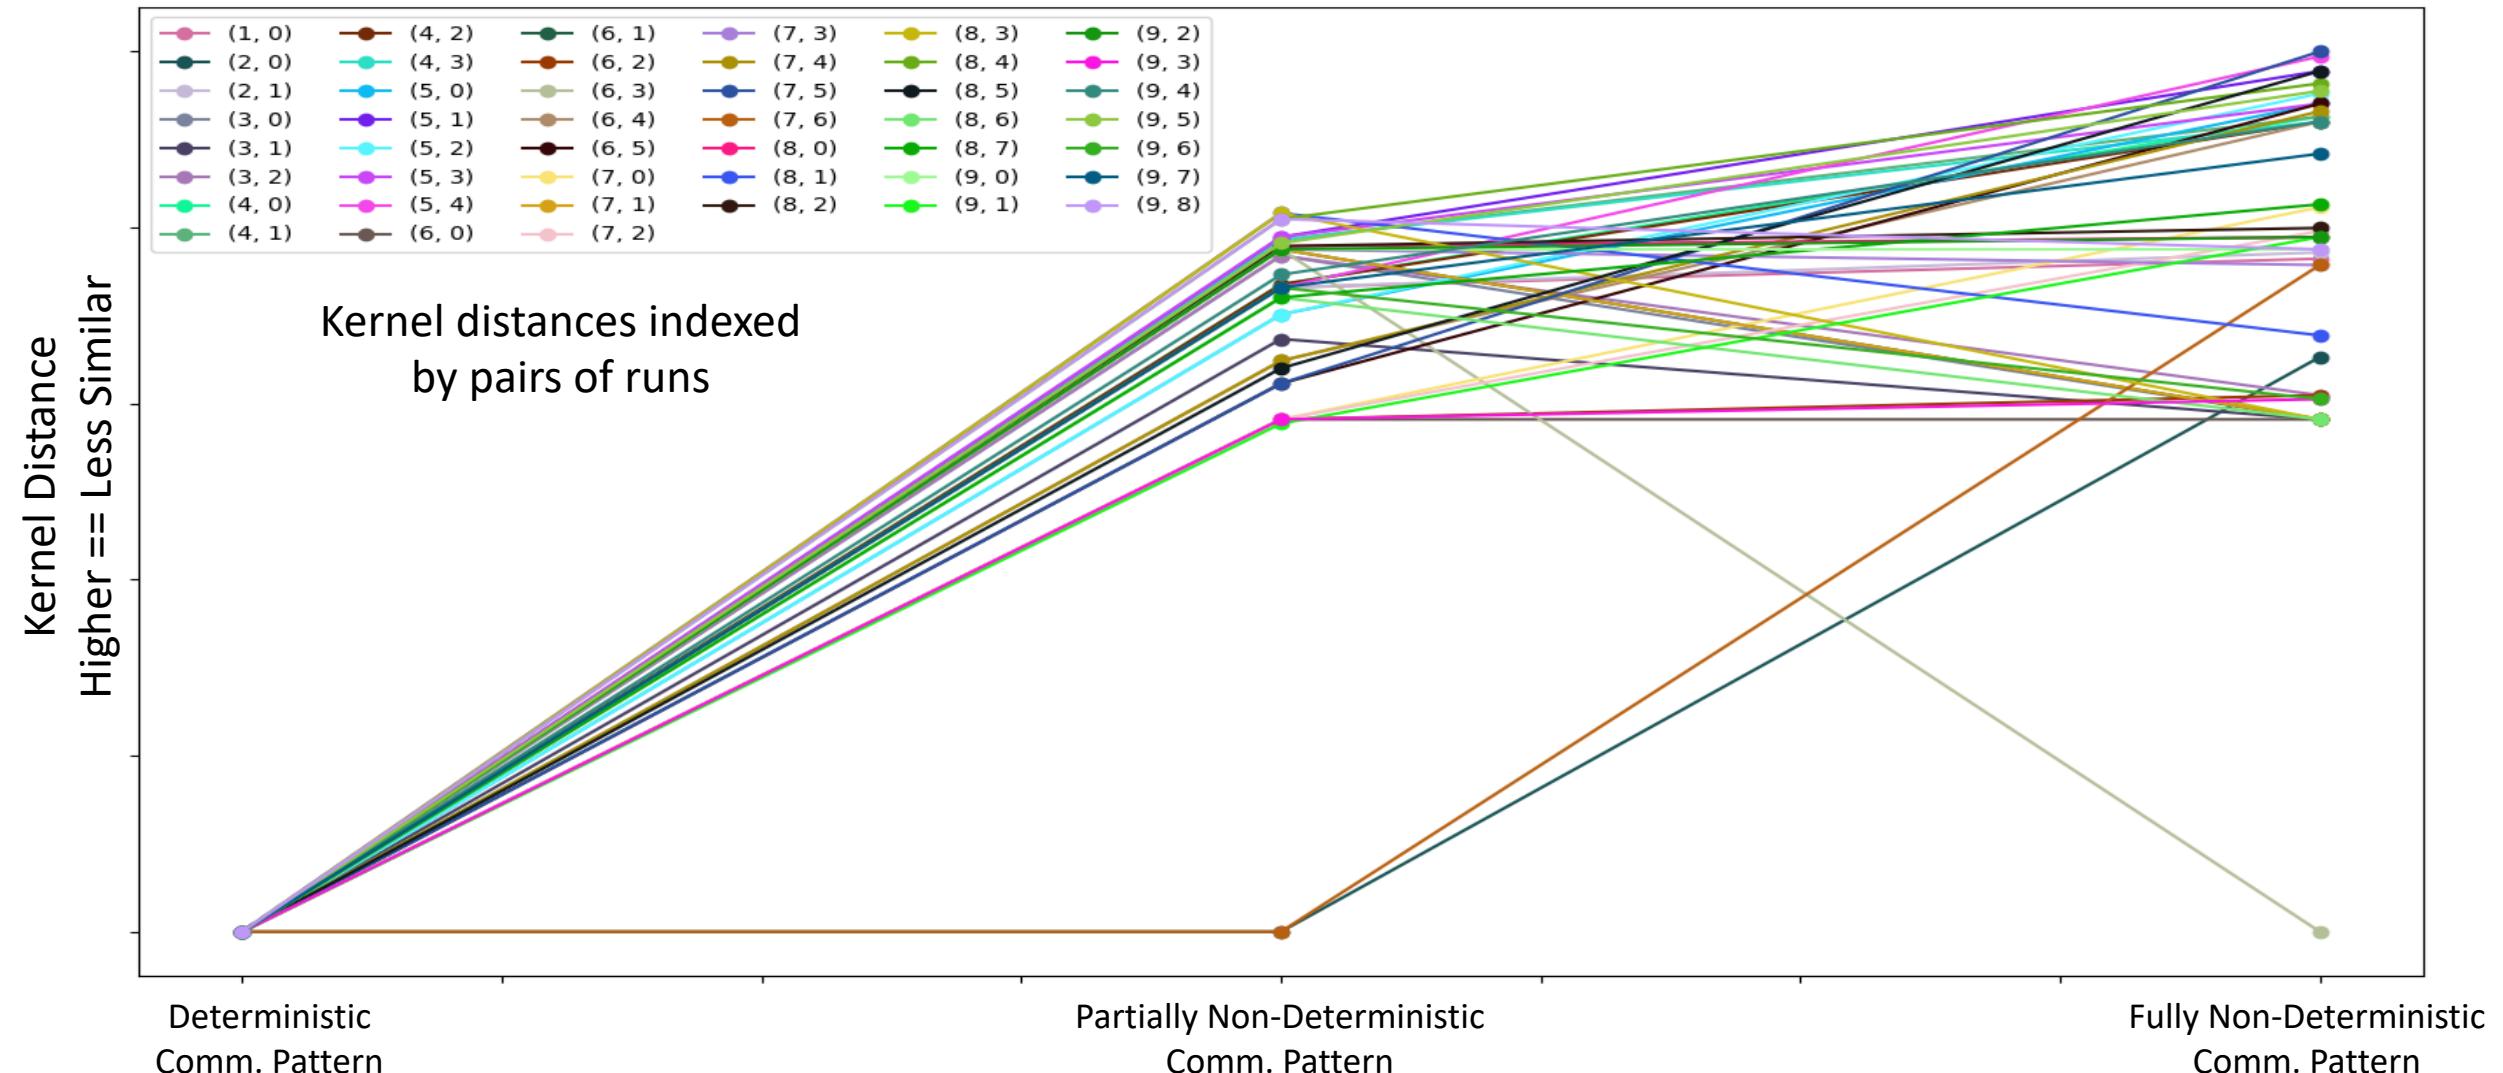
Kernel Distance Trends Across Multiple Runs

- Previous example only measured kernel distance between two executions
- But we want a statistical picture of the **trend in kernel distance over time across many executions**

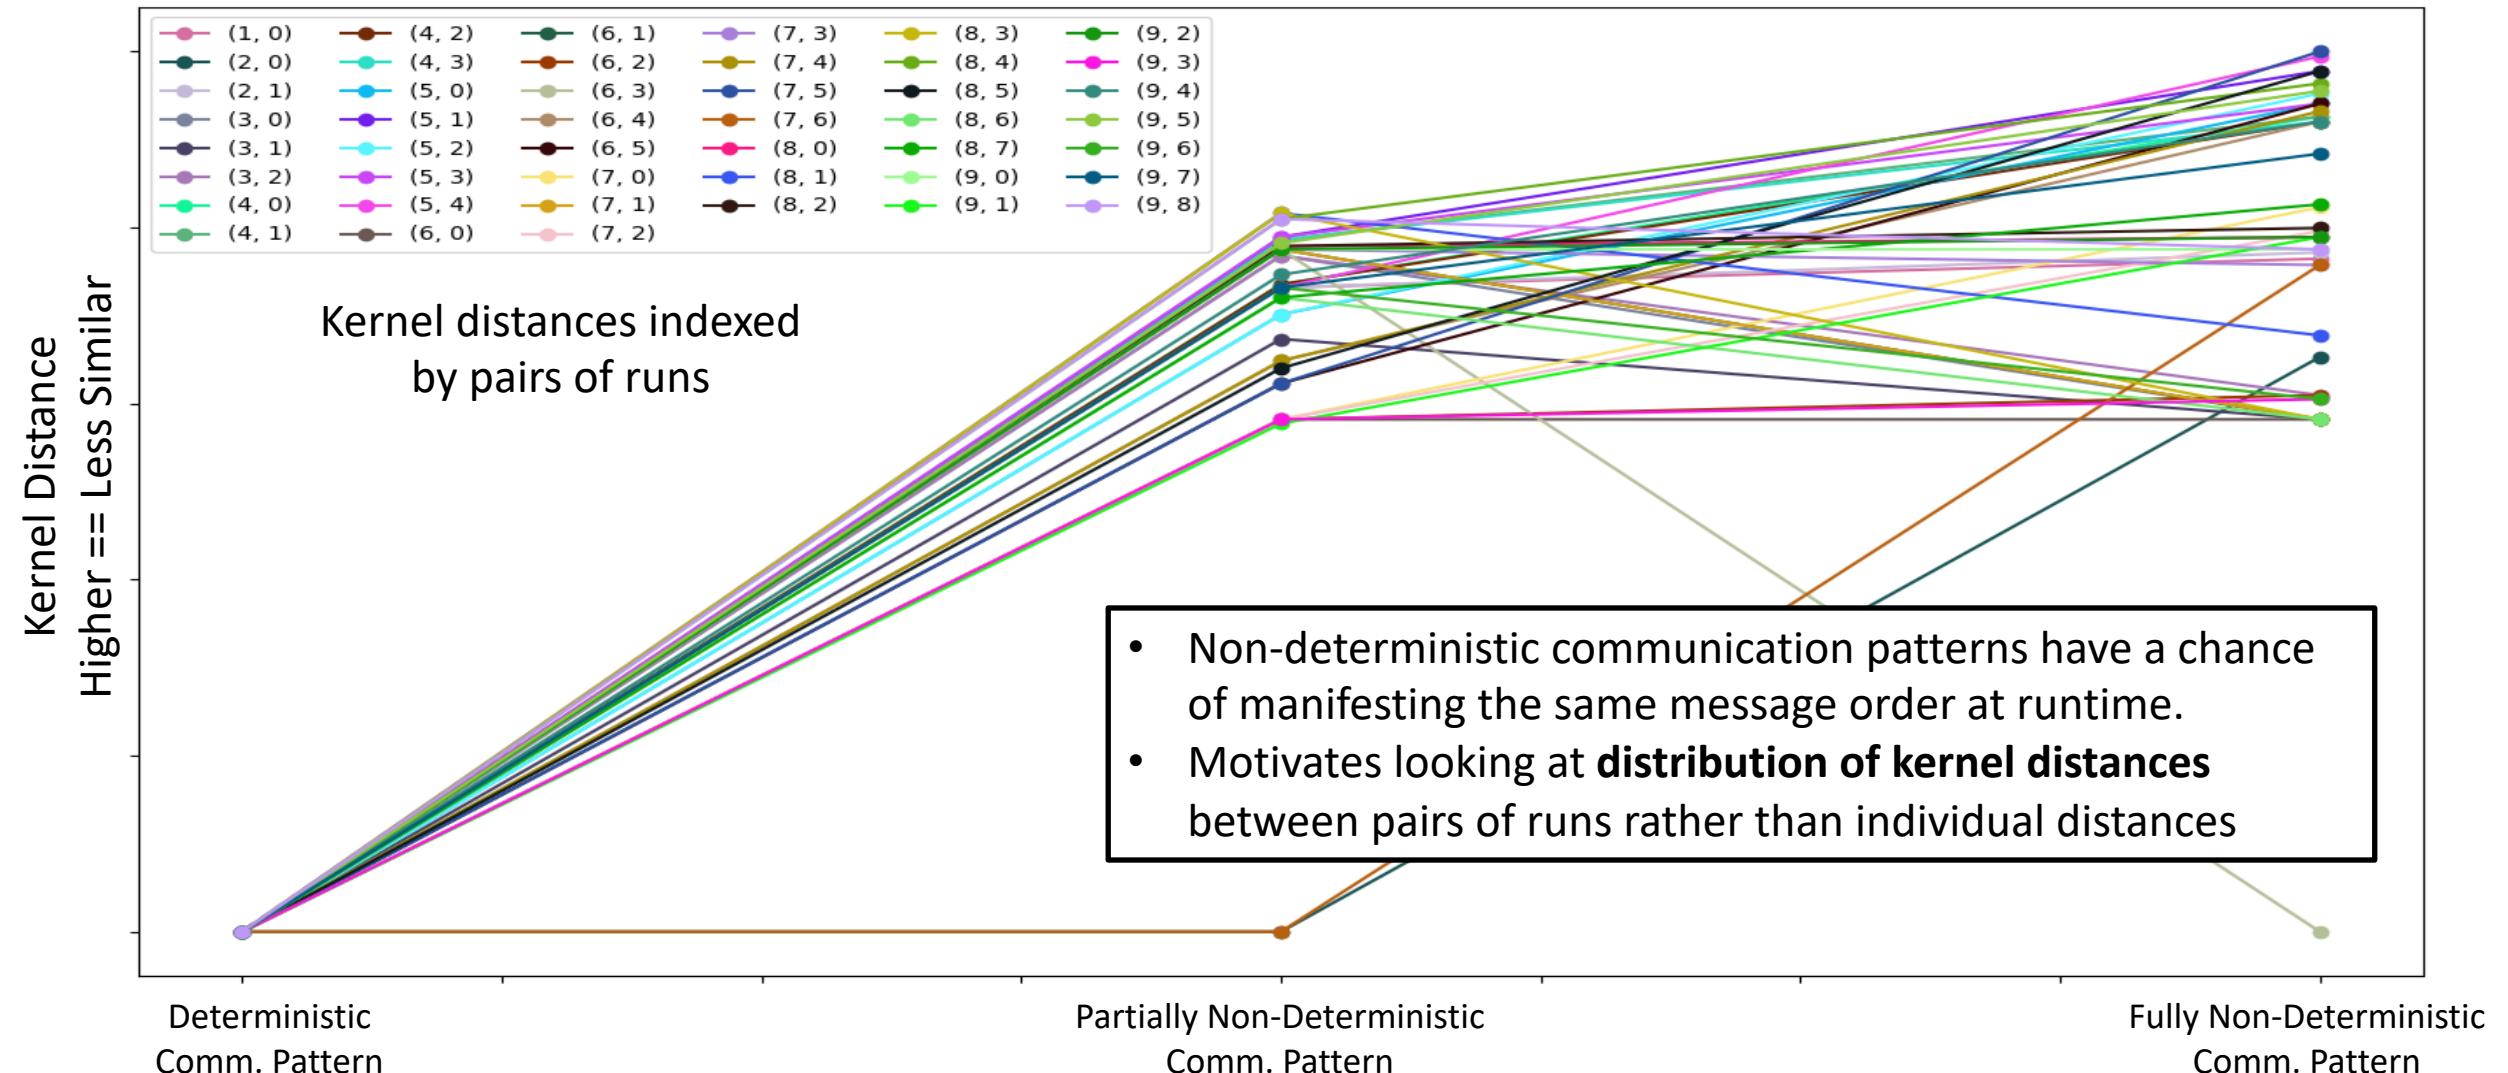


- = kernel distance between graphs i and j
- = kernel distance from graph to itself, always 0
- = redundant, due to symmetry

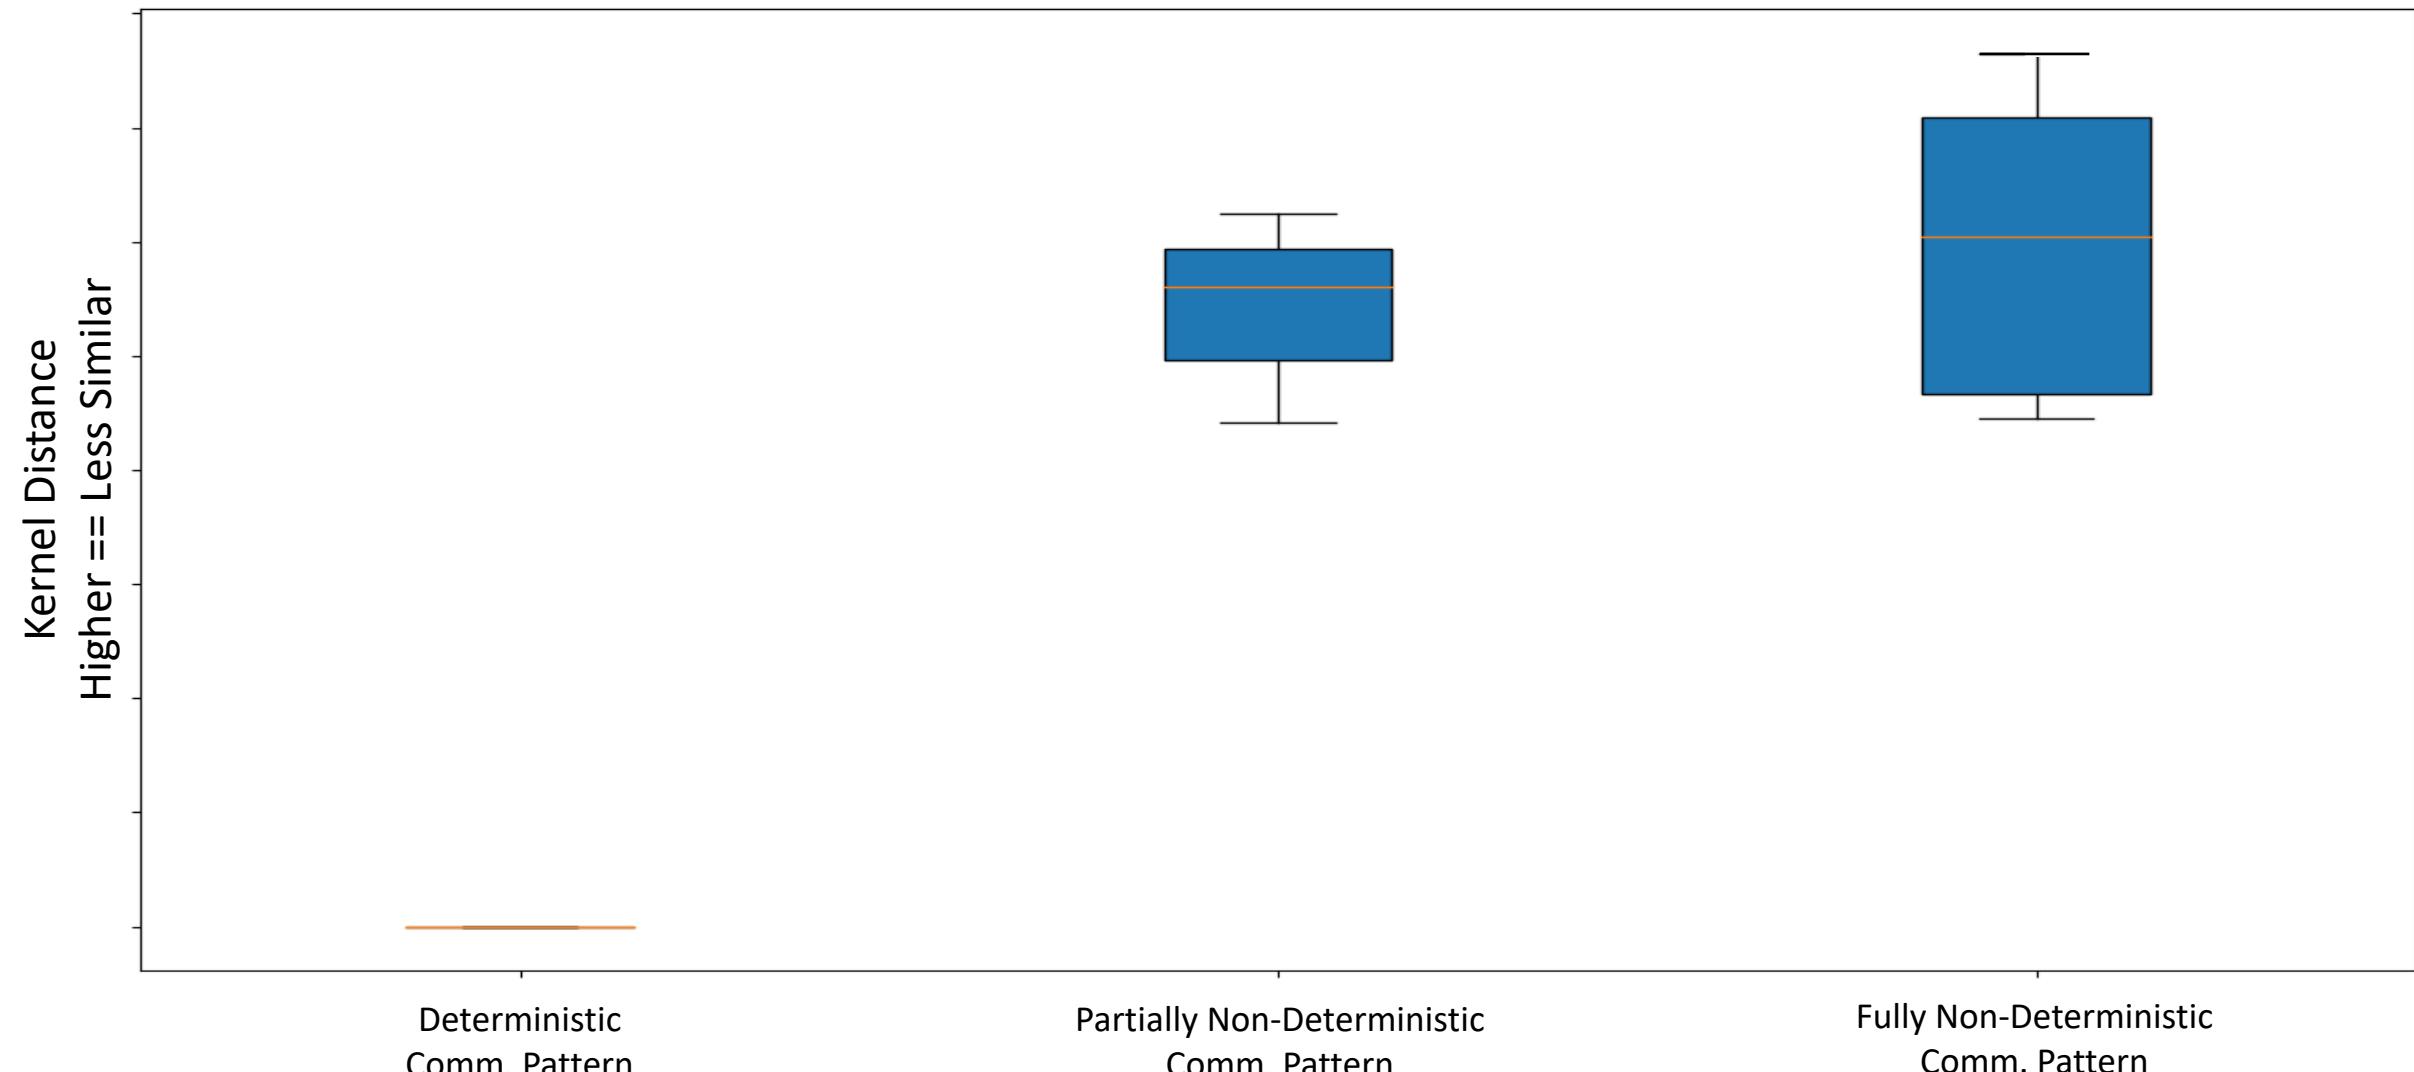
Kernel Distance Trends Across Multiple Runs



Kernel Distance Trends Across Multiple Runs



Kernel Distance Trends Across Multiple Runs



Evaluation Against a Realistic Communication Pattern

- We now evaluate against a communication pattern extracted from the AMG benchmark from the CORAL-2 Benchmark Suite

Evaluation Against a Realistic Communication Pattern

- We now evaluate against a communication pattern extracted from the AMG benchmark from the CORAL-2 Benchmark Suite
- Known to exhibit both receiver-side non-determinism and sender-side non-determinism [1]

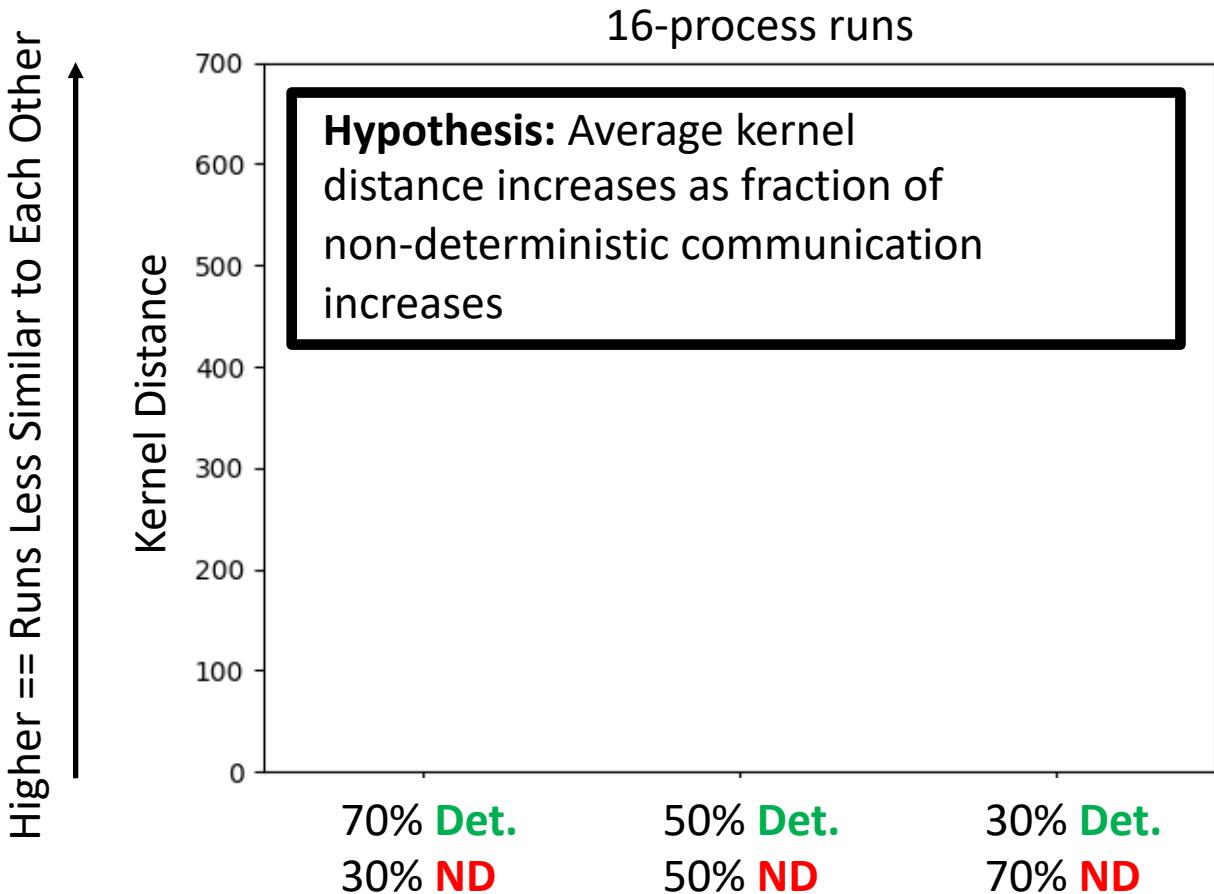


Evaluation Against a Realistic Communication Pattern

- We now evaluate against a communication pattern extracted from the AMG benchmark from the CORAL-2 Benchmark Suite
- Known to exhibit both receiver-side non-determinism and sender-side non-determinism [1]
- For our evaluation, we control proportion of non-deterministic communication volume by splitting ranks into two groups:
 - One group performs the actual non-deterministic AMG pattern
 - One group performs a determinized version of the pattern

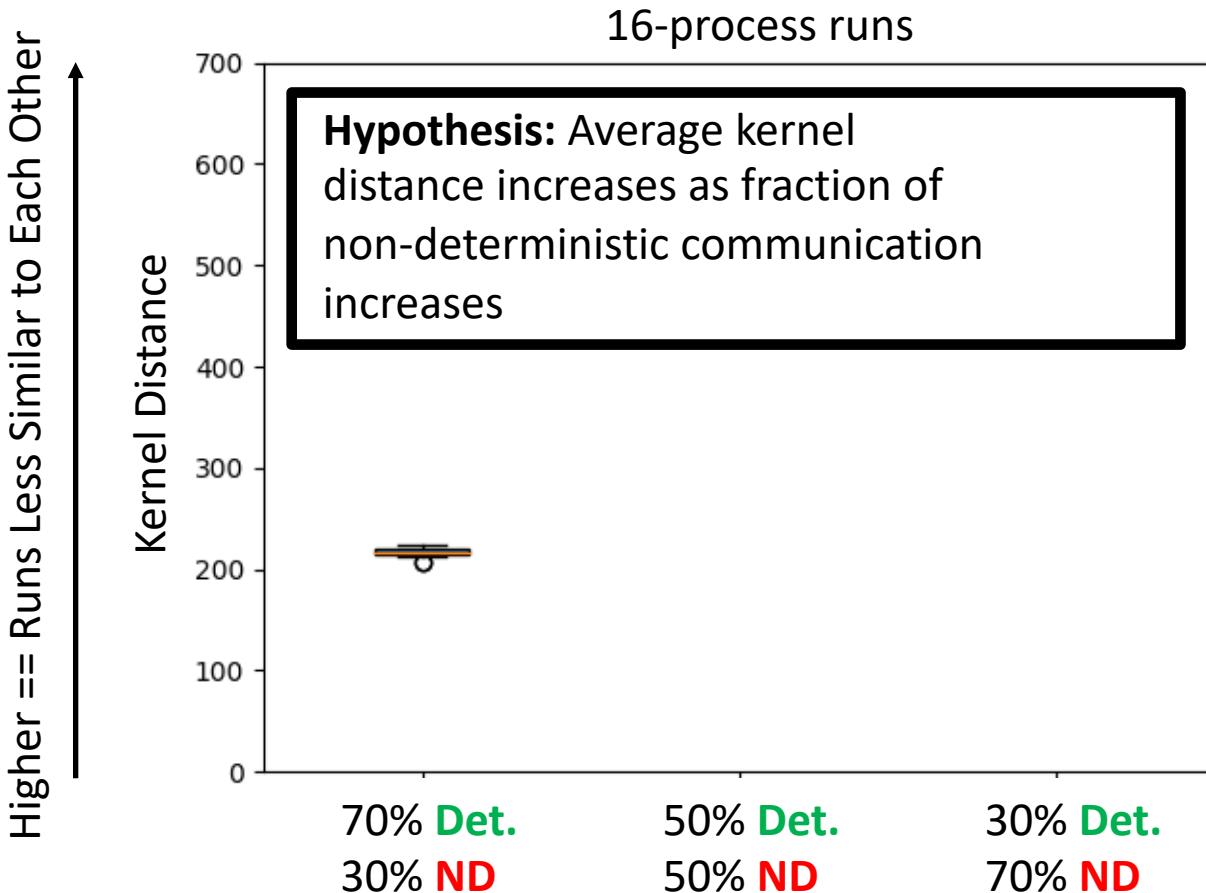


Kernel Distance Between Runs of Sequoia-AMG



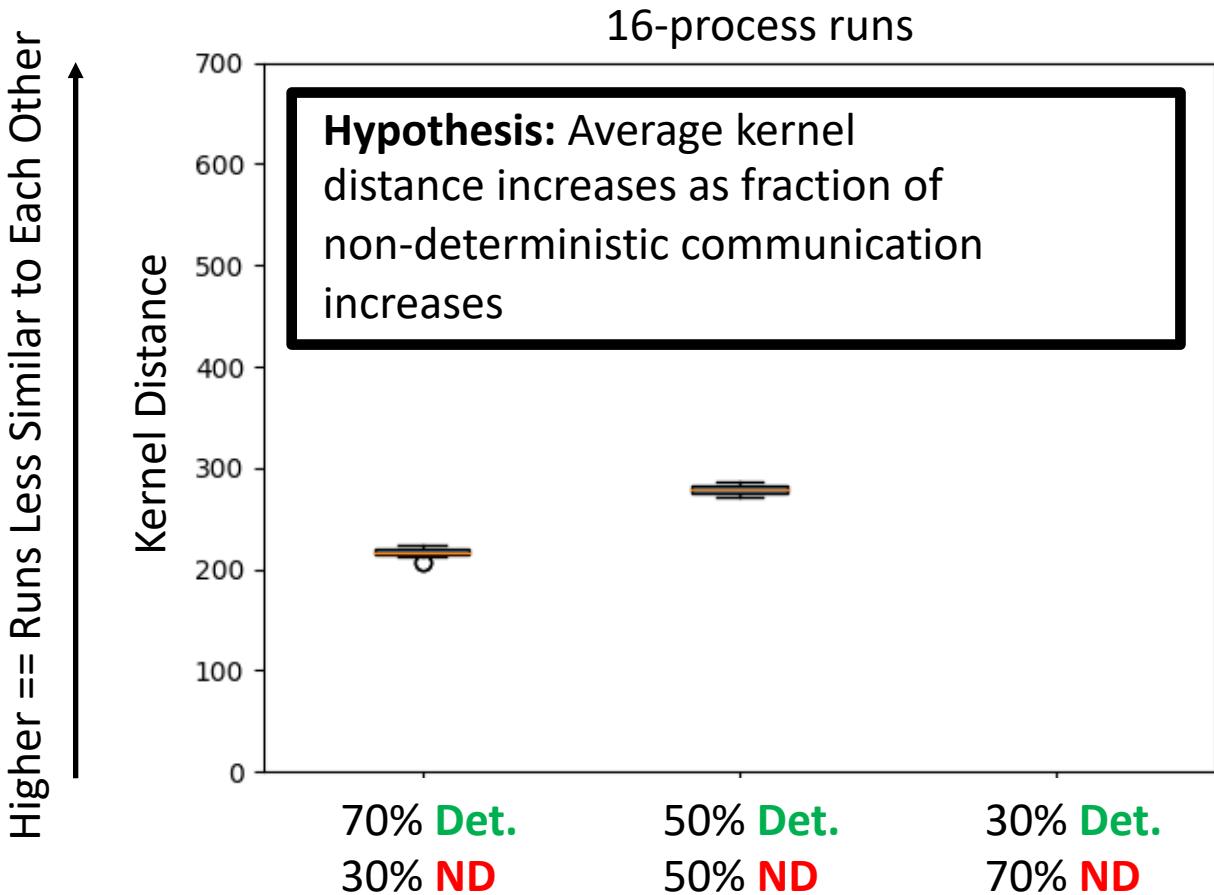
Graph Kernel Used: Weisfeiler-Lehman Subtree Kernel [1]

Kernel Distance Between Runs of Sequoia-AMG



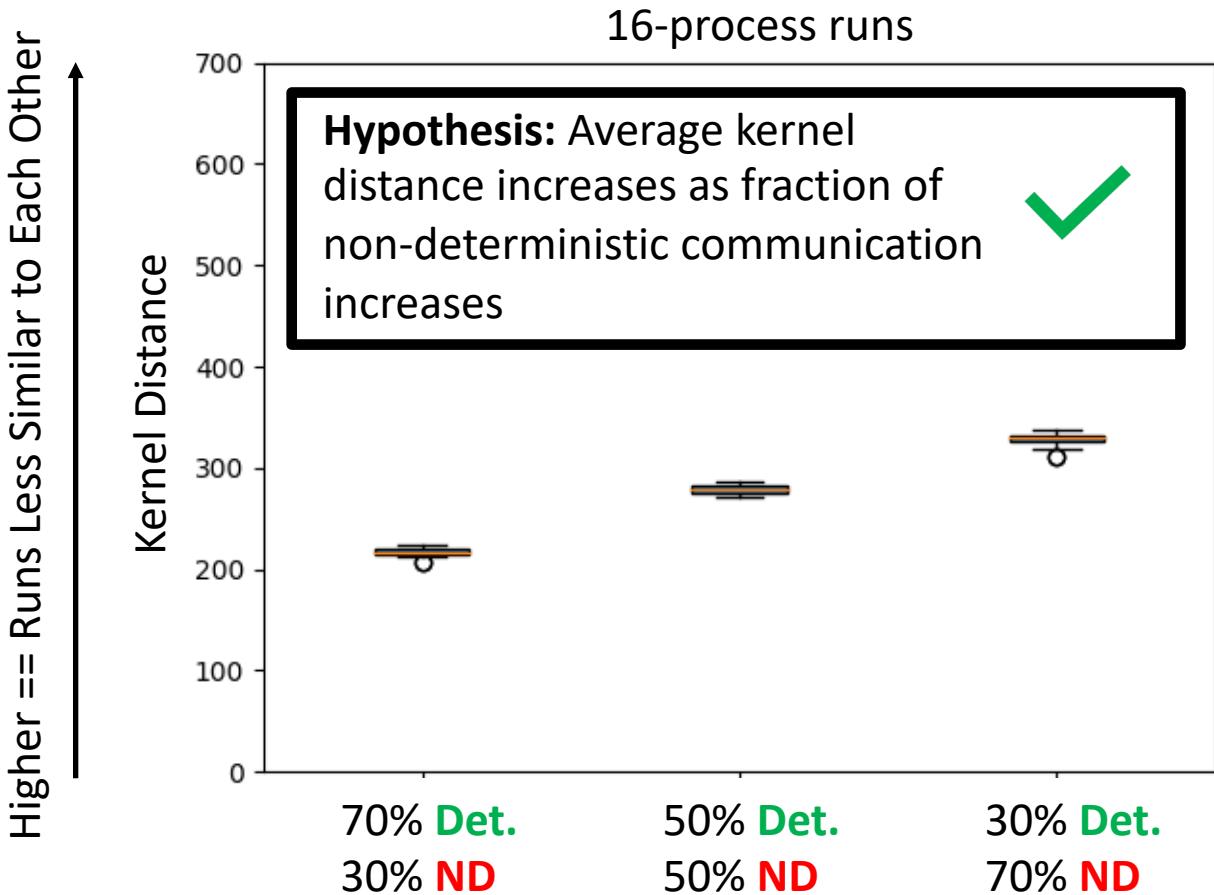
Graph Kernel Used: Weisfeiler-Lehman Subtree Kernel [1]

Kernel Distance Between Runs of Sequoia-AMG



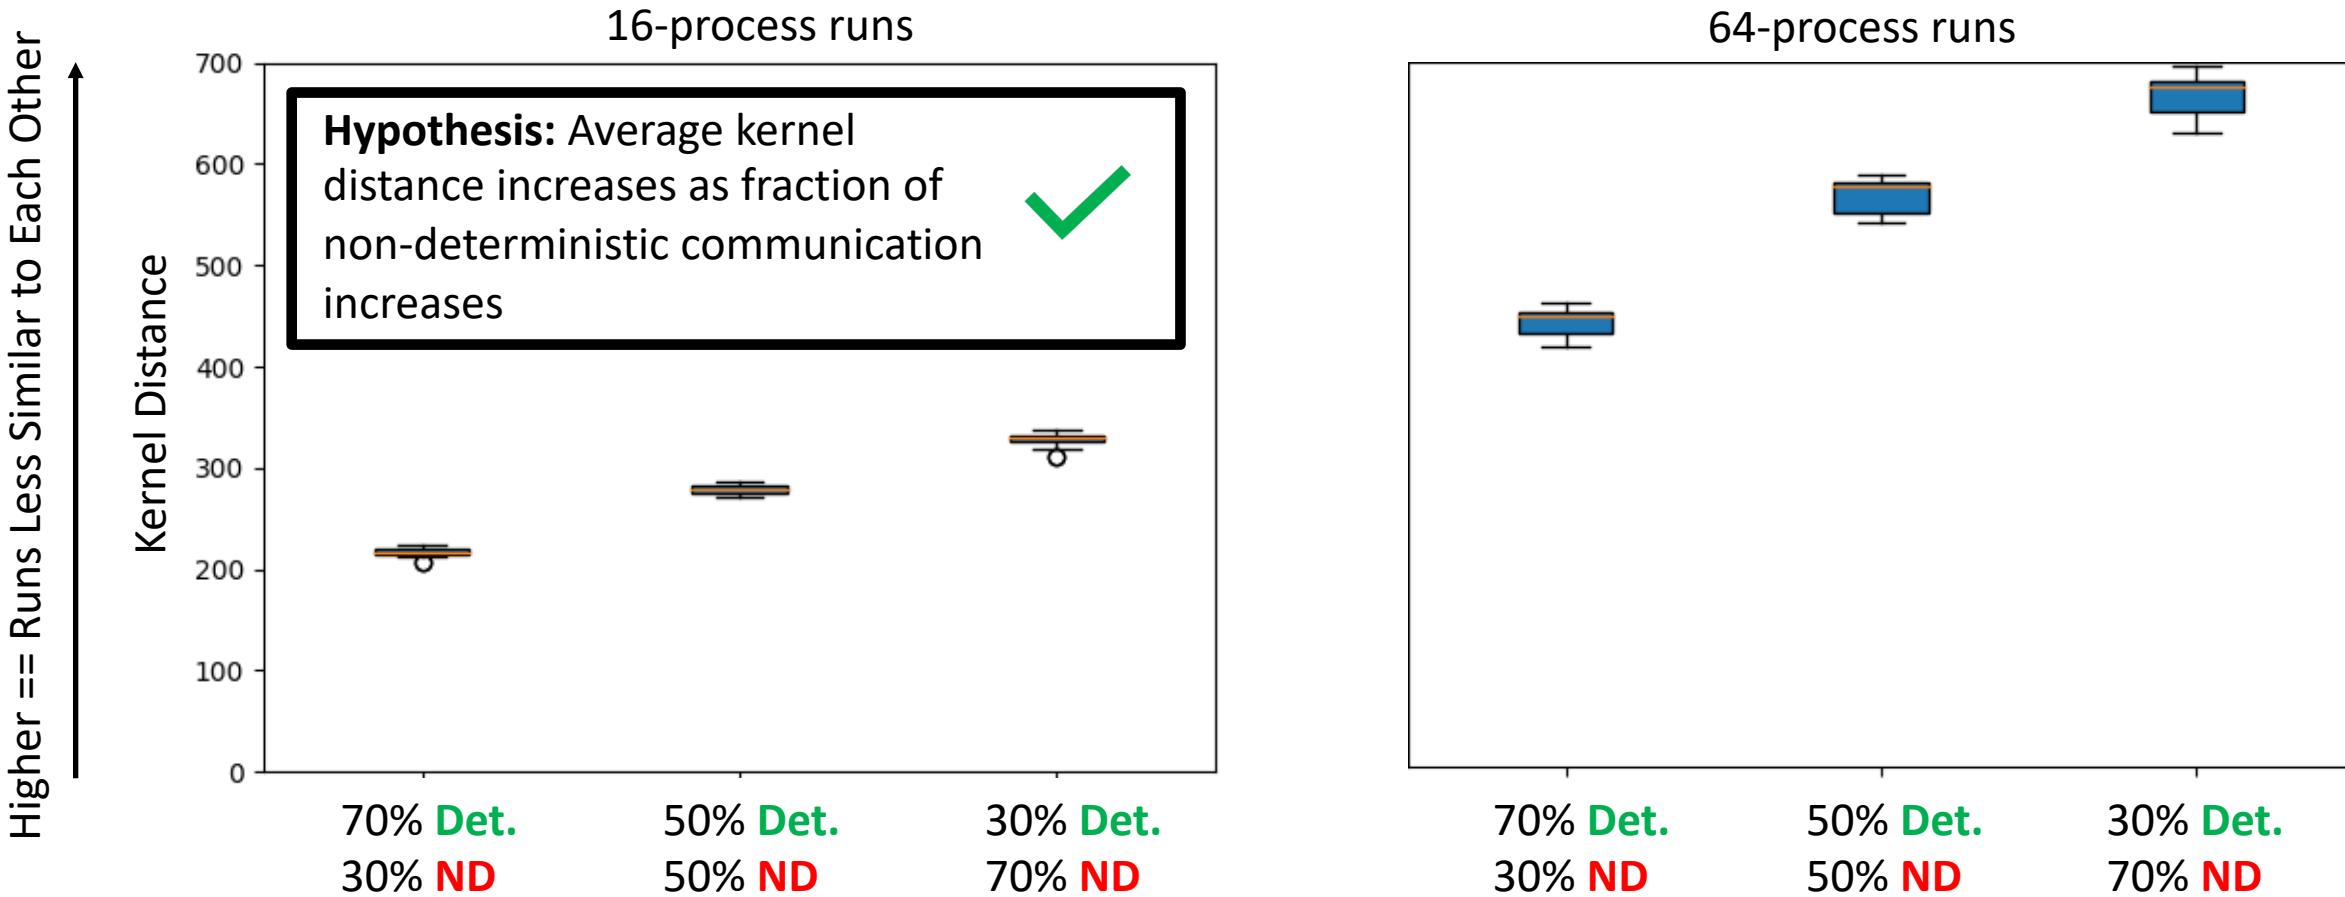
Graph Kernel Used: Weisfeiler-Lehman Subtree Kernel [1]

Kernel Distance Between Runs of Sequoia-AMG



Graph Kernel Used: Weisfeiler-Lehman Subtree Kernel [1]

Kernel Distance Between Runs of Sequoia-AMG



Graph Kernel Used: Weisfeiler-Lehman Subtree Kernel [1]

Lessons Learned

- Graph kernel distance is a useful proxy for degree of non-determinism in a communication pattern

Lessons Learned

- Graph kernel distance is a useful proxy for degree of non-determinism in a communication pattern
- This holds for both:
 - Simple communication patterns with only receiver-side non-determinism (e.g., naïve reduce pattern)
 - More complex patterns with mixed receiver/sender-side non-determinism (e.g., Sequoia-AMG Iprobe pattern)

Workflow for Non-Determinism Characterization

- **Phase 1:** Build graph-structured models of executions
- **Phase 2:** Quantify cross-execution trends in non-deterministic communication via graph similarity
- **Phase 3:** Detect periods of anomalous execution dissimilarity and localize potential root causes

Applying our Workflow to miniAMR

- Adaptive mesh refinement (AMR) code from the Matevo Benchmark Suite [1]

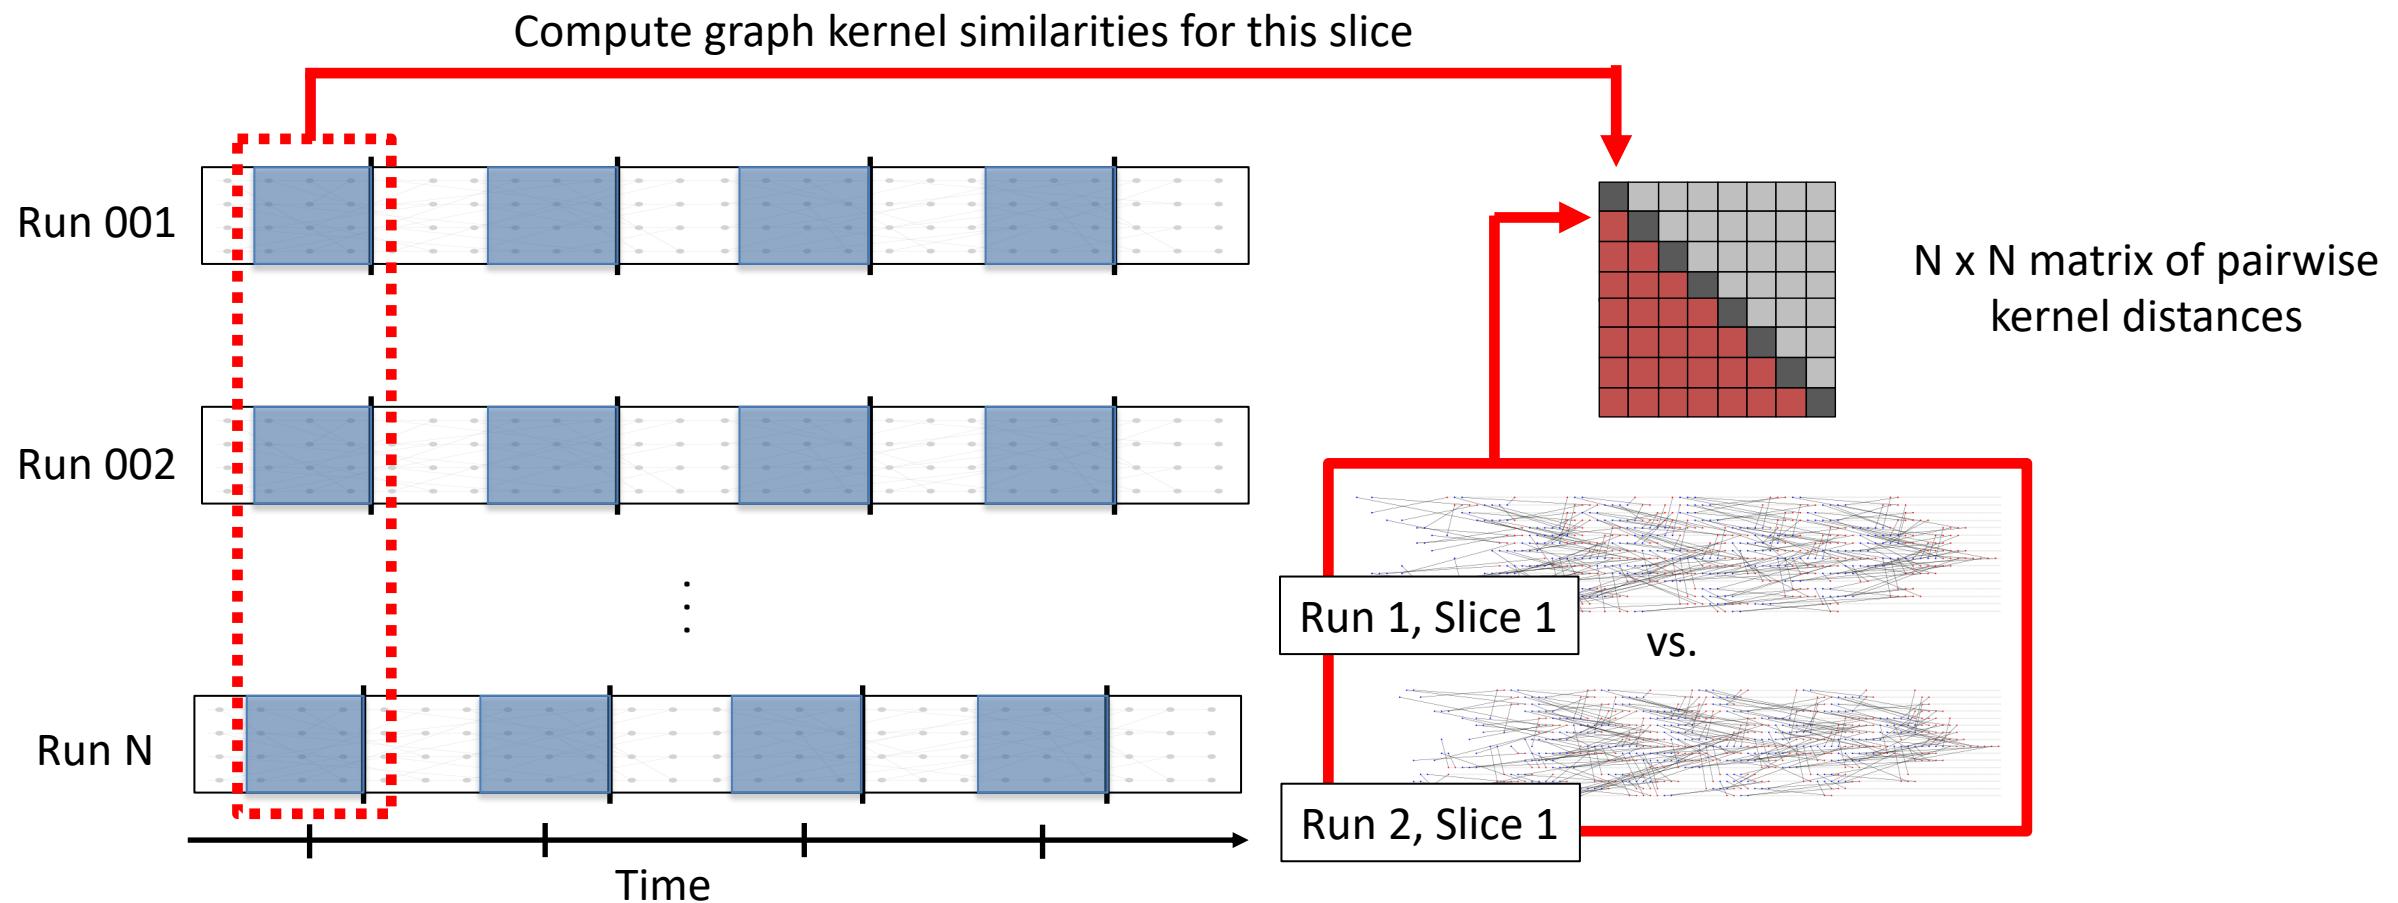


Applying our Workflow to miniAMR

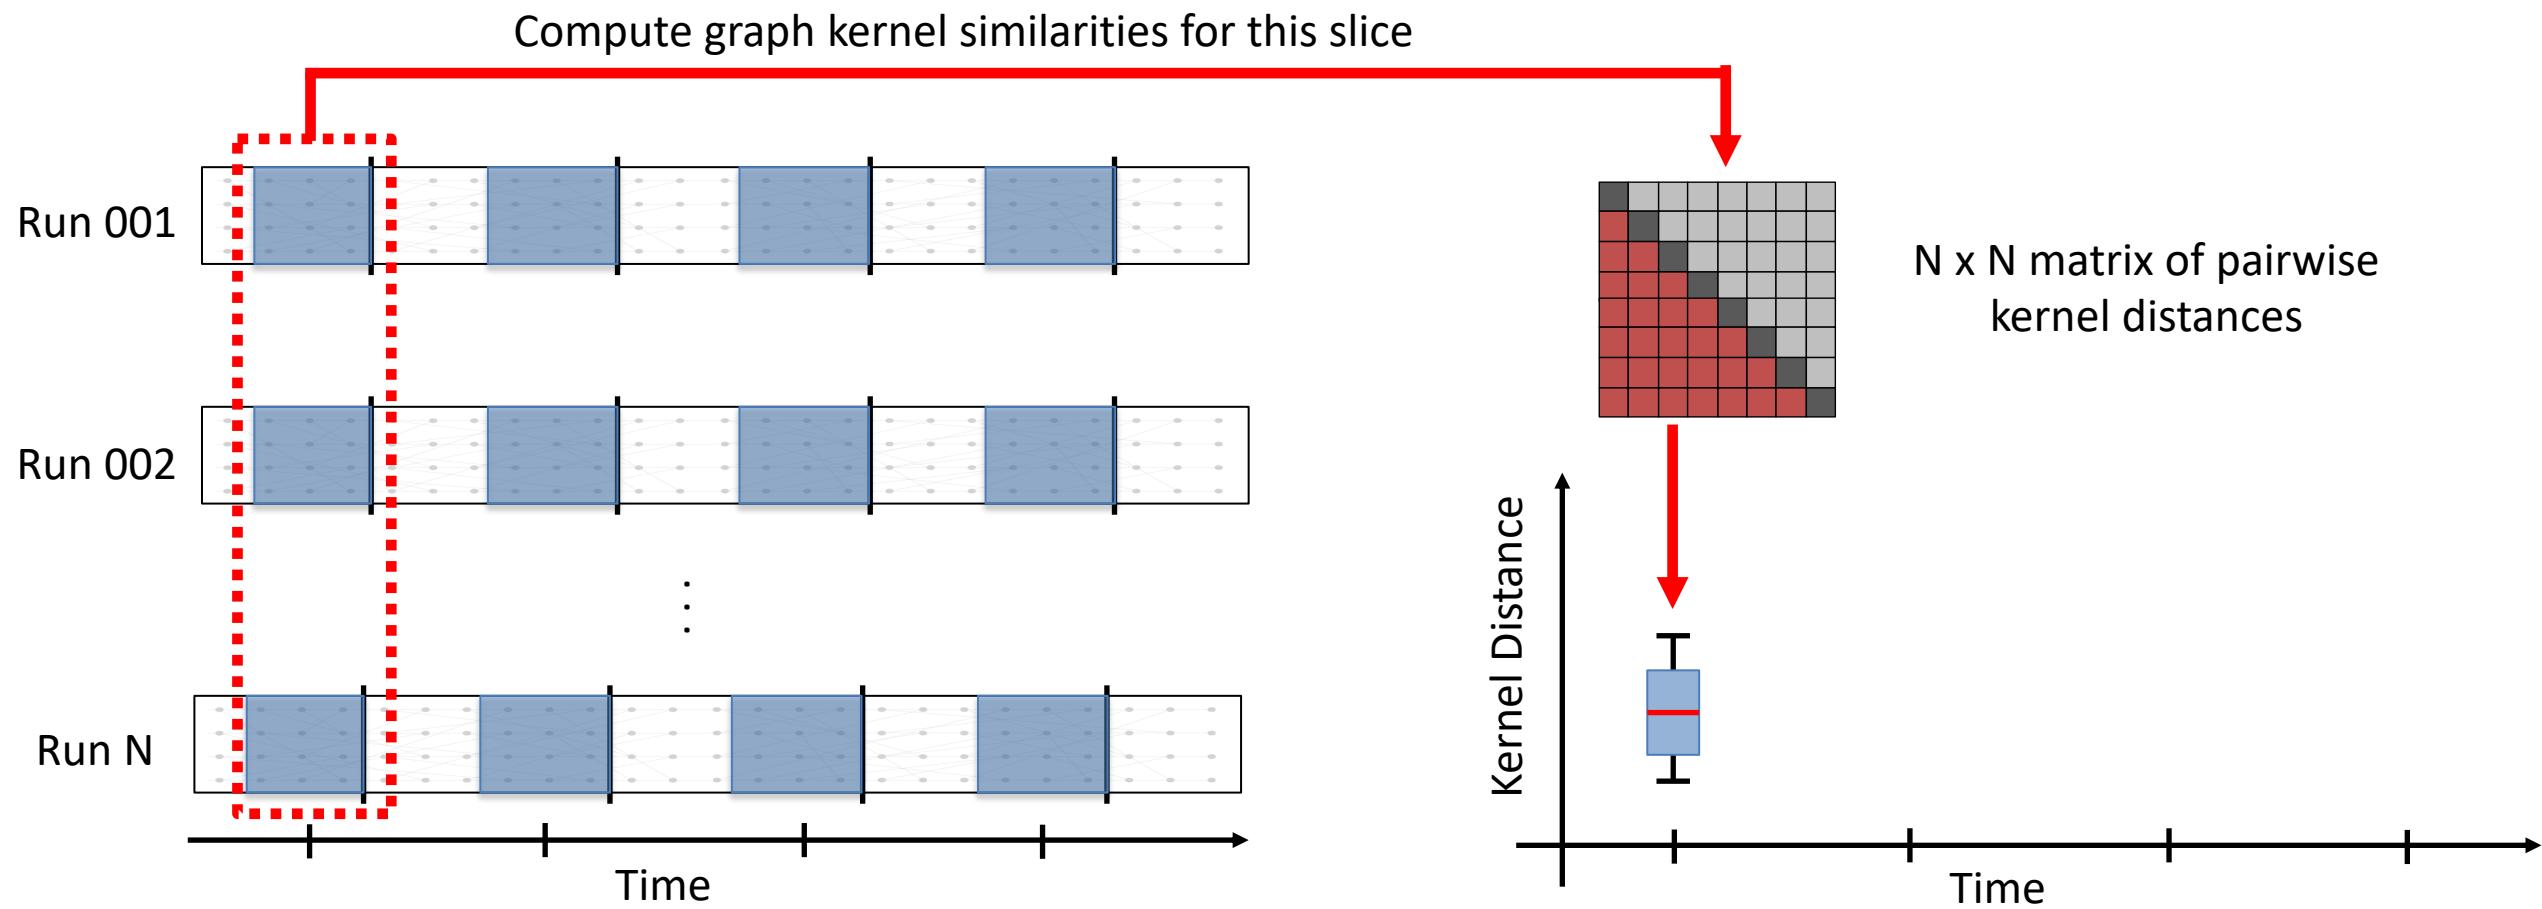
- Adaptive mesh refinement (AMR) code from the Matevo Benchmark Suite [1]
- We target miniAMR based on the following criteria:
 - MPI application exhibiting communication non-determinism
 - Root cause of non-determinism known *a priori*



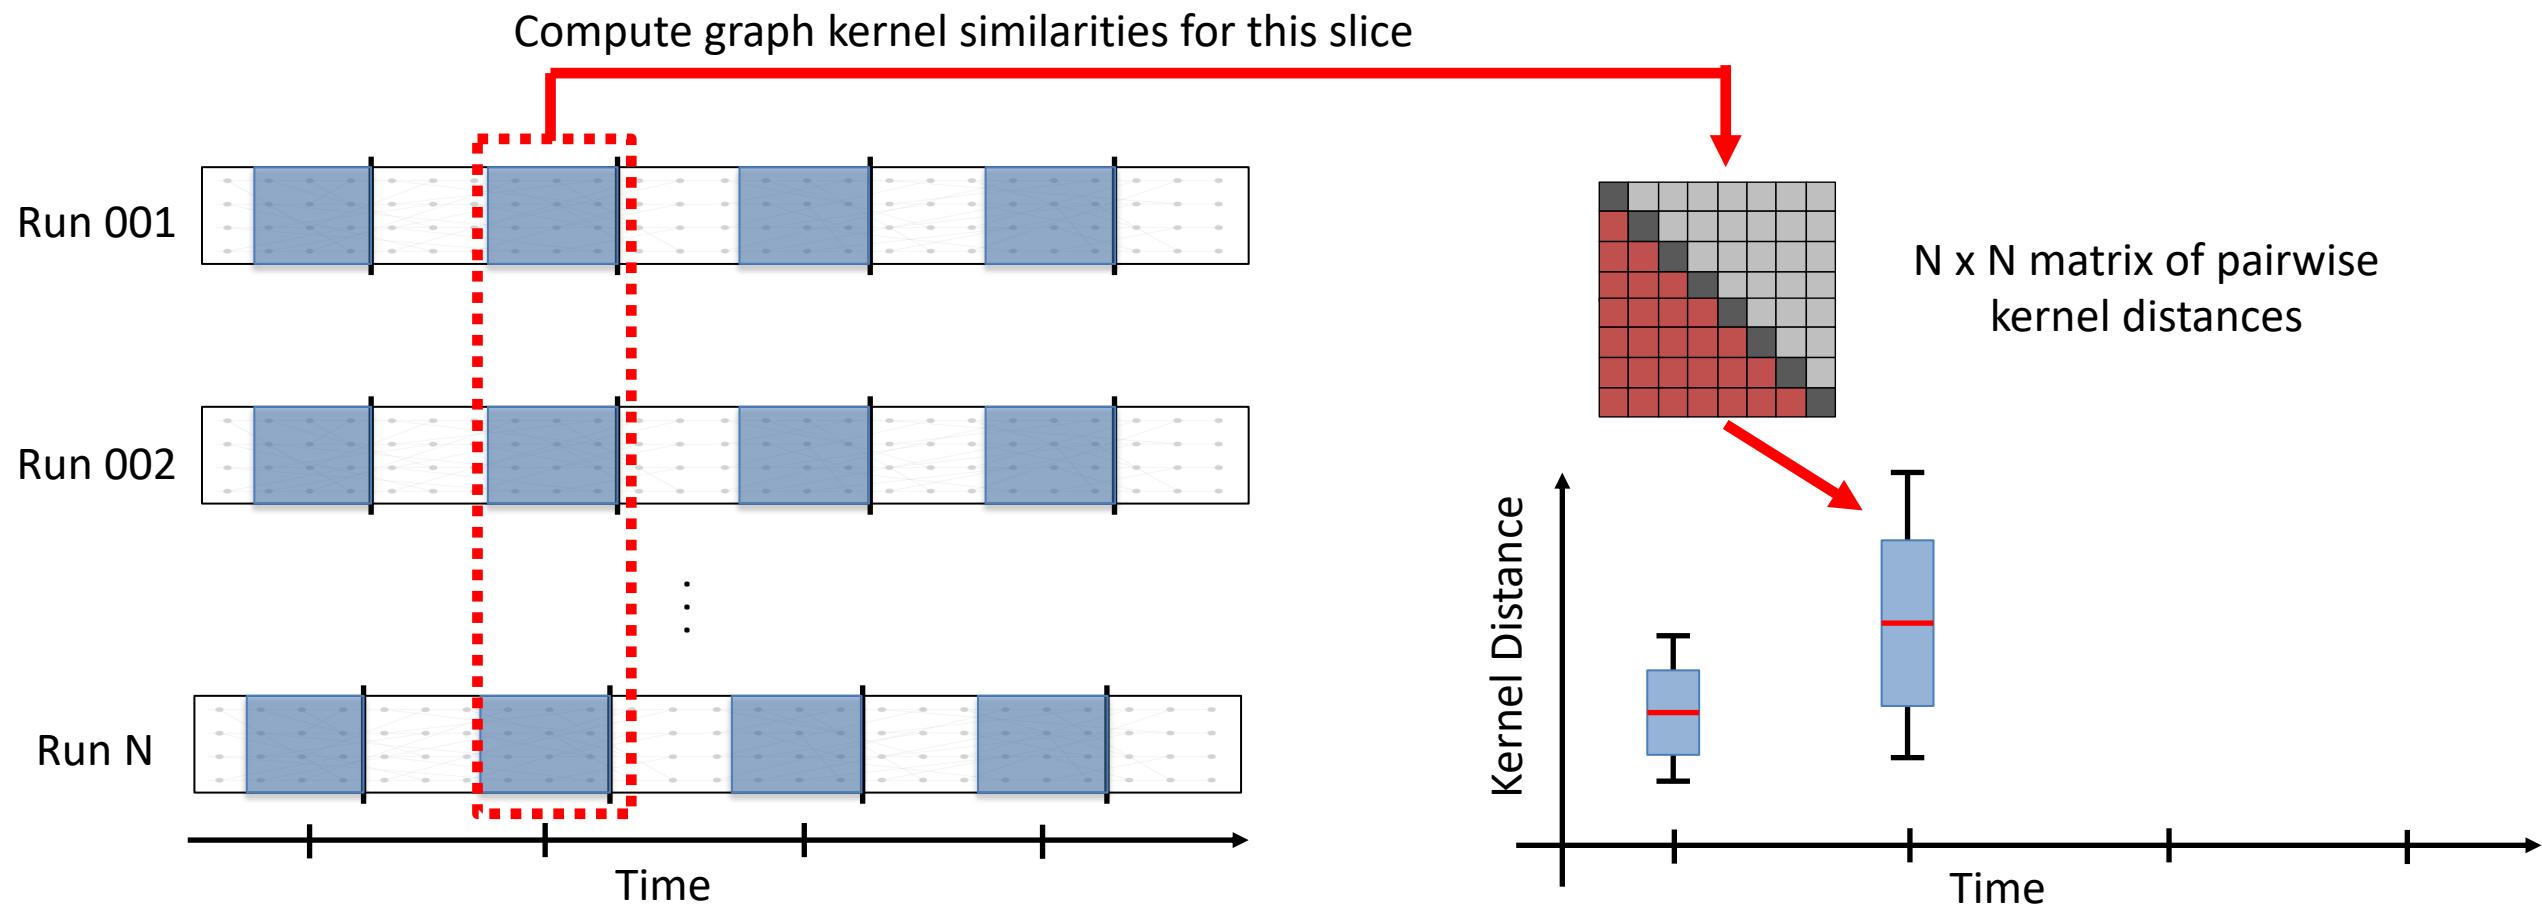
Event Graph Slicing



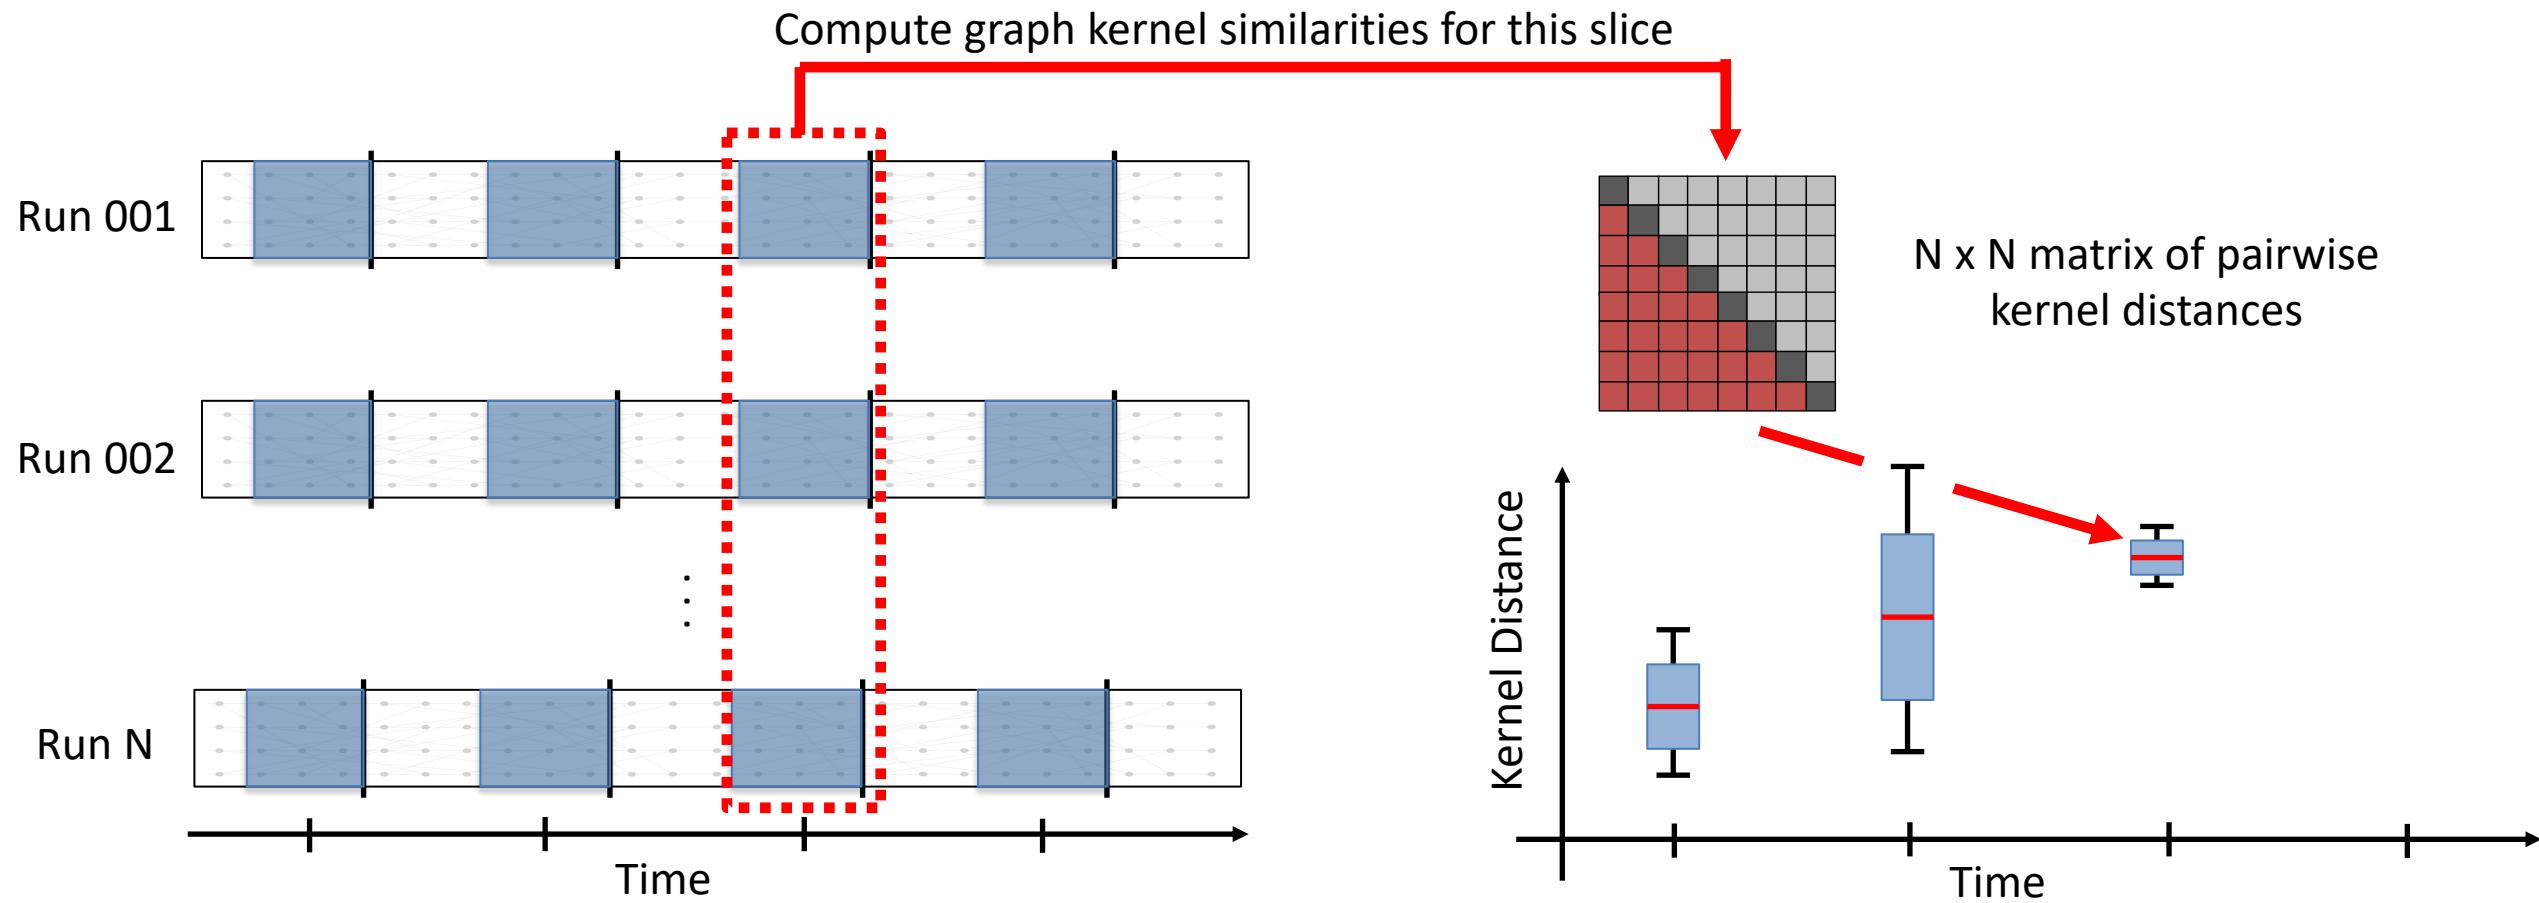
Distributions of Kernel Distances



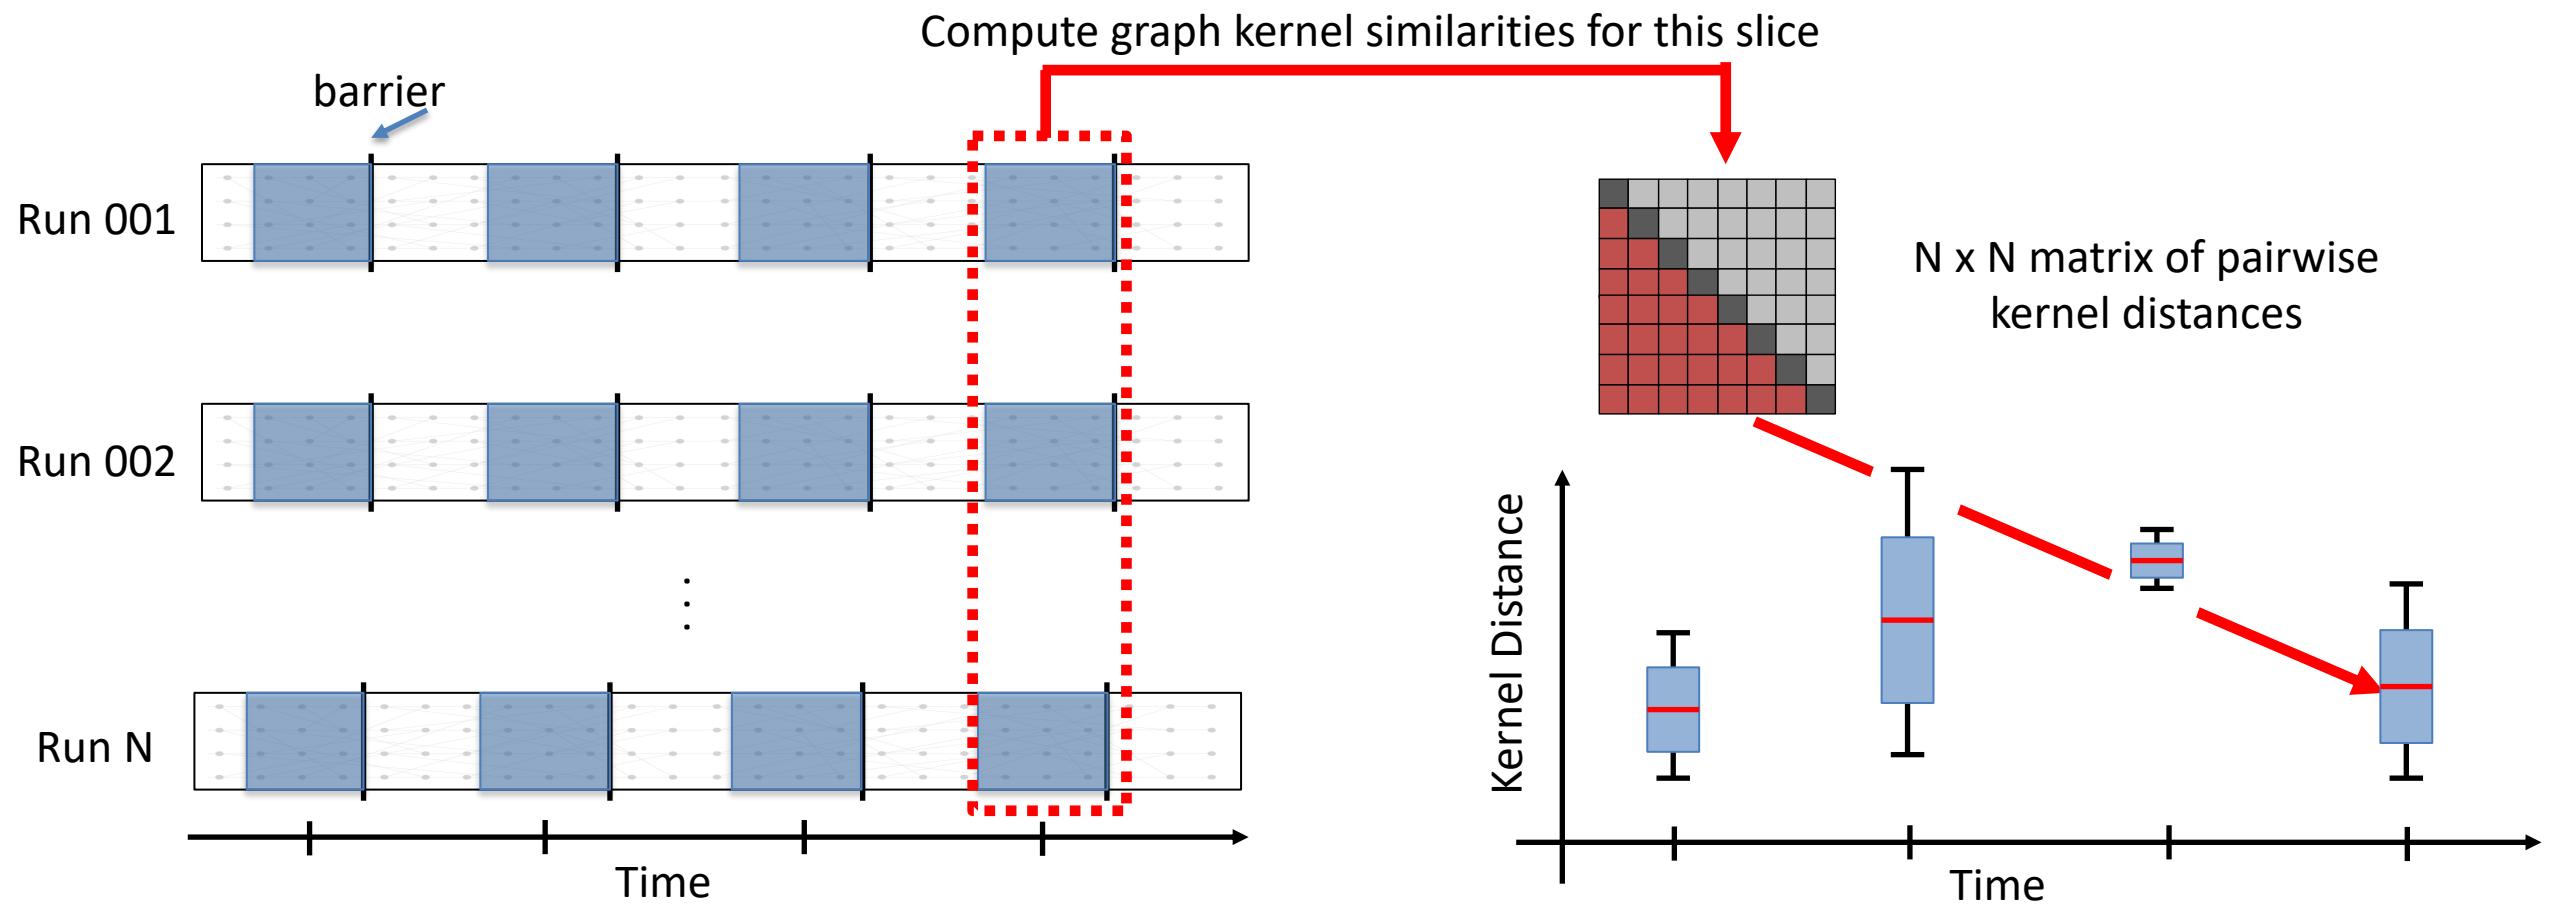
Kernel Distance Trends Over Time



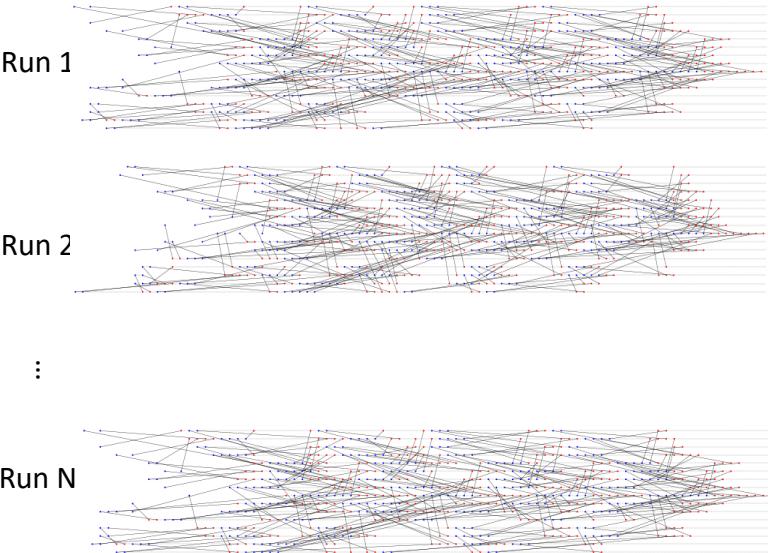
Kernel Distance Trends Over Time



Kernel Distance Trends Over Time

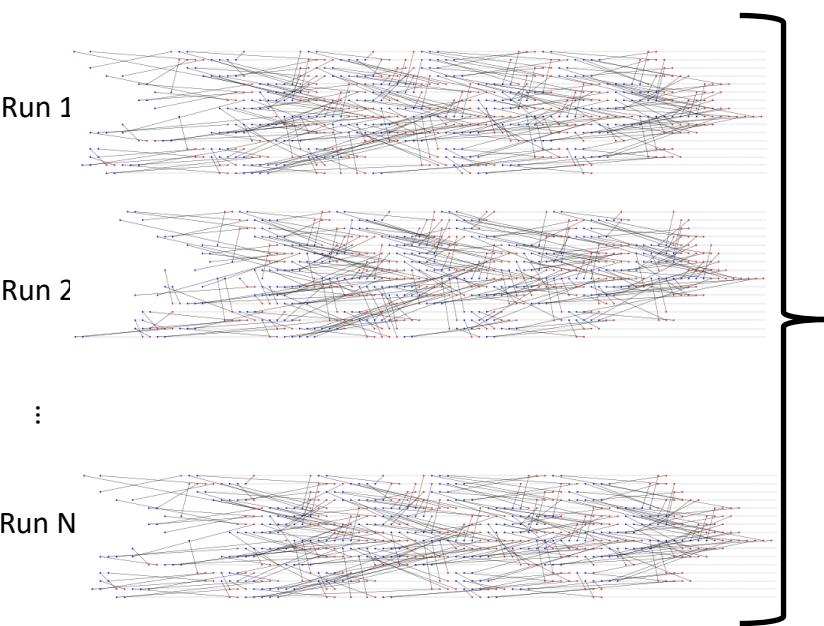


Linking Observed Non-Determinism to Root Causes

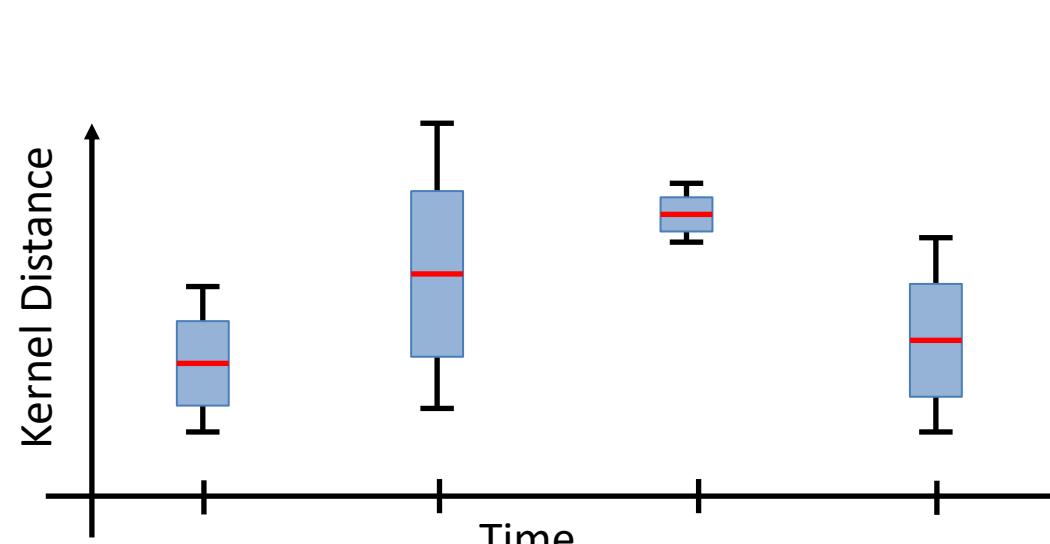


A set of N event graphs
modeling N runs of a
non-deterministic application

Linking Observed Non-Determinism to Root Causes

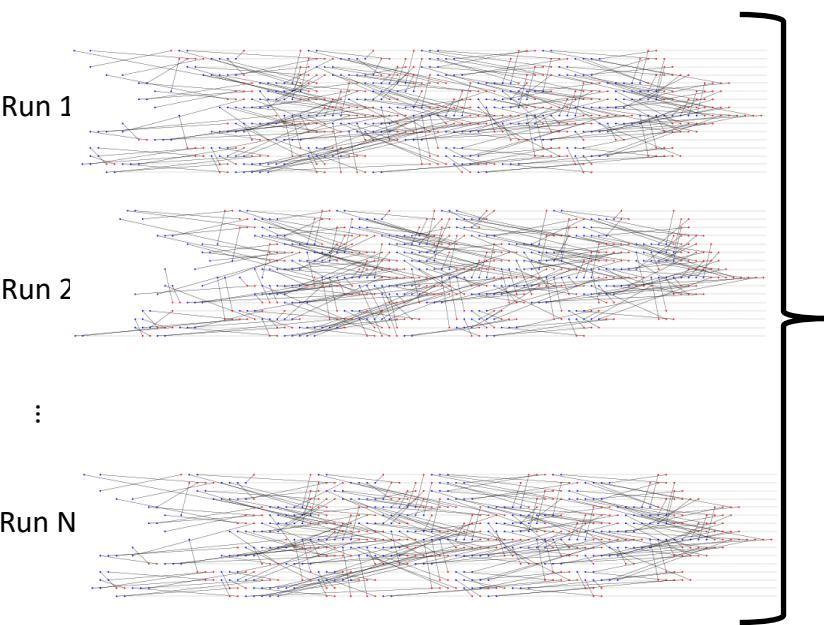


A set of N event graphs
modeling N runs of a
non-deterministic application

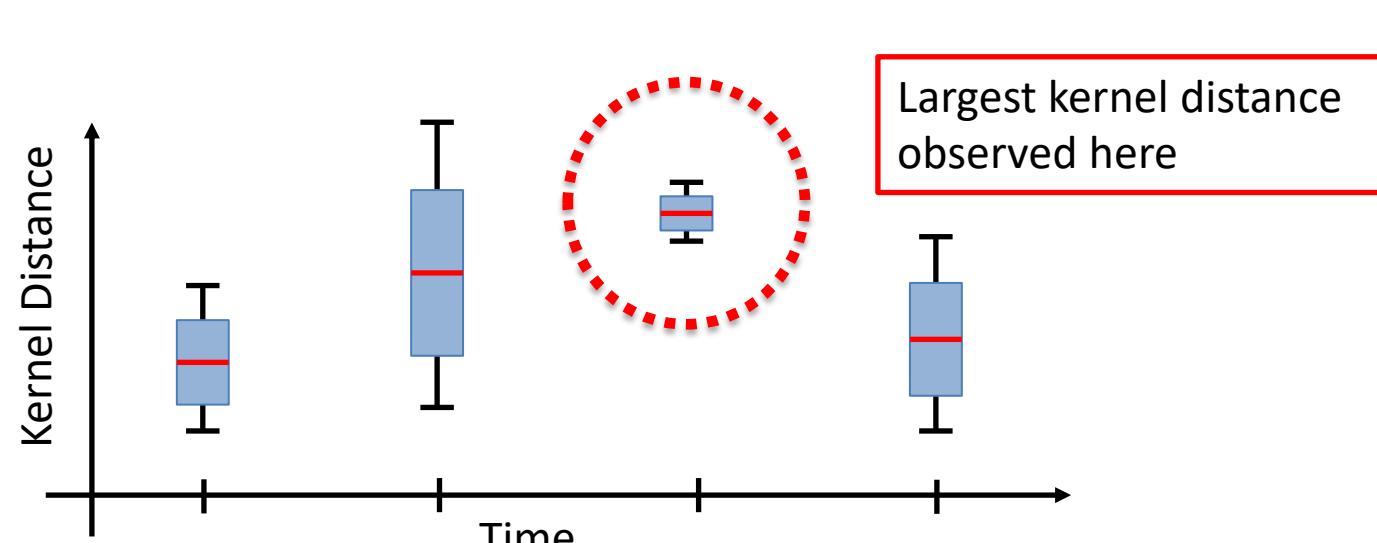


A time-series of distributions of graph kernel distances
between corresponding subgraphs
from pairs of runs

Linking Observed Non-Determinism to Root Causes

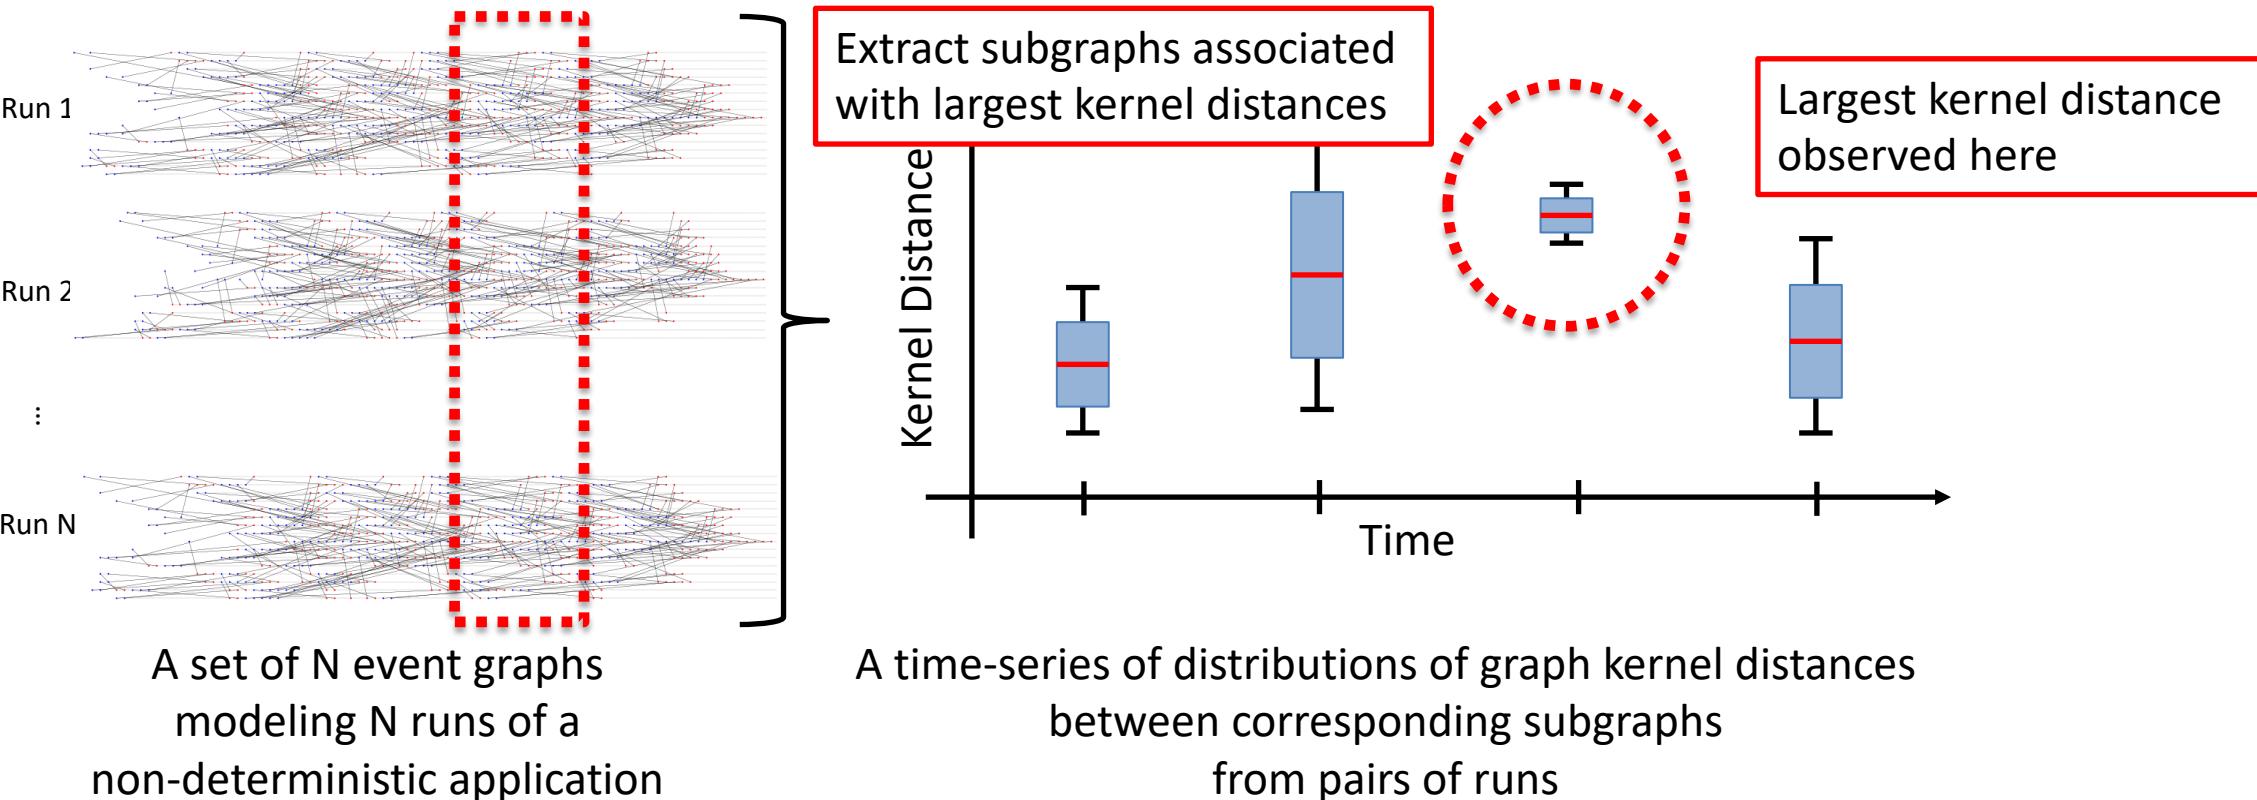


A set of N event graphs
modeling N runs of a
non-deterministic application

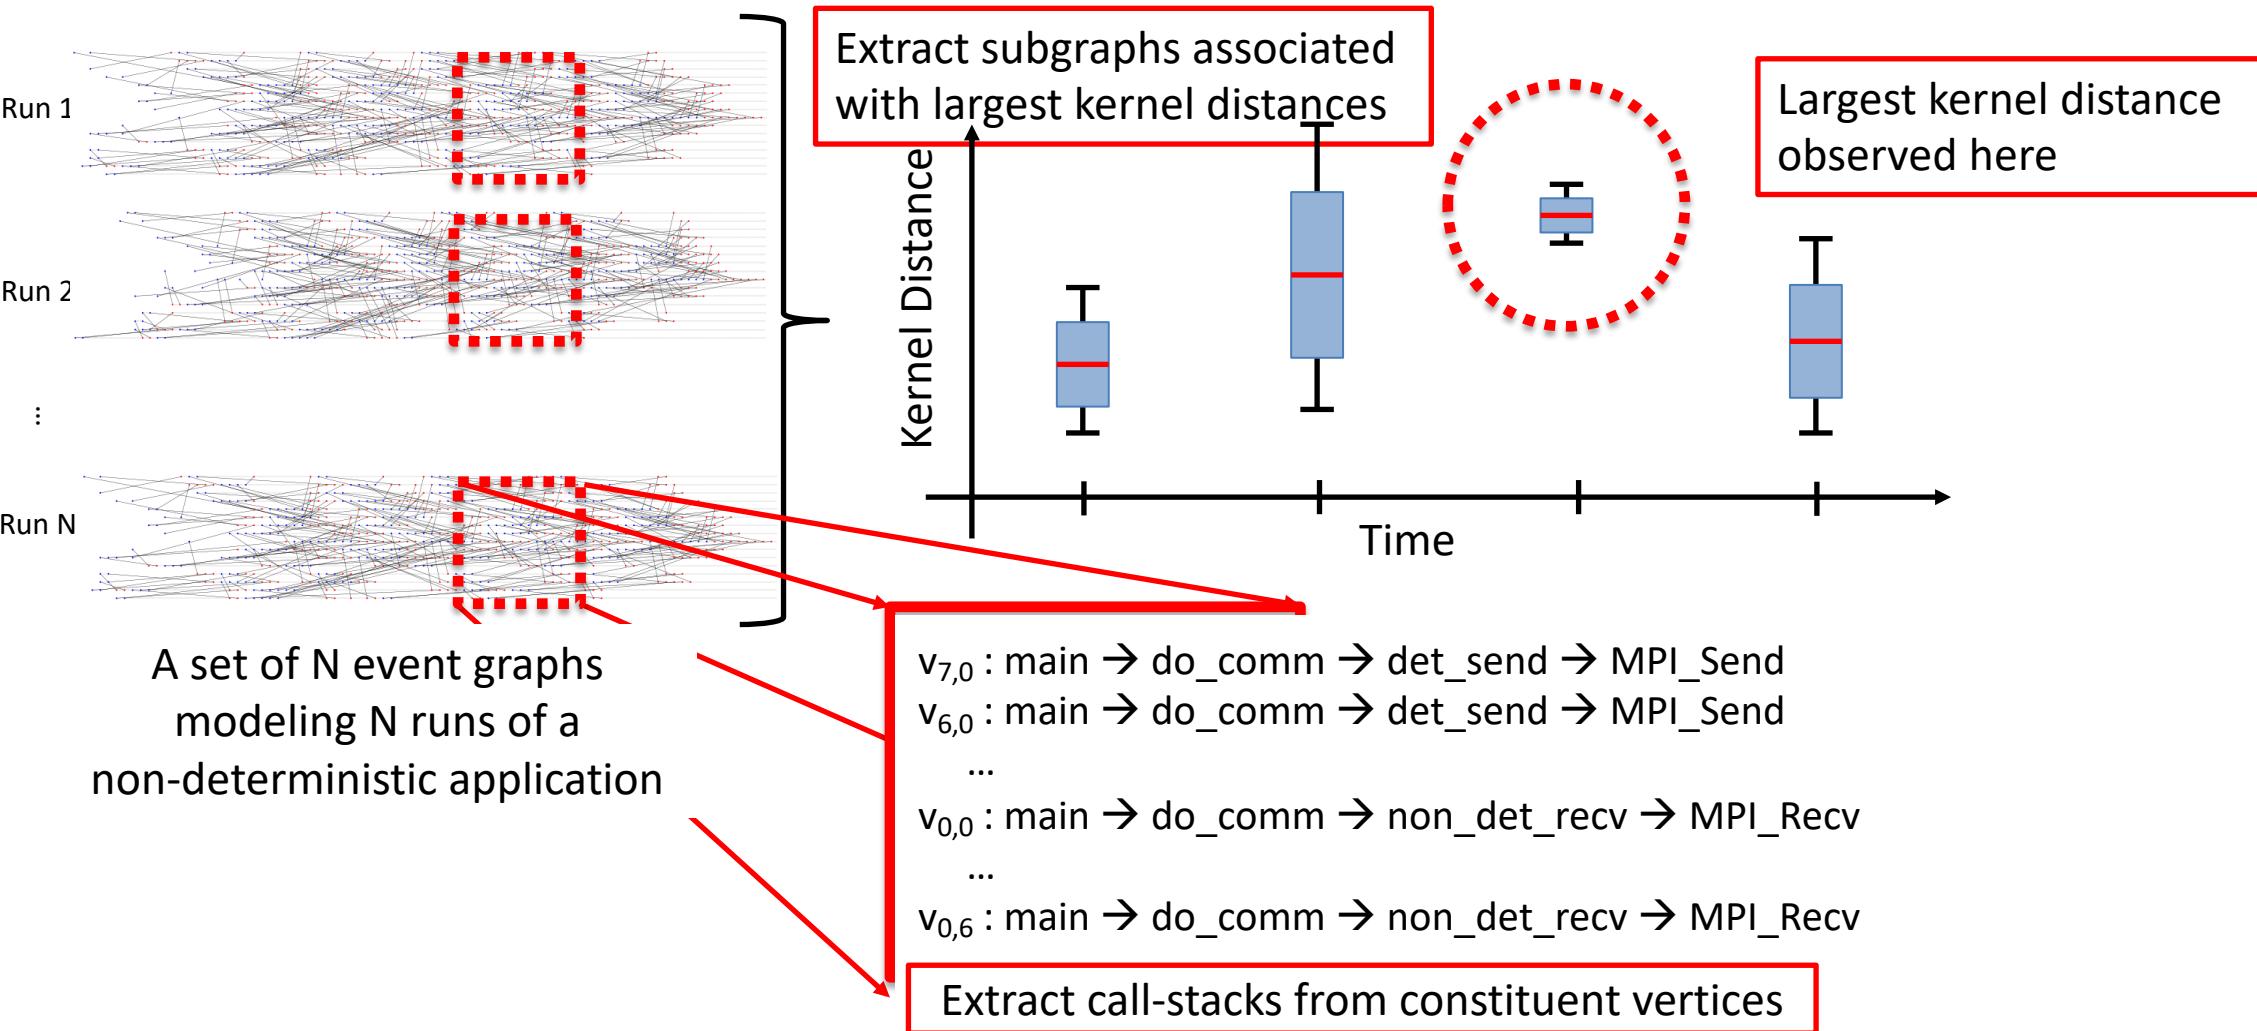


A time-series of distributions of graph kernel distances
between corresponding subgraphs
from pairs of runs

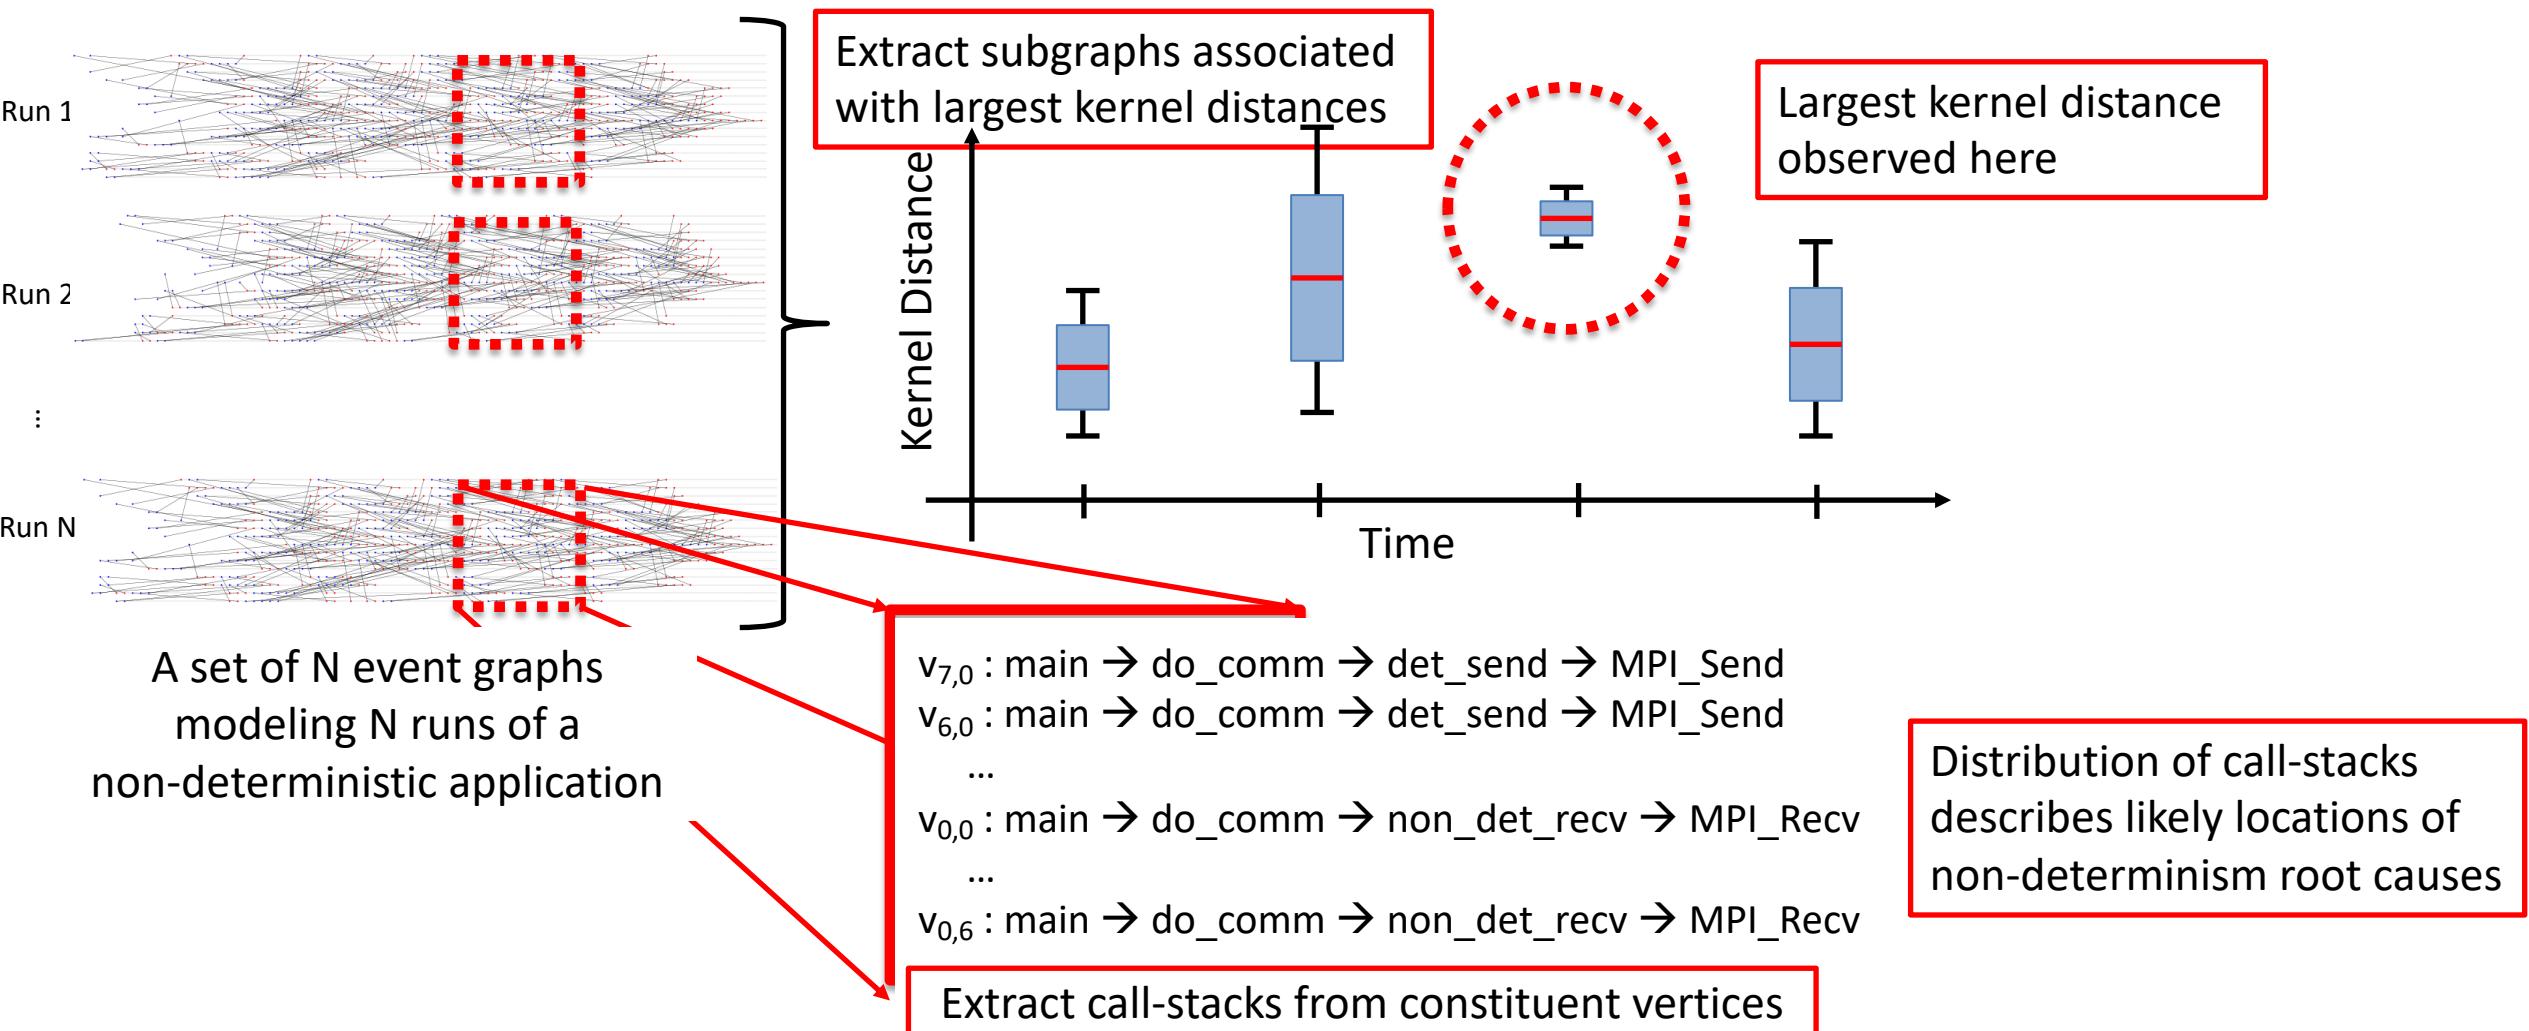
Linking Observed Non-Determinism to Root Causes



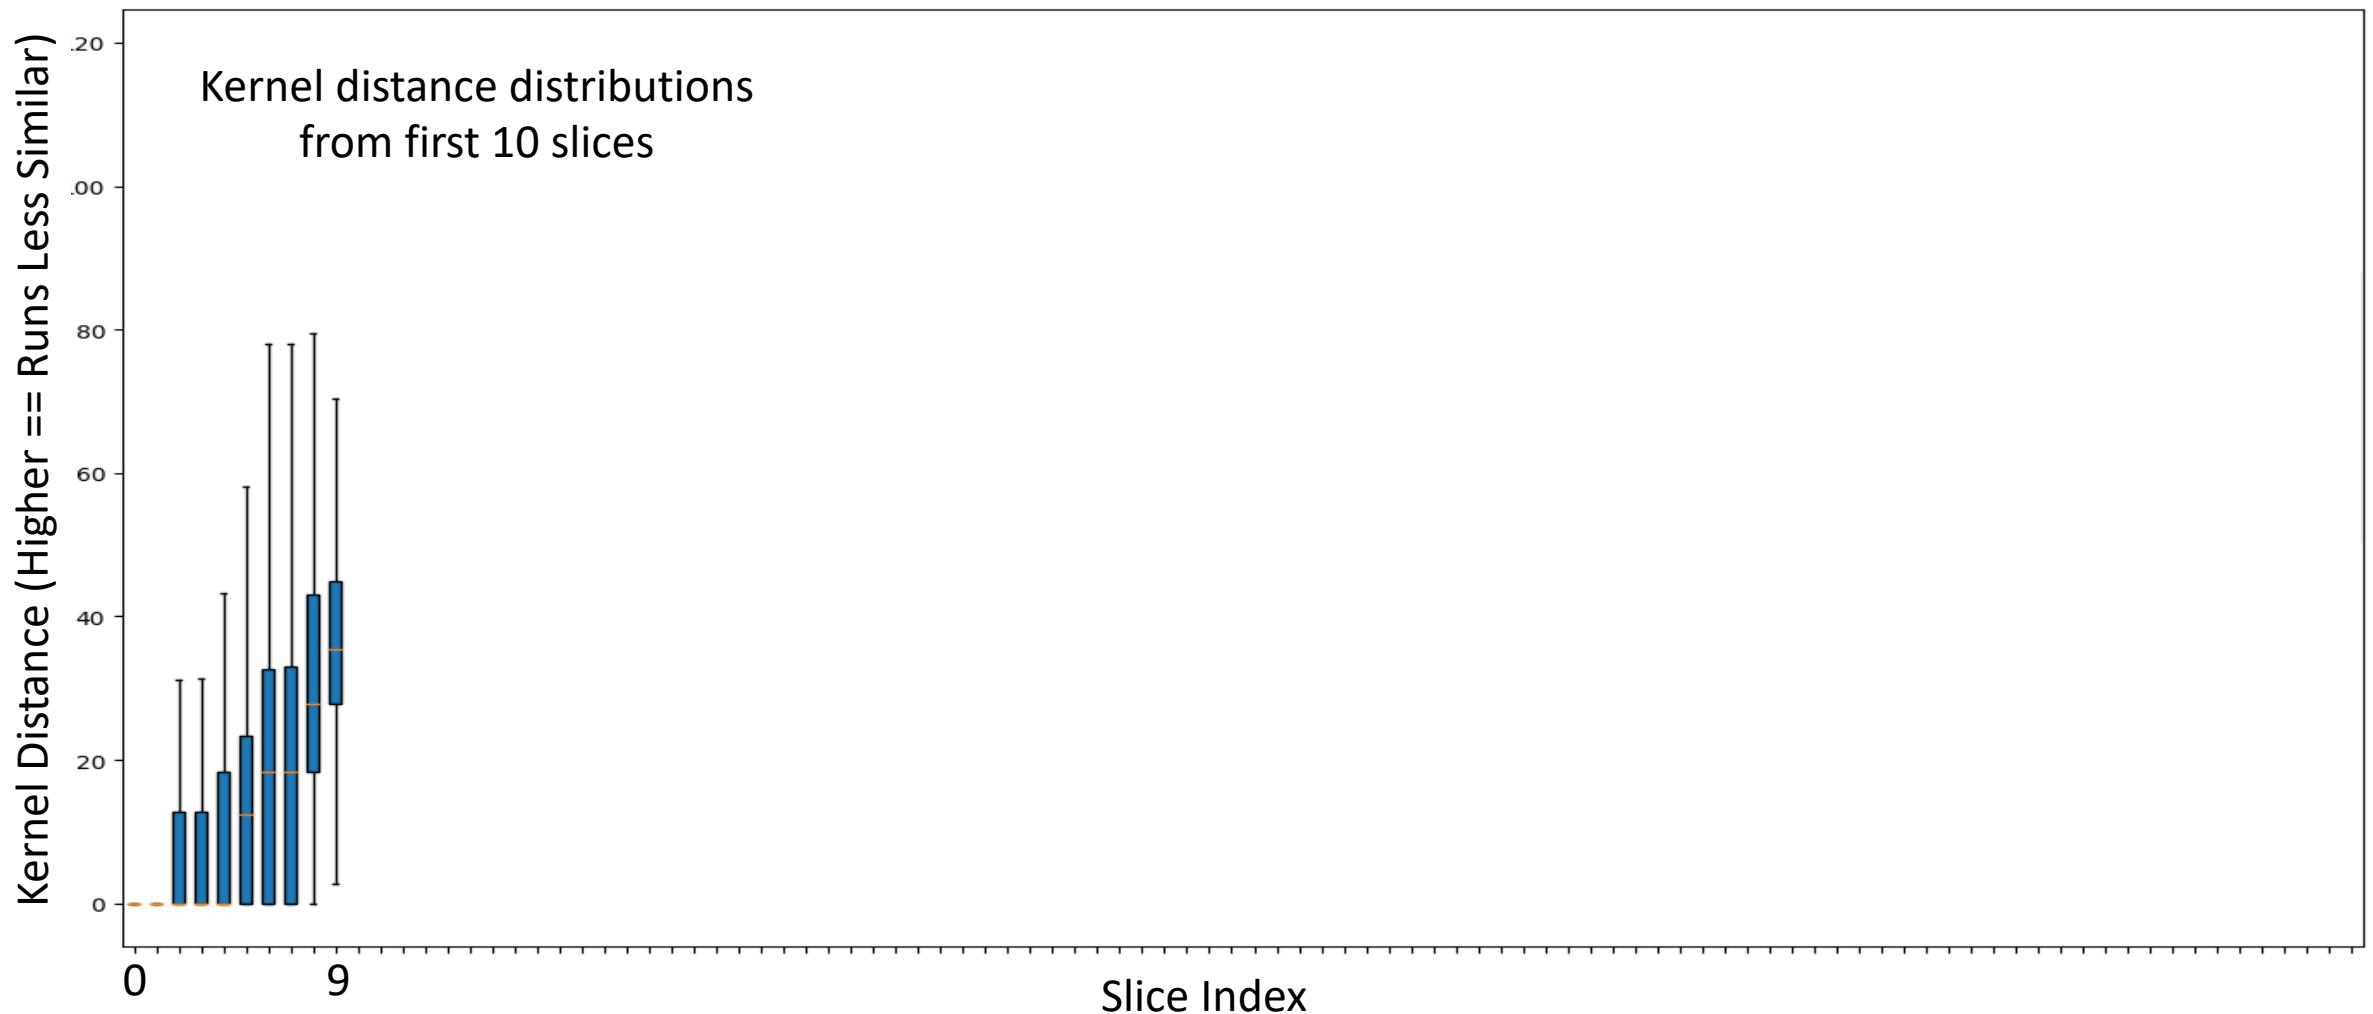
Linking Observed Non-Determinism to Root Causes



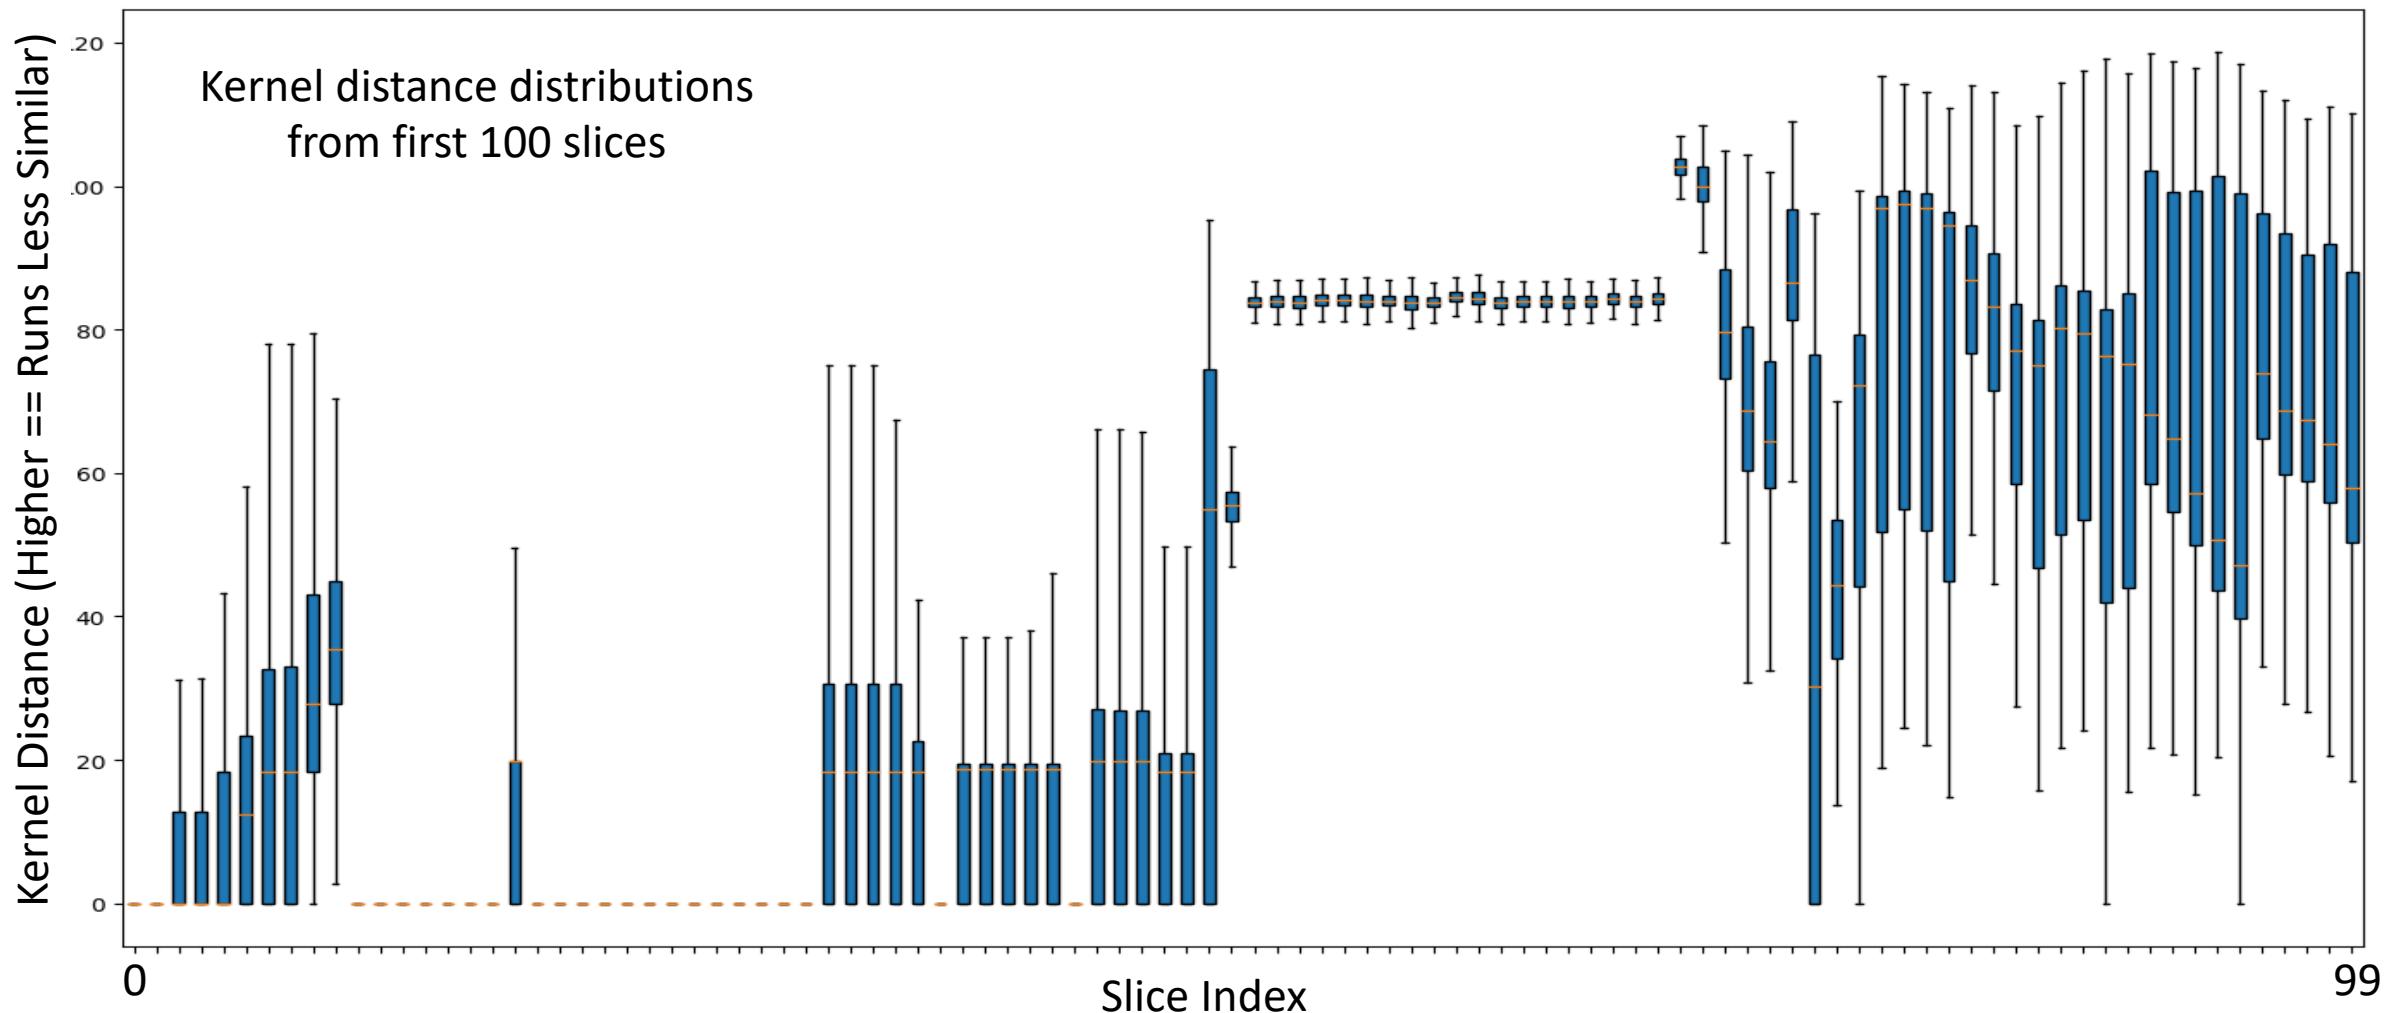
Linking Observed Non-Determinism to Root Causes



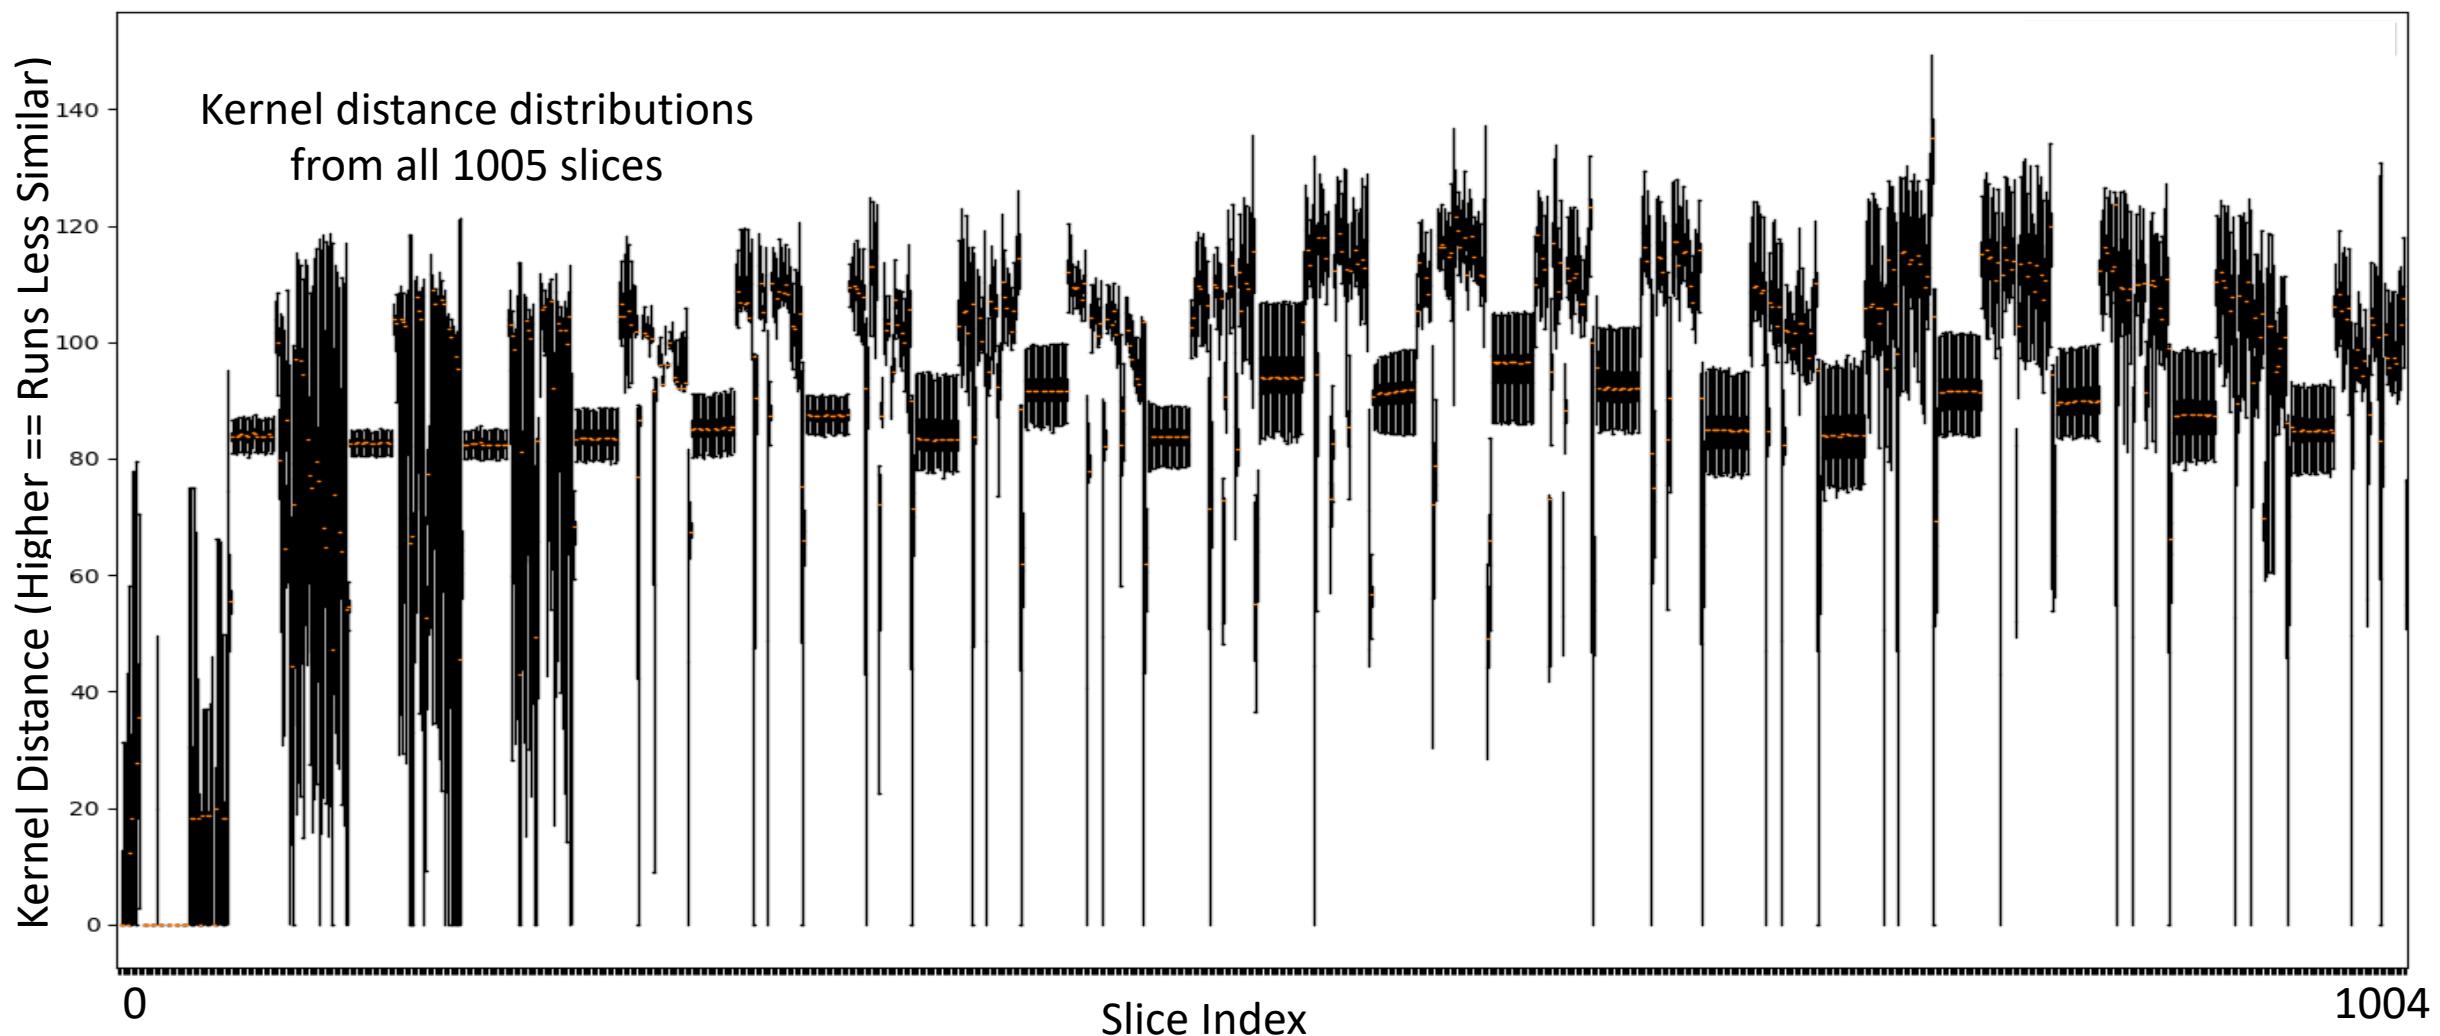
Kernel Distance Time Series for miniAMR



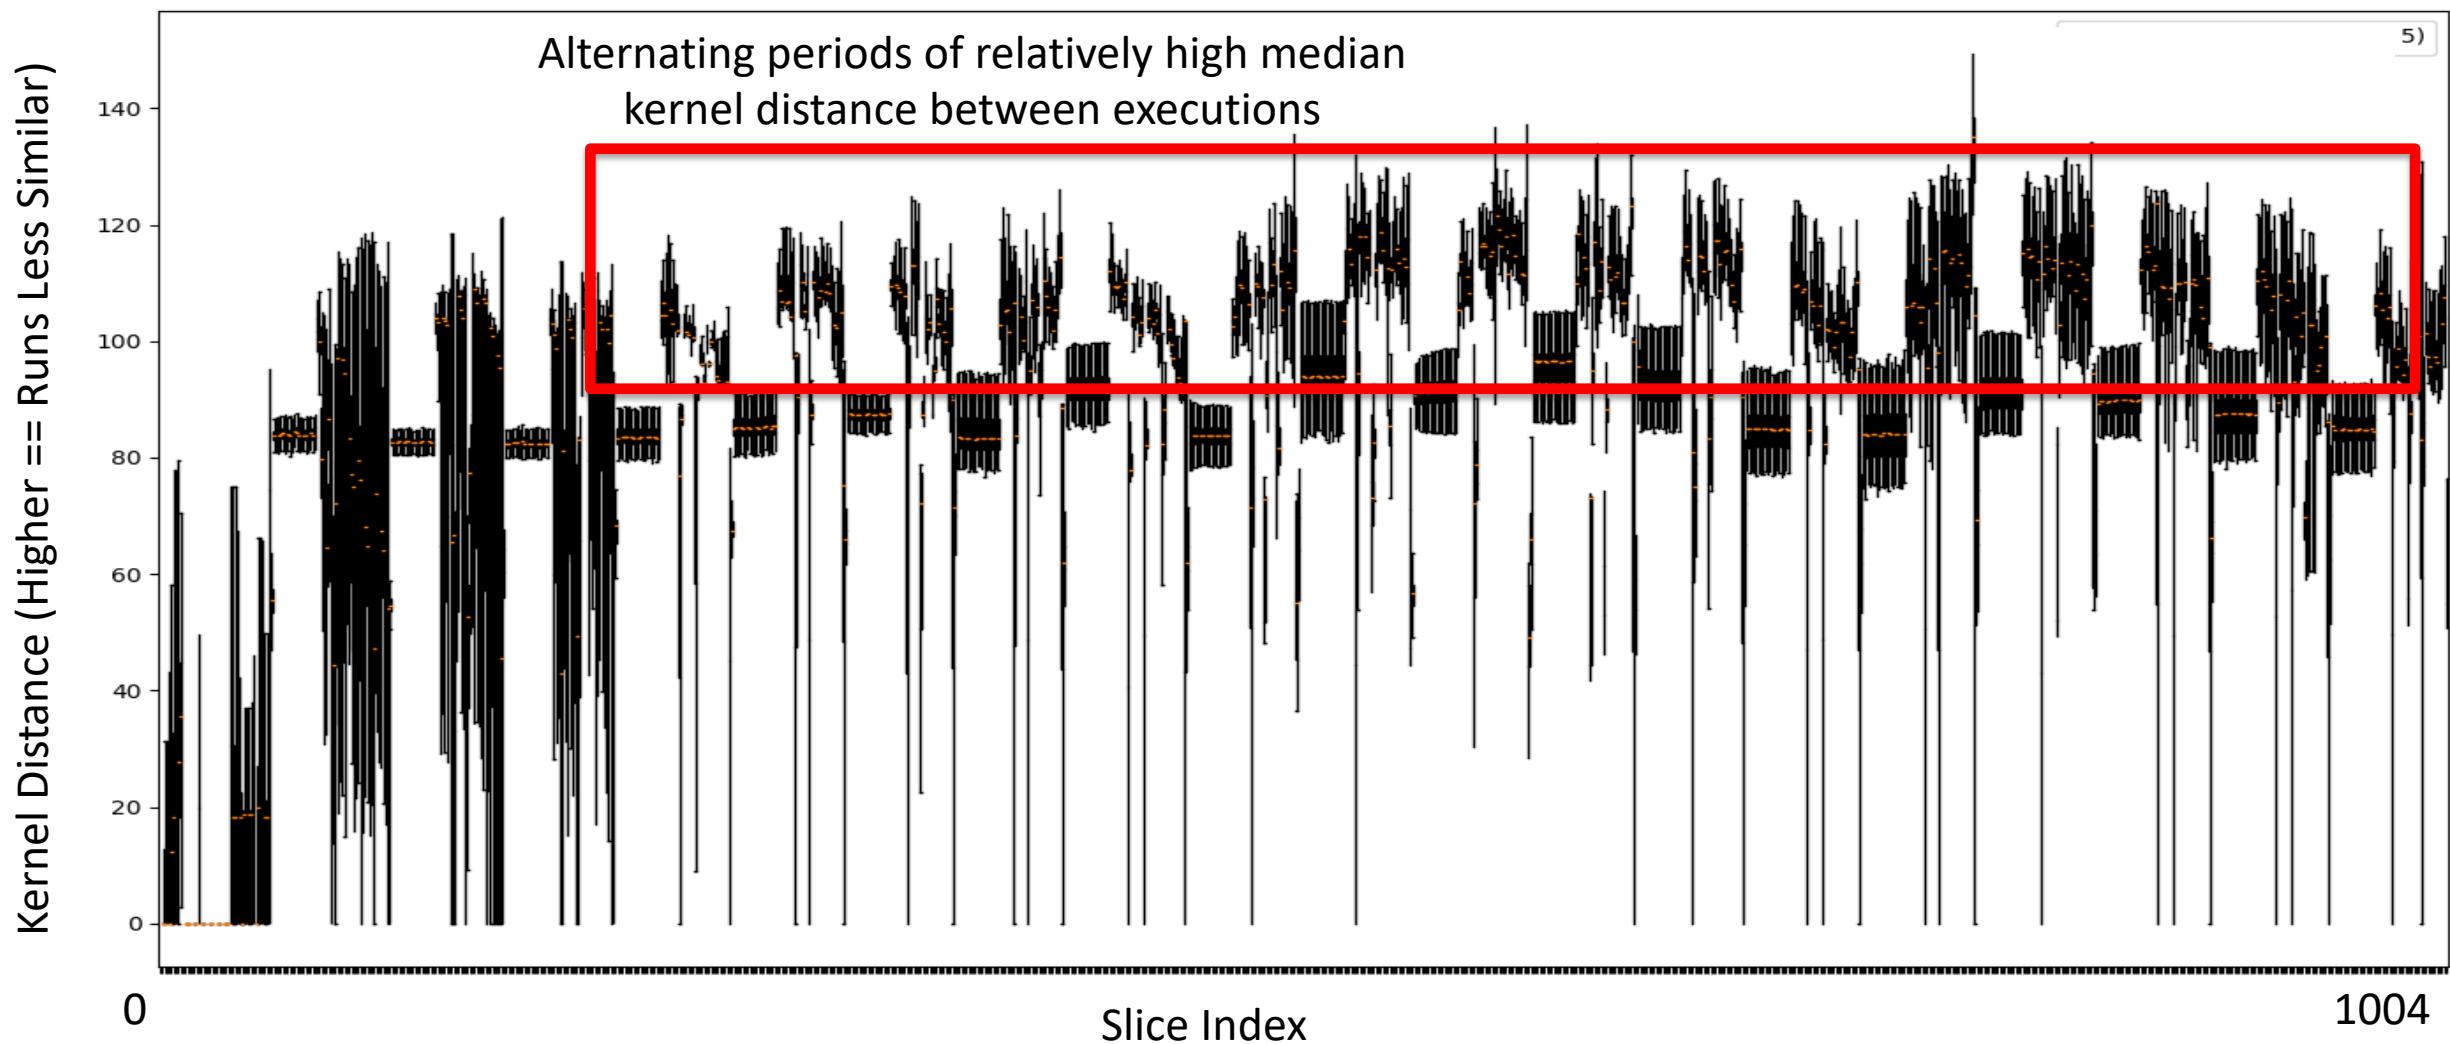
Kernel Distance Time Series for miniAMR



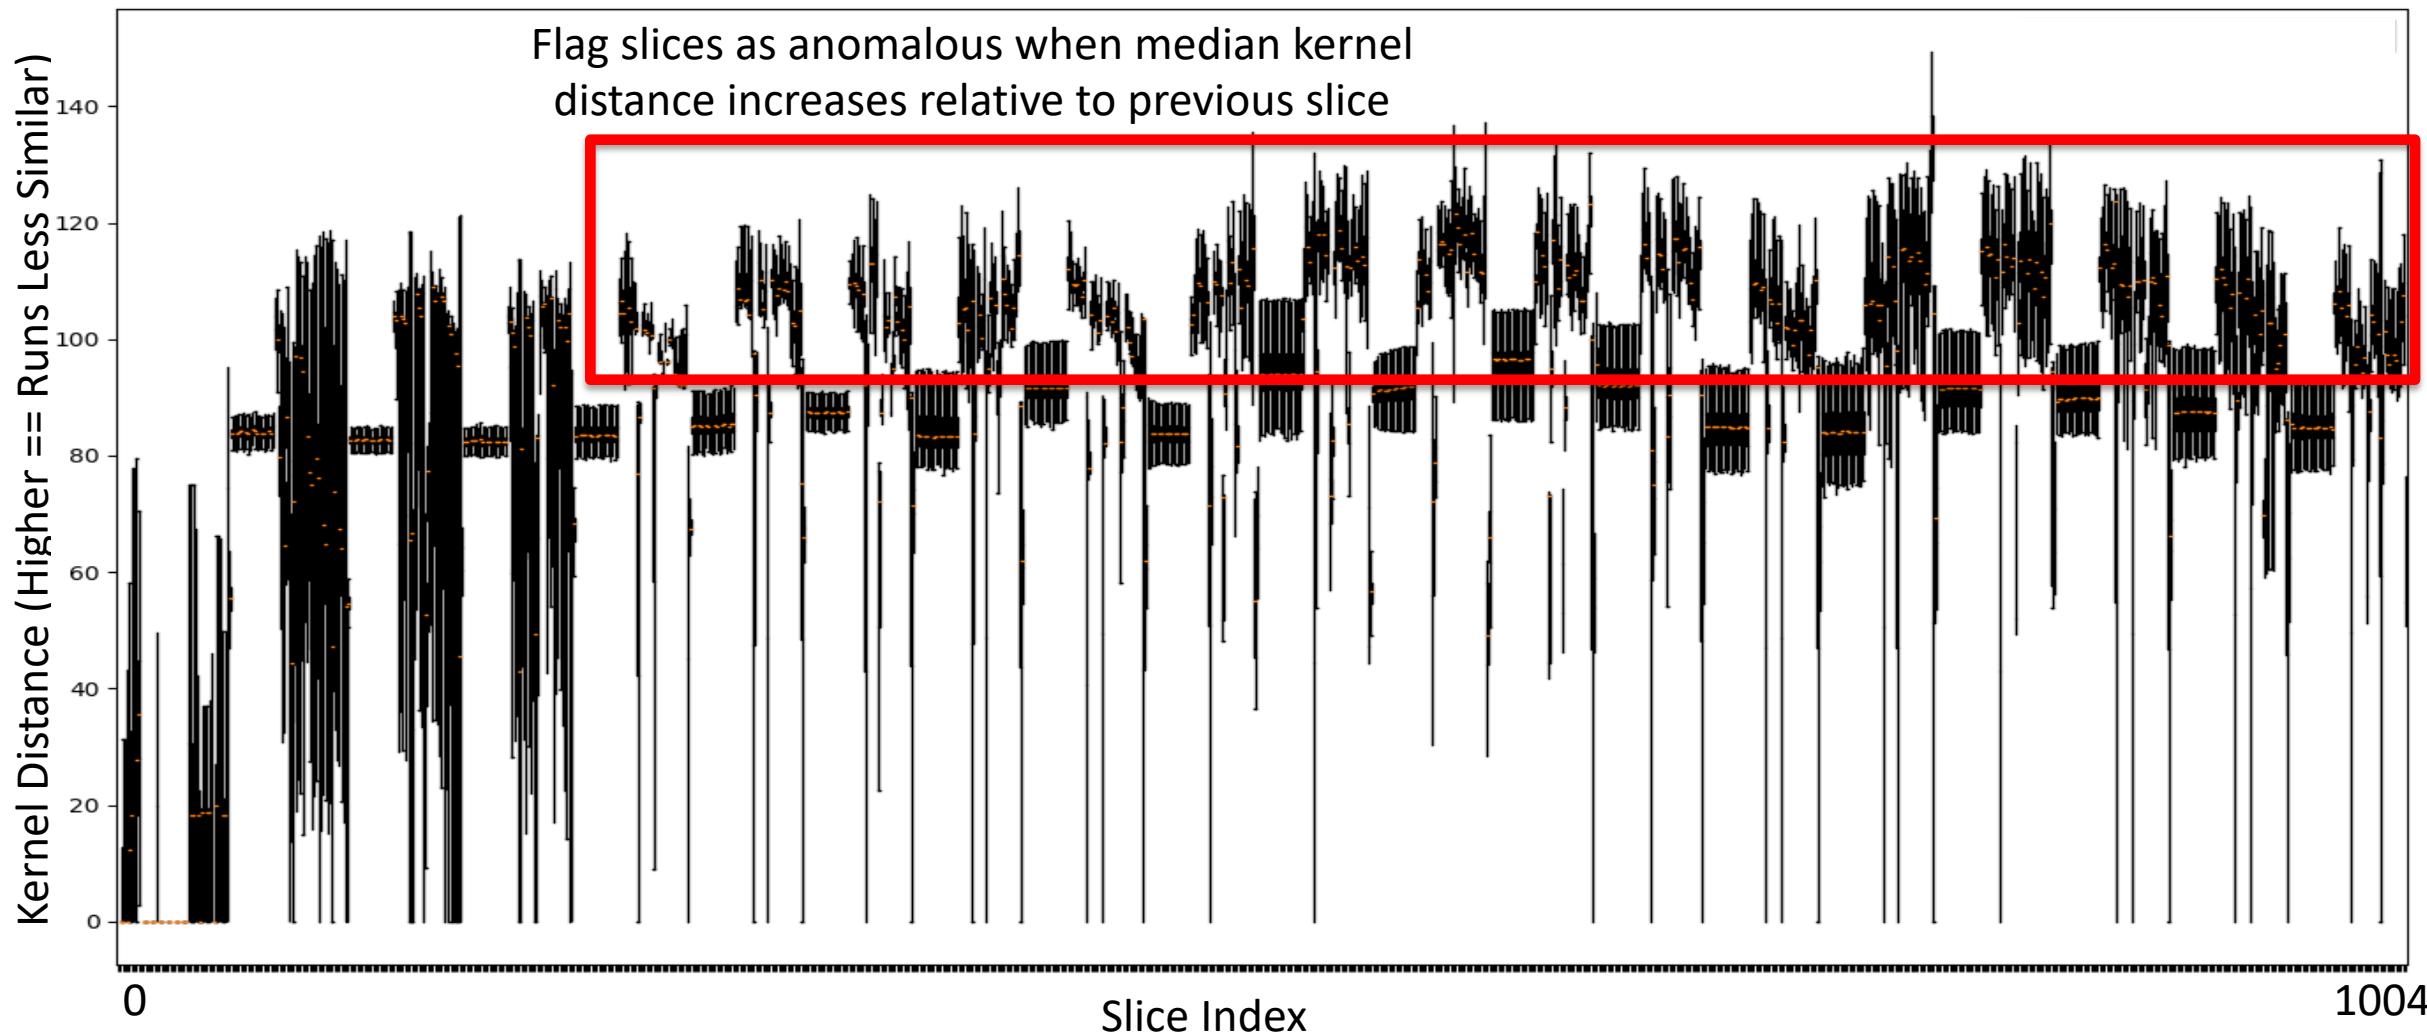
Kernel Distance Time Series for miniAMR



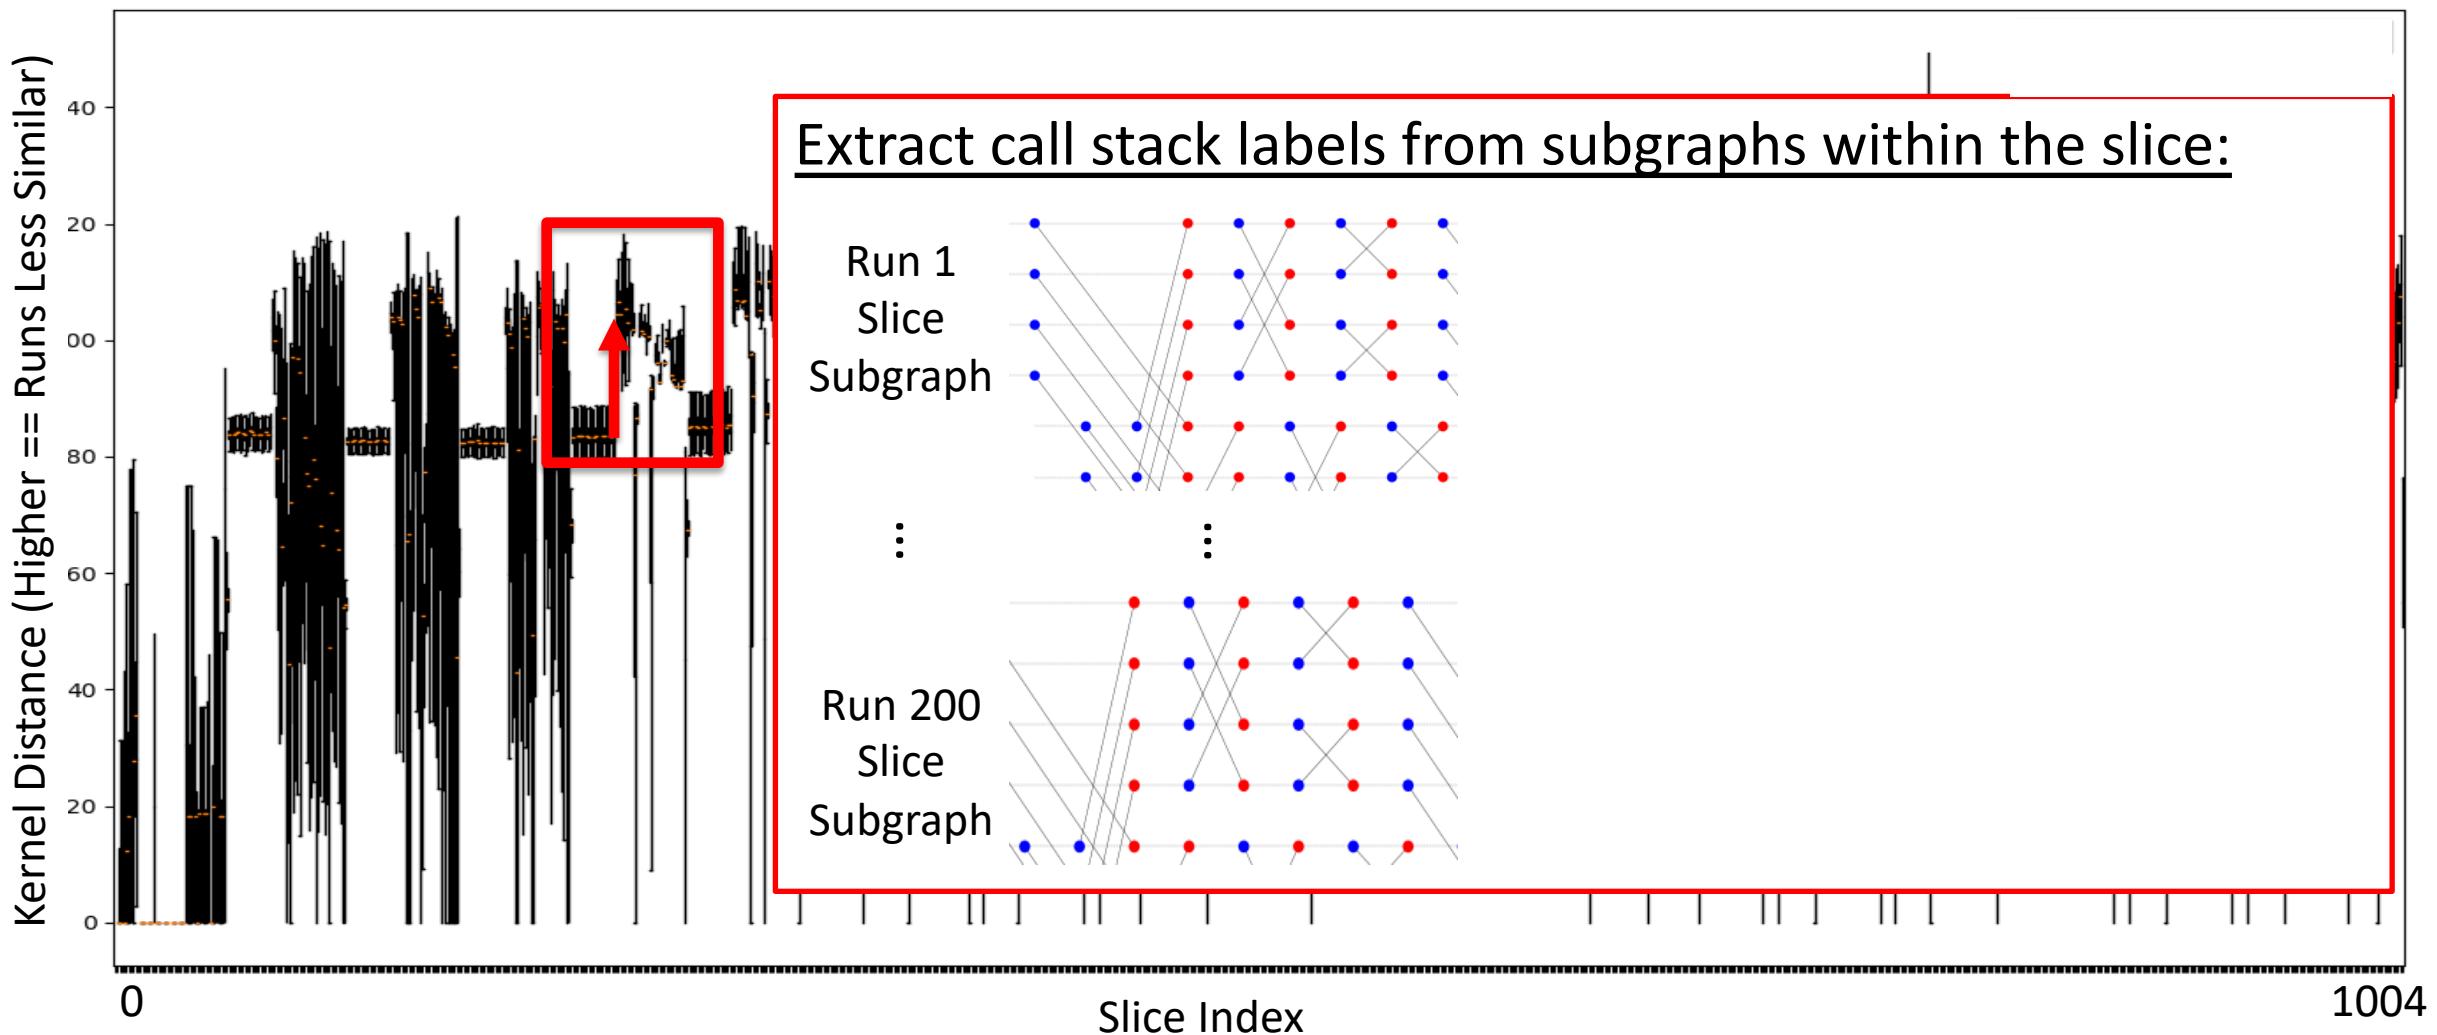
Identifying Anomalous Slices



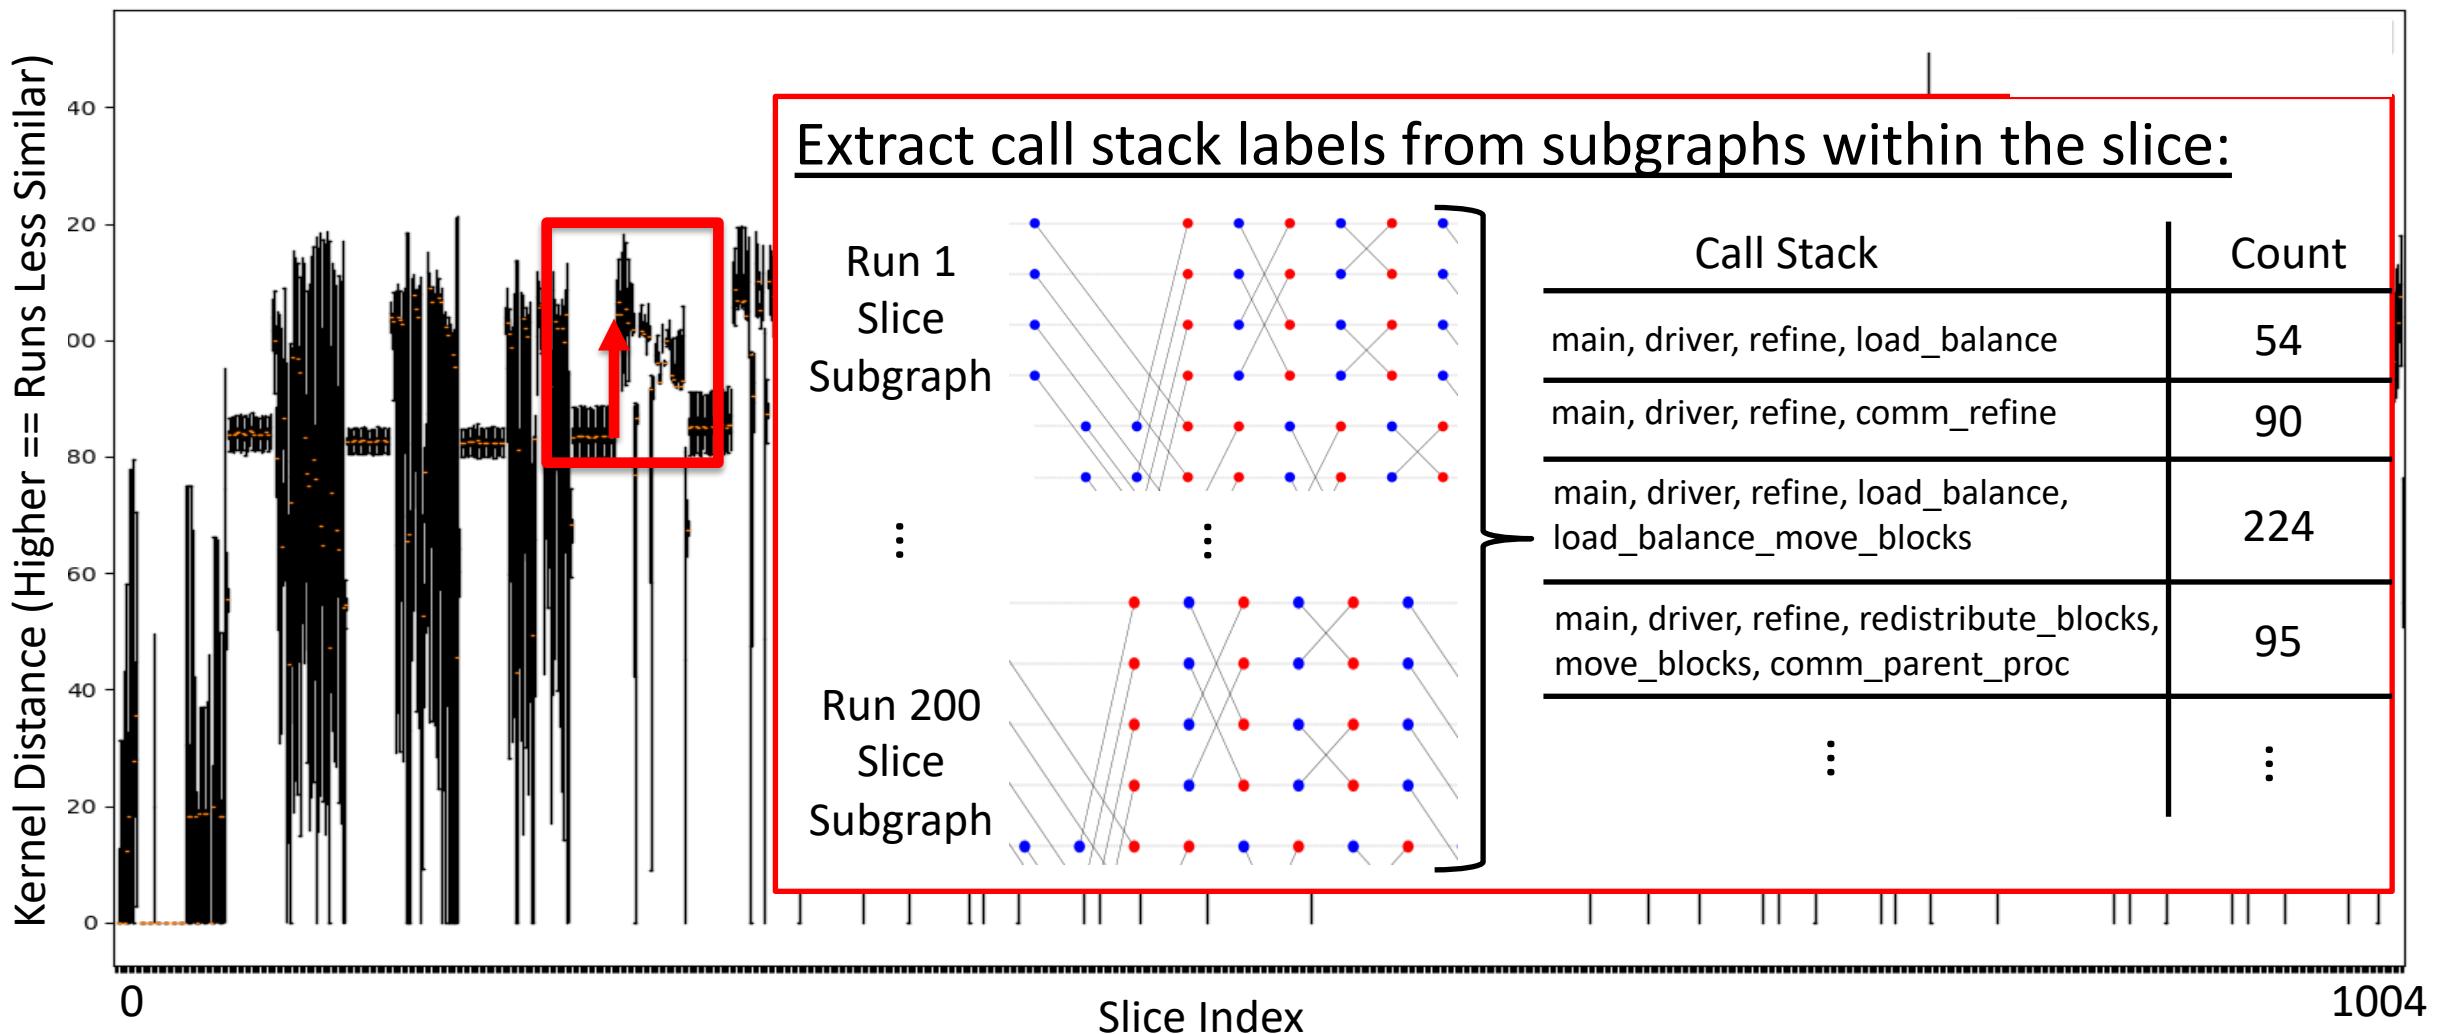
Identifying Anomalous Slices



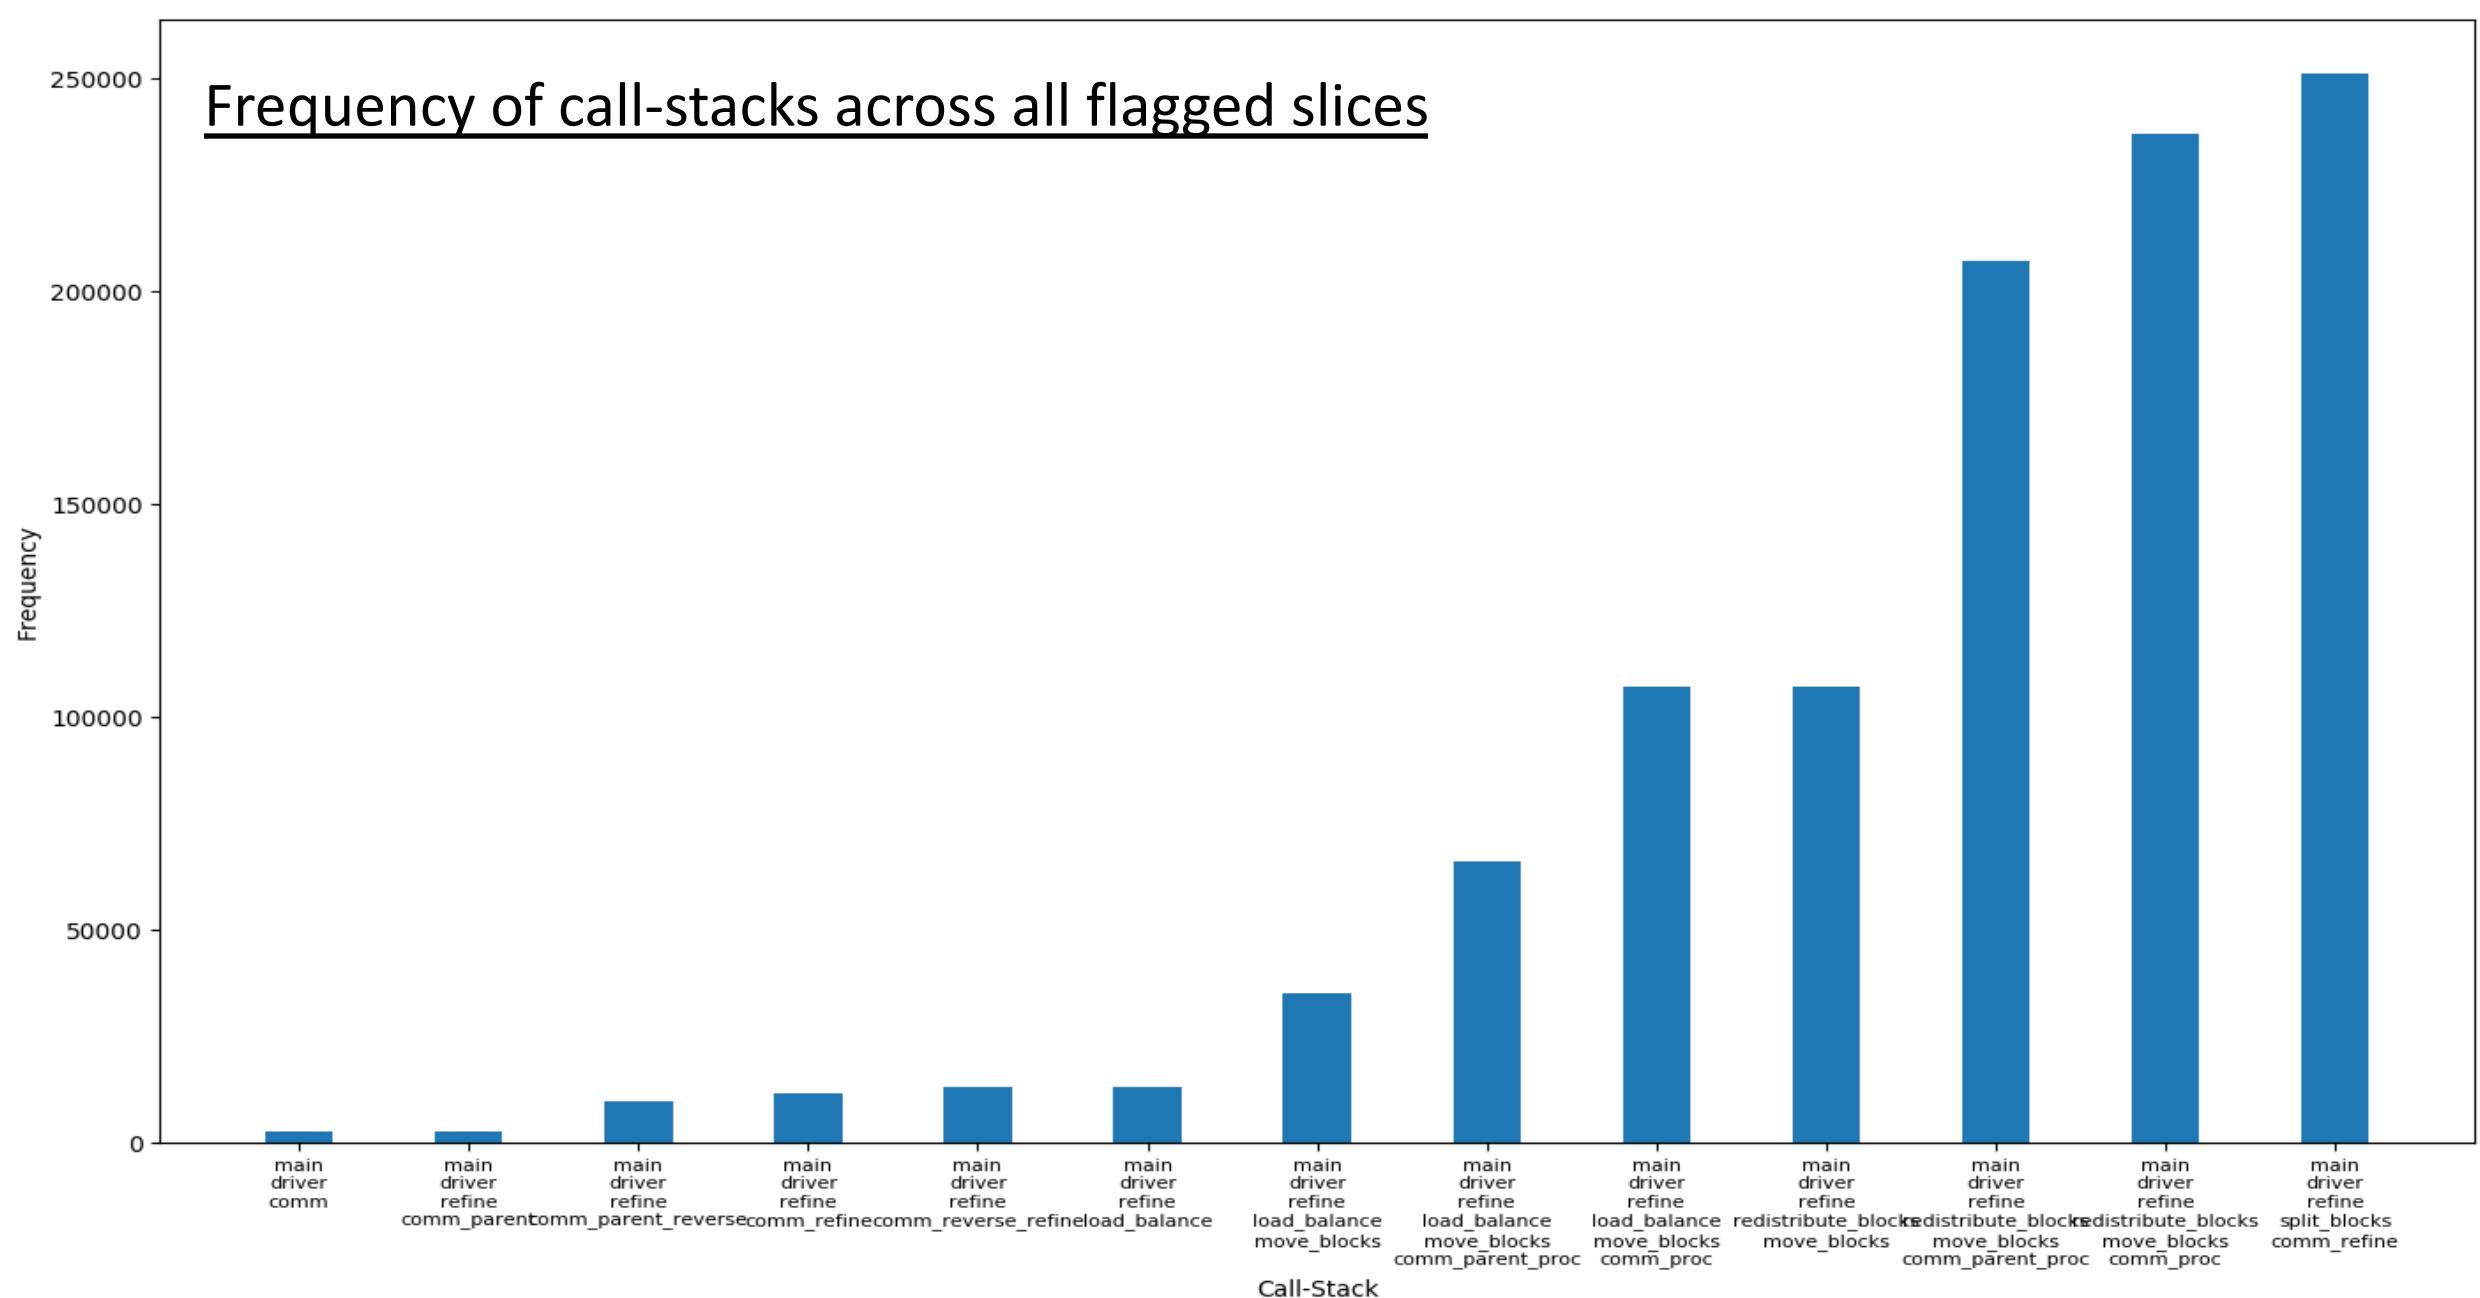
Linking to Potential Root Causes

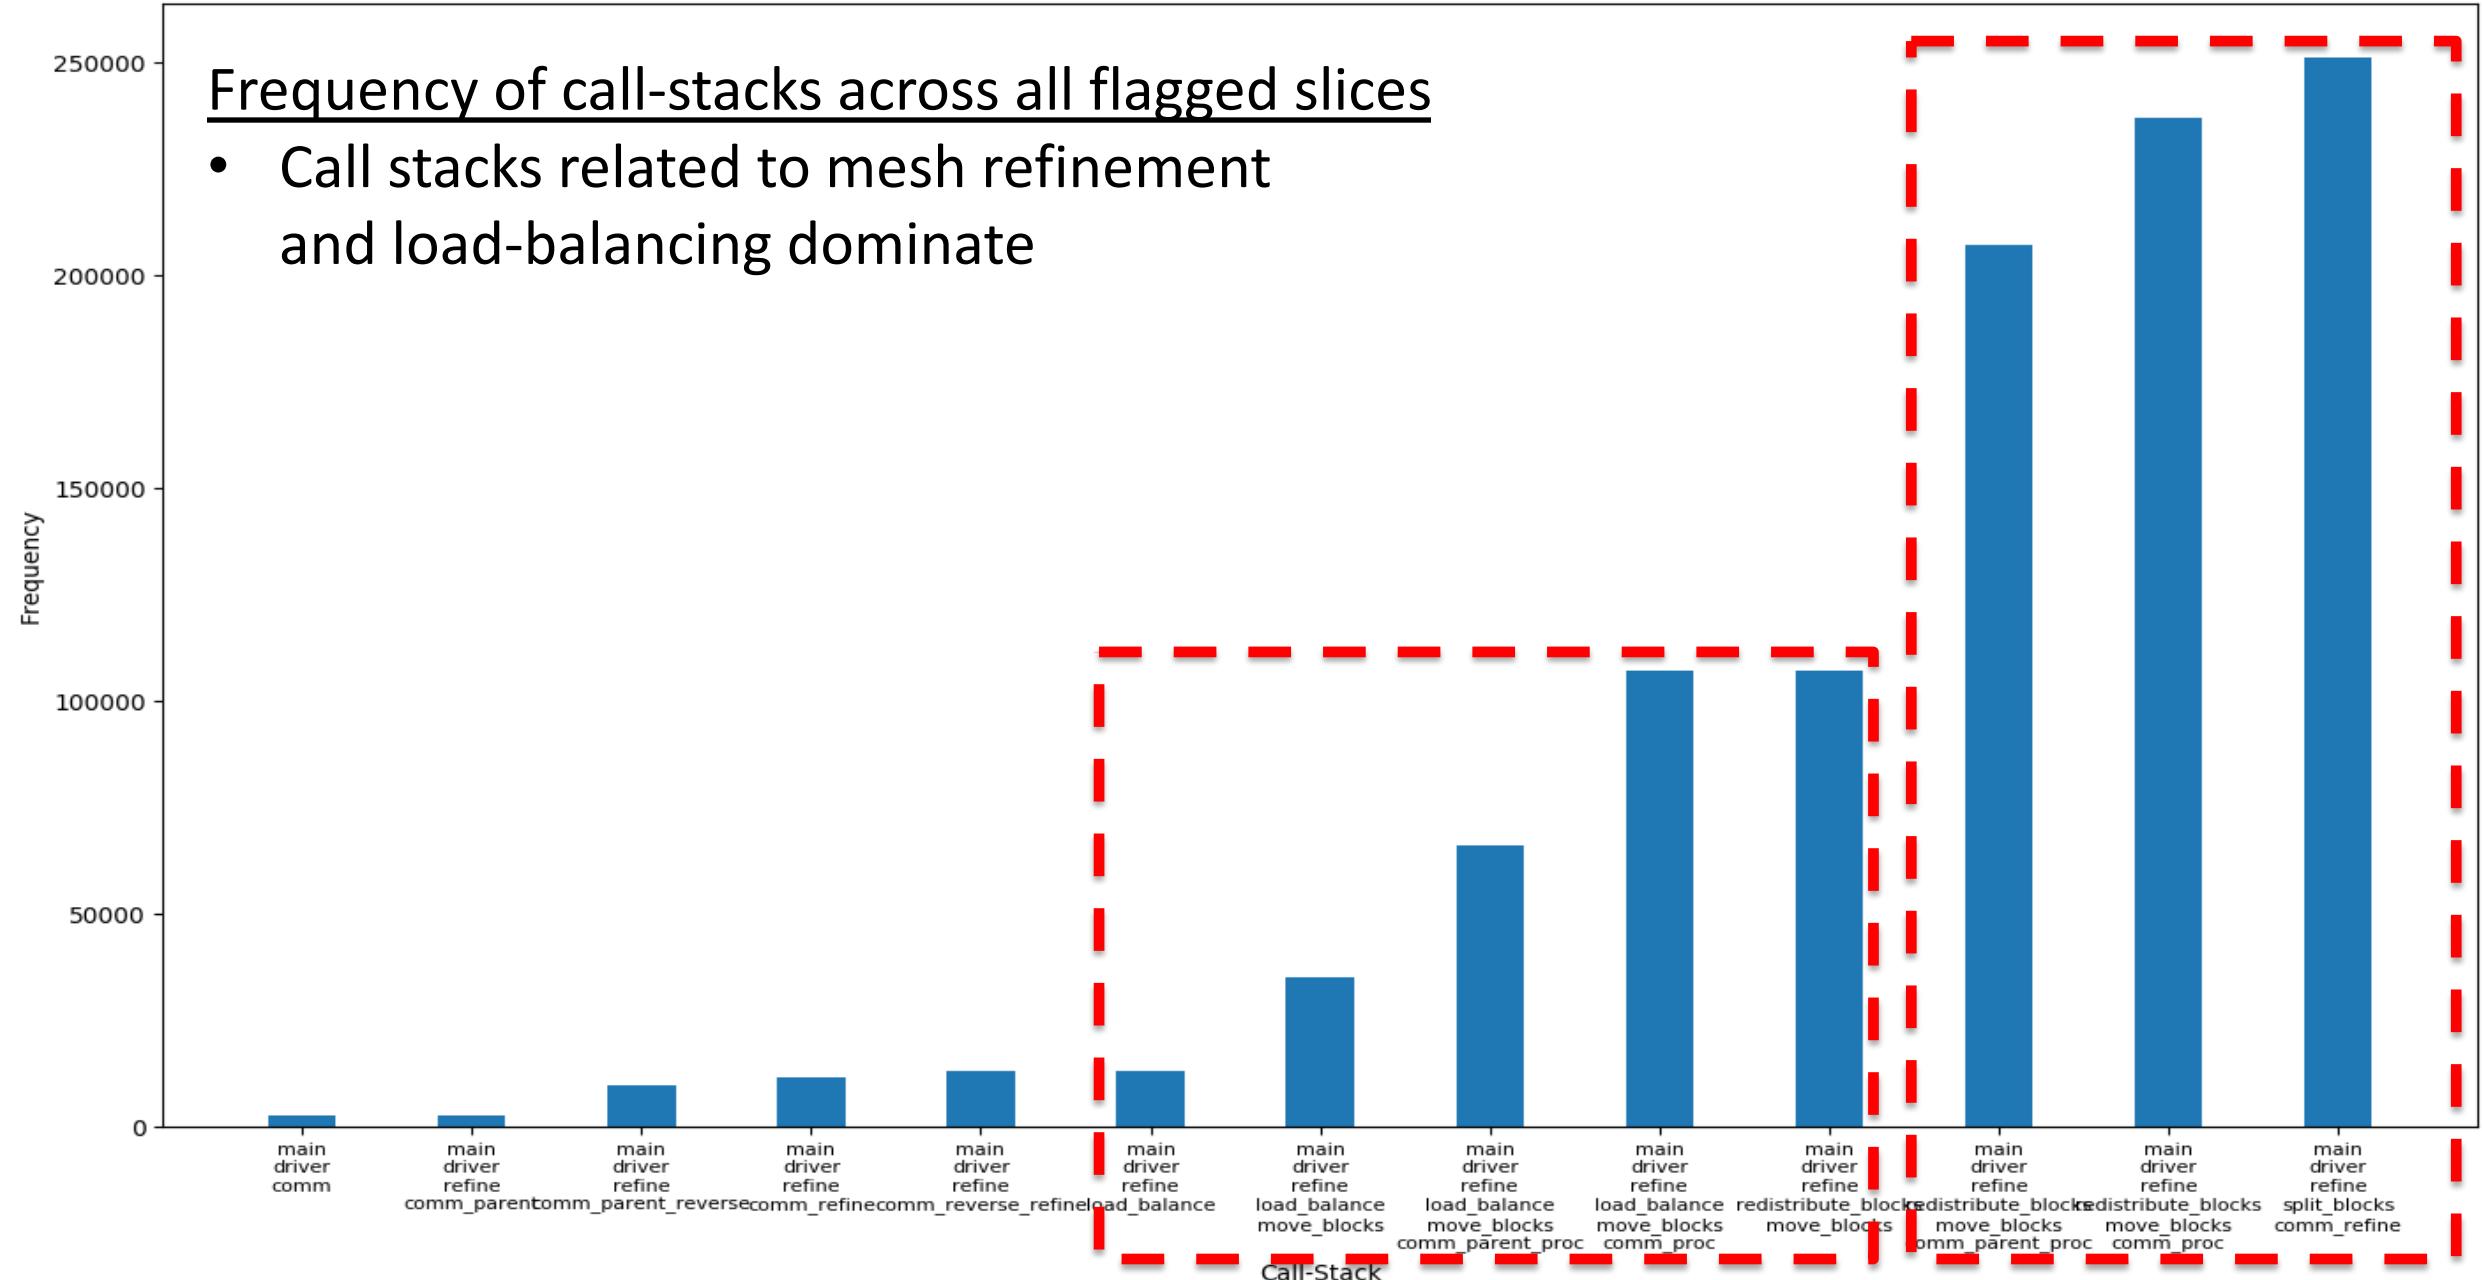


Linking to Potential Root Causes



Frequency of call-stacks across all flagged slices





Conclusion

- Our workflow identifies potential root causes of non-determinism in HPC applications by:
 - Building event graph models of executions
 - Quantifying trends in non-determinism via graph kernel distance
 - Linking runtime non-determinism to root causes

Conclusion

- Our workflow identifies potential root causes of non-determinism in HPC applications by:
 - Building event graph models of executions
 - Quantifying trends in non-determinism via graph kernel distance
 - Linking runtime non-determinism to root causes
- Demonstrated viability against:
 - Isolated communication patterns (CORAL-2 AMG)
 - Representative mini-app (miniAMR)

Conclusion

- Our workflow identifies potential root causes of non-determinism in HPC applications by:
 - Building event graph models of executions
 - Quantifying trends in non-determinism via graph kernel distance
 - Linking runtime non-determinism to root causes
- Demonstrated viability against:
 - Isolated communication patterns (CORAL-2 AMG)
 - Representative mini-app (miniAMR)
- Future Work:
 - Target full-fledged production AMR applications (e.g., Enzo)



Questions?



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE
BIG ORANGE. BIG IDEAS.®